

Go (Golang)

SDE2 Interview Guide

100 Questions & Answers with Code Examples

Level: Mid (2–4 yrs)	Stack: Go / Golang	Roles: SDE2 / Backend
----------------------	--------------------	-----------------------

#	Topic	Questions
1	Goroutines & Concurrency	Q1–Q15
2	Channels & Select	Q16–Q28
3	Runtime & Memory	Q29–Q40
4	Interfaces & Types	Q41–Q52
5	Error Handling	Q53–Q60
6	Context	Q61–Q67
7	Generics	Q68–Q74
8	Testing & Benchmarking	Q75–Q82
9	Go Patterns	Q83–Q92
10	Miscellaneous & Advanced	Q93–Q100

> **Focus:** Cache Fundamentals, Redis Internals, Eviction, Patterns, Distributed Caching, Go Integration | **Level:** SDE2

Table of Contents

[Redis Data Structures & Internals](#2-redis-data-structures--internals) — Q21–Q40
 [Eviction Policies & TTL](#3-eviction-policies-ttl) — Q41–Q50
 [Cache Patterns & Strategies](#4-cache-patterns--strategies) — Q51–Q65
 [Distributed Caching & Redis Cluster](#5-distributed-caching--redis-cluster) — Q66–Q80
 [Advanced Topics & Go Integration](#6-advanced-topics--go-integration) — Q81–Q100

Cache Fundamentals

Q1. What is a cache and why do we use it?

Difficulty: Easy

A cache stores copies of data in fast-access storage to reduce latency and load on slower backing stores (databases, APIs, disk).

Without cache:

User → App → DB (50ms) → User

With cache:

User → App → Cache (0.5ms) → User

↓ miss

App → DB (50ms) → Cache → User

Benefits: Reduce DB load, lower latency (sub-millisecond vs tens of ms), absorb traffic spikes, reduce cost (fewer DB ops).

When NOT to cache: Financial balances (stale = wrong), user-specific secrets, rapidly changing data with strict consistency requirements.

Q2. What are cache hit, miss, and hit ratio?

Difficulty: Easy

Cache hit: request served from cache (fast, cheap)

Cache miss: cache doesn't have it → fetch from DB → populate cache

Hit ratio: $\text{hits} / (\text{hits} + \text{misses}) \times 100\%$

Target hit ratio: >90% for most production caches

If hit ratio < 80%: cache is not helping much, reconsider strategy

Monitoring in Redis:

```
redis-cli INFO stats | grep keyspace
```

```
# keyspace_hits: 10000000
```

```
# keyspace_misses: 100000
```

```
# hit_ratio = 10M / (10M + 100K) ≈ 99%
```

Cold start: after deploy/restart, cache is empty → all misses

→ pre-warm critical keys before switching traffic

→ or accept brief latency spike on startup

Q3. What is cache latency and how does it compare to DB?

Difficulty: Easy

Storage latency comparison:

CPU L1 cache: 0.5 ns

CPU L2 cache: 7 ns

RAM: 100 ns

Redis (local): 0.1–0.5 ms ← in-process or same host

Redis (remote): 1–5 ms ← cross-AZ

```

SSD (local):    0.1 ms
PostgreSQL:    5-50 ms      ← indexed query
PostgreSQL:    50-500 ms   ← complex query / cold cache
HDD:           5-10 ms
S3:            20-100 ms

```

Rule of thumb: Redis $\approx$ 10-100 $\times$ faster than DB for cached reads

At 1M QPS: 50ms DB per query = 50,000 DB-seconds/sec = 50,000 cores needed

0.5ms Redis per query = 500 Redis-cores needed

Q4. What are the layers of caching in a web application?

Difficulty: Easy

```

Browser cache    → HTML, JS, CSS, images (304 Not Modified, Cache-Control)
CDN edge cache   → static assets, public API responses (global PoPs)
API Gateway cache → response caching per route
Application cache → in-process (sync.Map, groupcache, bigcache)
Distributed cache → Redis / Memcached (shared across instances)
DB query cache   → PostgreSQL shared_buffers, buffer pool
OS page cache    → file system reads
Hardware cache   → CPU + SSD caches

```

In-process vs distributed:

In-process: $\sim$ 0ms (no network), not shared, lost on restart

Distributed: $\sim$ 1ms (network), shared, survives restarts

L1 (in-process) + L2 (Redis) = best of both worlds

Q5. What is the difference between cache-aside, read-through, and write-through?

Difficulty: Medium

Cache-Aside (Lazy Loading):

Read: check cache $\rightarrow$ miss $\rightarrow$ read DB $\rightarrow$ store in cache $\rightarrow$ return

Write: write to DB $\rightarrow$ delete/update cache

Pros: cache only what's needed, resilient to cache failure

Cons: first read always cold, possible stale data

Read-Through:

App always reads from cache

Cache fetches from DB on miss automatically (cache handles it)

Pros: transparent to app

Cons: cold start, cache lib must know DB schema

Write-Through:

App writes to cache $\rightarrow$ cache synchronously writes to DB

Pros: cache always up to date

Cons: write latency (2 writes), cache stores data never read

Write-Behind (Write-Back):

Write to cache only $\rightarrow$ async flush to DB

Pros: very fast writes, absorbs bursts

Cons: data loss if cache crashes before flush

Most common production: Cache-Aside reads + Delete on writes

Q6. What is cache invalidation and why is it hard?

Difficulty: Medium

"There are only two hard things in Computer Science:
cache invalidation and naming things." — Phil Karlton

Invalidation strategies:

1. TTL (Time-To-Live): expire after N seconds
Simple, always eventually consistent
Downside: stale for up to TTL duration
2. Delete on write: when DB is updated, delete cache key
Consistent (cache repopulated on next read)
Downside: thundering herd on popular keys
3. Event-driven: publish cache invalidation event to Kafka/Redis Pub/Sub
Other services subscribe and delete local cache
Downside: async, brief inconsistency window
4. Versioned keys: key = "user:42:v7" where v7 = current version
Old versions auto-expire via TTL
Downside: old versions waste memory until TTL

Best practice: delete key on write (not update) + short TTL as safety net
Never update cached values directly – race conditions

Q7. What is the thundering herd problem?

Difficulty: Hard

Thundering herd (cache stampede):

Popular key expires → 1000 concurrent requests → all miss
→ 1000 requests hit DB simultaneously → DB overload → slow → more misses

Solutions:

1. Mutex / single-flight (most common):
Only first requester fetches from DB; others wait for result
Go: golang.org/x/sync/singleflight


```
var group singleflight.Group
result, err, _ := group.Do(key, func() (interface{}, error) {
    return db.GetUser(id)
})
```
2. Jittered TTL:
TTL = base_ttl + random(0, jitter) e.g., 300s ± 30s
Prevents all keys expiring simultaneously
Simple, always use this
3. Background refresh (proactive):
Before key expires: background goroutine refreshes it
Client always gets cached value (no miss latency)
4. Stale-while-revalidate:
Serve stale data immediately
Trigger async background refresh
Client gets fast (possibly old) response

Production: singleflight + jittered TTL + background refresh for hottest keys

Q8. What is cache warming and why is it important?**Difficulty:**
Medium

Cache warming: pre-populating cache before traffic hits
Prevents cold start performance degradation on deploys

When you need it:

- After Redis flush/restart
- New cluster node added
- After application deploy (if in-process cache cleared)
- After major schema change

Strategies:

1. Lazy warm (natural): first requests populate cache
Simple but first N requests are slow
OK for non-critical paths
2. Explicit pre-warm: load popular keys before serving traffic
Read top 10K products from DB → SET in Redis
Run before deploy completes
3. Replay warm: copy cache from production to staging
RDB snapshot → load into new Redis
4. Traffic replay: send shadow traffic before cutover

```
// Go: warm on startup
func warmCache(ctx context.Context, db *sql.DB, redis *redis.Client) error {
    products, _ := db.QueryContext(ctx, "SELECT id, data FROM products ORDER BY views DESC LIMIT 10000")
    pipe := redis.Pipeline()
    for products.Next() { pipe.Set(ctx, key, data, 10*time.Minute) }
    _, err := pipe.Exec(ctx)
    return err
}
```

Q9. What is a cache stampede vs a hotspot?**Difficulty:**
Medium

Cache stampede: many simultaneous misses for ONE expired key
→ burst of DB requests for same key
Fix: singleflight, mutex, background refresh

Hotspot: one key accessed far more than others
→ single Redis instance handles all traffic for that key
→ single-threaded Redis gets overloaded for that key

Example: celebrity tweet being read by 1M followers/sec

Hotspot solutions:

- Local in-process cache (L1): each app server caches hot key locally
→ 0ms read, no Redis load
→ short TTL (seconds) for freshness

Key replication: replicate hot key to multiple Redis slots
key:copy1, key:copy2, key:copy3 → read-only replicas
Route reads round-robin across copies

Redis read replicas: add Redis replica nodes for read-only scale
Primary: writes only
Replicas: reads only, replicate from primary

Q10. What is cache penetration?

Difficulty:
Medium

Cache penetration: querying for keys that don't exist in DB
→ every request misses cache (no value to cache)
→ every request hits DB
→ can be exploited as DoS attack (flood with random IDs)

Solutions:

1. Cache null values:

Key doesn't exist in DB → cache null with short TTL (30s)
SET user:99999 "null" EX 30
Next request: cache hit → null → return 404
Downside: wasted cache space

2. Bloom filter:

Probabilistic set membership: "is key 99999 in DB?"
No false negatives: if Bloom says NO → definitely not in DB → skip DB
False positives possible (0.1%): might check DB when not needed
Space-efficient: 10 bits per element at 1% false positive rate

Check Bloom filter → "definitely not" → return 404 (no DB hit)

Redis: RedisBloom module (BF.ADD, BF.EXISTS)

3. Rate limiting by IP/user:

Limit requests per key pattern to prevent DoS

Production: Bloom filter for large key spaces + cache nulls for small

Q11. What is a Bloom filter?

Difficulty: Hard

Bloom filter: space-efficient probabilistic set membership test
"Is element X in the set?" → "Definitely NO" or "Probably YES"
No false negatives, possible false positives

How it works:

Bit array of M bits, K hash functions
Add element: set bits at positions $h_1(x)$, $h_2(x)$, ..., $h_k(x)$
Query: check if ALL k positions are 1
All 1 → probably in set (false positive possible)
Any 0 → definitely NOT in set

Properties:

Cannot delete (use Counting Bloom Filter for deletions)
False positive rate $\approx (1 - e^{(-kn/m)})^k$
At 1% FPR: ~9.6 bits per element (1M elements = 1.2 MB)
At 0.1% FPR: ~14.4 bits per element

Applications:

Cache penetration: skip DB for keys not in Bloom filter
Cassandra: per-SSTable Bloom filter to avoid reading wrong file

Chrome: malicious URL detection
HBase, RocksDB: avoid disk reads for non-existent keys

```
// Go:
import "github.com/bits-and-blooms/bloom/v3"
filter := bloom.New(1000000, 5) // 1M items, 5 hash functions
filter.Add([]byte("user:42"))
exists := filter.Test([]byte("user:99999")) // false → skip DB
```

Q12. What is cache consistency and how do you handle it?

Difficulty: Hard

Cache consistency: cache reflects latest DB state

Levels of consistency:

Strong: cache always matches DB (too expensive)
Eventual: cache converges to DB state within TTL
Read-your-writes: user sees their own writes immediately

Practical approaches:

1. Write-through: update cache and DB atomically
Problems: write latency, race conditions
2. Delete-on-write: DB write → delete cache key
Next read: repopulate from DB
Safer than updating (avoids race conditions)

Race condition in cache update:

Thread1: read DB(v1) → Thread2: write DB(v2), delete cache
→ Thread1: write cache(v1) → STALE!

Race condition in cache delete:

Thread1: read DB(v1) → Thread2: write DB(v2), delete cache
→ Thread1: write cache(v1) → Thread3: reads cache(v1) STALE
But: solved by short TTL as safety net

3. TTL safety net: even if invalidation misses, TTL ensures eventual consistency
4. Cache-aside with version check:
Store version with cached data
Compare with DB version on read → invalidate if mismatch

Q13. What is the write-around pattern?

Difficulty: Easy

Write-Around: writes go directly to DB, bypassing cache
Cache populated lazily on first read

When to use:

Data written once, rarely read (logs, audit events)
Prevents cache pollution (caching data nobody will ever read)

Analogy: bulk import – don't cache 1M imported rows nobody will read

vs Write-Through:

Write-Through: every write → cache + DB (all writes cached)

Write-Around: write → DB only (only reads populate cache)

Implementation:

```
// Write-Around:
func createAuditLog(log AuditLog) error {
    return db.Insert(log) // no cache write
}

// Read (Cache-Aside, lazy population):
func getAuditLog(id string) (*AuditLog, error) {
    if v := cache.Get(id); v != nil { return v, nil }
    log := db.Get(id)
    cache.Set(id, log, 5*time.Minute)
    return log, nil
}
```

Q14. What are the trade-offs between in-process vs distributed cache?

Difficulty:
Medium

	In-Process	Distributed (Redis)
Latency	~0ms (no network)	1-5ms (network hop)
Shared	No (per-instance)	Yes (all instances)
Consistency	Each instance diverges	Single source of truth
Capacity	Limited by JVM/Go heap	Terabytes possible
Eviction	GC pressure	No GC, LRU/LFU
Persistence	Lost on restart	Optional (RDB/AOF)
Complexity	Simple	Connection management
Cost	Free (RAM already there)	Extra infrastructure
Best for	Hot computation, L1 cache	Session state, rate limiting

Recommended pattern: L1 (in-process) + L2 (Redis)
 L1: last 1K hot keys, 5s TTL (very short for freshness)
 L2: all cached data, minutes/hours TTL
 Cache L2 result in L1 on first access
 0ms for L1 hits, 1ms for L2 hits, 10ms for DB misses

Q15. What is a session cache and how do you implement it?

Difficulty: Easy

Session cache: store user session data in Redis
 Key: session:{session_id}
 Value: JSON of session data (user_id, roles, preferences)
 TTL: 24 hours (or session duration)

Advantages over DB sessions:
 Sub-millisecond lookup (every authenticated request)
 Easy horizontal scaling (any server reads from Redis)
 TTL handles expiry automatically

Implementation:

```
// Store session
sessionID := uuid.New().String()
sessionData, _ := json.Marshal(Session{UserID: 42, Roles: []string{"admin"}})
redis.Set(ctx, "session:"+sessionID, sessionData, 24*time.Hour)

// Retrieve session
data, err := redis.Get(ctx, "session:"+sessionID).Result()
var session Session
```



```

json.Unmarshal([]byte(data), &session)

// Extend on activity
redis.Expire(ctx, "session:"+sessionID, 24*time.Hour)

// Logout
redis.Del(ctx, "session:"+sessionID)

Security:
  Session ID = cryptographically random UUID
  Store in HttpOnly + Secure cookie
  HTTPS only

```

Q16. What is a rate limiter using Redis?

Difficulty:
Medium

Rate limiter: limit requests per user/IP per time window

Sliding window counter (most common):

```

// Token bucket with Redis (atomic via Lua script)
local count = redis.call('INCR', KEYS[1])
if count == 1 then
    redis.call('EXPIRE', KEYS[1], ARGV[1])
end
return count

// Go:
func isAllowed(ctx context.Context, rdb *redis.Client, key string, limit int, window time.Duration) (bool, error) {
    luaScript := redis.NewScript(`
        local count = redis.call('INCR', KEYS[1])
        if count == 1 then redis.call('EXPIRE', KEYS[1], ARGV[1]) end
        return count
    `)
    count, err := luaScript.Run(ctx, rdb, []string{key},
        int(window.Seconds())).Int()
    return count <= limit, err
}

// Usage:
allowed, _ := isAllowed(ctx, rdb, "rate:ip:1.2.3.4", 100, time.Minute)
if !allowed { return http.StatusTooManyRequests }

// Response headers:
X-RateLimit-Limit: 100
X-RateLimit-Remaining: 45
X-RateLimit-Reset: 1609459200

```

Q17. What is a distributed lock using Redis?

Difficulty: Hard

Distributed lock: ensure only one process executes a critical section
Across multiple machines (unlike sync.Mutex which is process-local)

Redis SETNX-based lock:

```

// Acquire: SET key value NX PX ttl_ms (atomic!)
// NX = only set if not exists
// PX = expire in milliseconds

```

```
func acquireLock(ctx context.Context, rdb *redis.Client, resource string, ttl time.Duration) (string, bool) {
    token := uuid.New().String()
    ok, _ := rdb.SetNX(ctx, "lock:"+resource, token, ttl).Result()
    return token, ok
}
```

```
// Release: Lua script – only delete if WE own it
func releaseLock(ctx context.Context, rdb *redis.Client, resource, token string) {
    luaScript := redis.NewScript(`
        if redis.call('GET', KEYS[1]) == ARGV[1] then
            return redis.call('DEL', KEYS[1])
        end
        return 0
    `)
    luaScript.Run(ctx, rdb, []string{"lock:" + resource}, token)
}
```

Redlock (multi-node):

- Acquire lock on $N/2+1$ Redis nodes (majority)
- Use when single Redis node failure is unacceptable

Library: github.com/bsm/redislock (recommended)

Q18. What is cache-aside with singleflight in Go?

Difficulty:
Medium

```
import "golang.org/x/sync/singleflight"

type UserCache struct {
    rdb    *redis.Client
    db     *sql.DB
    group singleflight.Group
}

func (c *UserCache) GetUser(ctx context.Context, id int64) (*User, error) {
    // Check cache first
    key := fmt.Sprintf("user:%d", id)
    data, err := c.rdb.Get(ctx, key).Bytes()
    if err == nil {
        var u User
        return &u, json.Unmarshal(data, &u)
    }
    if !errors.Is(err, redis.Nil) {
        return nil, err
    }

    // Cache miss: singleflight deduplicates concurrent DB requests
    result, err, _ := c.group.Do(key, func() (interface{}, error) {
        var u User
        if err := c.db.QueryRowContext(ctx,
            "SELECT * FROM users WHERE id=$1", id).Scan(&u); err != nil {
            return nil, err
        }
        data, _ := json.Marshal(u)
        // Jittered TTL prevents stampede on popular keys
        ttl := 5*time.Minute + time.Duration(rand.Intn(60))*time.Second
        c.rdb.Set(ctx, key, data, ttl)
    })
    return result.(*User), err
}
```

```

        return &u, nil
    })
    if err != nil { return nil, err }
    return result.(*User), nil
}

```

Q19. What is content-addressable caching?

Difficulty:
Medium

Content-addressable: cache key = hash of content/request
 Same request parameters → same key → deterministic cache hit

Examples:

```

key = sha256(sql_query + params) → cache query result
key = sha256(template + data) → cache rendered HTML
key = sha256(image_url + dimensions) → cache resized image

```

Advantages:

- Natural deduplication: identical requests always hit same cache entry
- No need to manually construct cache keys
- Immutable content: can cache indefinitely (content never changes for same hash)

```

// Go:
import "crypto/sha256"

func queryKey(query string, args ...interface{}) string {
    h := sha256.New()
    fmt.Fprintf(h, "%s:%v", query, args)
    return fmt.Sprintf("query:%x", h.Sum(nil))
}

// CDN: Content-Addressable Storage
// Vite/Webpack: bundle.abc123.js (hash in filename → infinite cache TTL)
// Docker layers: content-addressed → layer shared between images

```

Q20. What are cache metrics you should monitor?

Difficulty: Easy

Key Redis metrics:

```
redis-cli INFO all | grep -E "keyspace|memory|clients|stats"
```

Memory:

```

used_memory: total memory used
used_memory_peak: peak memory
mem_fragmentation_ratio: >1.5 = fragmentation issue

```

Performance:

```

keyspace_hits: successful lookups
keyspace_misses: failed lookups
hit_ratio: hits / (hits + misses) → alert if < 90%
instantaneous_ops_per_sec: current OPS

```

Connections:

```

connected_clients: active connections
blocked_clients: clients waiting (BLPOP, etc.)

```

Eviction:

```
evicted_keys: keys removed due to memory limit
expired_keys: keys expired by TTL
```

Replication:

```
master_repl_offset vs slave_repl_offset → replication lag
```

Prometheus metrics: redis_exporter

Alert thresholds:

```
hit_ratio < 90% → cache not effective
used_memory > 80% of maxmemory → about to evict
evicted_keys increasing → memory pressure
replication lag > 100MB → replica falling behind
```

Redis Data Structures & Internals

Q21. What are Redis data structures and their use cases?

Difficulty: Easy

String: any binary-safe value up to 512MB

```
SET key value [EX seconds] [NX|XX]
INCR key → atomic counter
Use: cache values, counters, sessions, flags
```

List: doubly-linked list (fast head/tail ops)

```
LPUSH/RPUSH, LPOP/RPOP, LRANGE, LLEN
Use: queues (RPUSH+LPOP), recent items, timelines
```

Hash: field-value map within a key

```
HSET key field value, HGET key field, HGETALL
Use: user profiles, object fields (memory-efficient for small objects)
```

Set: unordered unique elements

```
SADD, SMEMBERS, SINTER, SUNION, SCARD, SISMEMBER
Use: unique visitors, tags, friends, online users
```

Sorted Set (ZSet): set ordered by float score

```
ZADD key score member, ZRANGE, ZREVRANGE, ZRANK, ZSCORE
Use: leaderboards, rate limiting, delayed queues, timelines
```

HyperLogLog: approximate cardinality estimation (~0.81% error)

```
PFADD, PFCOUNT, PFMERGE
Use: unique visitor count (1M+ items, tiny memory ~12KB)
```

Streams: append-only log with consumer groups

```
XADD, XREAD, XREADGROUP, XACK
Use: event sourcing, message queue with reply tracking
```

Bitmap: bit operations on strings

```
SETBIT, GETBIT, BITCOUNT, BITOP
Use: user activity tracking, feature flags per user, bloom filters
```

Q22. What is Redis Sorted Set and how does it work internally?

Difficulty: Hard

Sorted Set: elements with associated float score, ordered by score

Unique elements, non-unique scores allowed
 Supports: $O(\log N)$ insert/update, $O(\log N + M)$ range queries

Internal implementation:

Small (≤ 128 elements, ≤ 64 bytes each): ziplist (compact array)
 Large: skiplist + hashtable
 skiplist: $O(\log N)$ for most operations, $O(1)$ amortized insert
 hashtable: $O(1)$ ZSCORE lookup

Operations:

ZADD leaderboard 1500 "alice"	$O(\log N)$
ZSCORE leaderboard "alice"	$O(1)$
ZRANK leaderboard "alice"	$O(\log N)$ (position from low)
ZREVRANK leaderboard "alice"	$O(\log N)$ (position from high)
ZRANGE leaderboard 0 9	$O(\log N + M)$ top 10 by score
ZREVRANGE leaderboard 0 9 WITHSCORES	
ZRANGEBYSCORE lb 1000 2000	$O(\log N + M)$ score range
ZRANGEBYLEX lb "[a" "[z"	$O(\log N + M)$ lexicographic range
ZINCRBY leaderboard 10 "alice"	$O(\log N)$ atomic score increment
ZREM leaderboard "alice"	$O(\log N)$

Use cases:

Leaderboard: ZADD score player → ZREVRANGE top10
 Rate limiter: sorted set of timestamps per user
 Delayed queue: score = execute_at timestamp
 Timeline: score = created_at timestamp

Q23. What is the Redis Pub/Sub model?

Difficulty:
Medium

Redis Pub/Sub: fire-and-forget broadcast messaging
 No message persistence (missed if no subscriber)
 No acknowledgment

Commands:

```
SUBSCRIBE channel [channel ...]
PSUBSCRIBE pattern [pattern ...] (glob patterns)
PUBLISH channel message
UNSUBSCRIBE
```

Example (Go):

```
// Subscriber:
pubsub := rdb.Subscribe(ctx, "order-updates")
defer pubsub.Close()
for msg := range pubsub.Channel() {
    processUpdate(msg.Payload)
}

// Publisher:
rdb.Publish(ctx, "order-updates", jsonPayload)
```

Limitations:

Messages lost if subscriber is disconnected
 No message history / replay
 No consumer groups
 Fire-and-forget only

Use cases:

Real-time notifications (user online/offline)
 Cache invalidation signals (broadcast "invalidate user:42")
 Live dashboard updates

For reliable messaging: use Redis Streams (with XREADGROUP)

Q24. What are Redis Streams?

Difficulty: Hard

Redis Streams (Redis 5.0+): append-only log with consumer groups
 Persistent messages (unlike Pub/Sub)
 Consumer groups (multiple consumers, each gets different messages)
 Acknowledgment support

Commands:

XADD stream * field1 val1 field2 val2	append message
XREAD COUNT 10 STREAMS stream 0	read from beginning
XLEN stream	message count

Consumer groups:

XGROUP CREATE stream group1 \$	create group
XREADGROUP GROUP group1 consumer1 COUNT 10 STREAMS stream >	
XACK stream group1 message_id	acknowledge

// Go example:

```
rdb.XAdd(ctx, &redis.XAddArgs{
    Stream: "orders",
    Values: map[string]interface{}{"order_id": 42, "amount": "99.99"},
})
```

// Consumer group read

```
msgs, _ := rdb.XReadGroup(ctx, &redis.XReadGroupArgs{
    Group:    "processors",
    Consumer: "worker-1",
    Streams:  []string{"orders", ">"},
    Count:    10,
})
for _, msg := range msgs[0].Messages {
    processOrder(msg.Values)
    rdb.XAck(ctx, "orders", "processors", msg.ID)
}
```

vs Kafka: Redis Streams simpler but lower throughput, no partition-level parallelism

Q25. What is Redis persistence (RDB vs AOF)?

Difficulty: Medium

RDB (Redis Database Snapshot):

Periodic full snapshot to disk (binary format)
 Default: save every 60s if 10K changes (configurable)
 Fast restart (load snapshot in seconds)
 Risk: lose up to 60s of data on crash

```
# redis.conf
save 900 1    # save after 900s if at least 1 key changed
save 300 10   # save after 300s if at least 10 keys changed
save 60 10000 # save after 60s if at least 10000 keys changed
```

AOF (Append-Only File):

Log every write command
 appendfsync: always (safe, slow), everysec (balanced), no (fast, risky)
 Durability: lose at most 1s of data (everysec)
 Larger files, slower restart (replay all commands)

```
appendonly yes
appendfsync everysec # recommended: sync to disk every second
```

Both (hybrid, recommended):

Load RDB on startup (fast)
 Use AOF for ongoing durability
 aof-use-rdb-preamble yes # AOF file starts with RDB snapshot

For cache-only (no persistence):

Comment out all save lines
 appendonly no
 → Fastest, data lost on restart (acceptable for pure cache)

Q26. What is Redis memory optimization?**Difficulty: Hard****Memory-efficient encodings (automatic):**

Small hash (≤ 128 fields, ≤ 64 bytes each) → ziplist (compact)
 Small sorted set → ziplist
 Small set → intset (integers) or listpack
 Small list → listpack

Configure thresholds:

```
hash-max-listpack-entries 128
hash-max-listpack-value 64
zset-max-listpack-entries 128
zset-max-listpack-value 64
```

Memory savings:

ziplist vs hashtable: 4-10× less memory for small objects
 intset vs hashtable: 2-5× less memory for integer sets

Compression:

```
no: no compression (default)
lz4: compress RDB and AOF
```

Key naming:

Short keys: "u:42:sess" vs "user:42:session" → saves memory
 Shared prefix in hash: HSET u:42 name "Alice" vs SET u:42:name "Alice"
 Hash approach: 10× more memory efficient for many small values

Memory analysis:

```
redis-cli --bigkeys # find largest keys
redis-cli MEMORY USAGE key # bytes for specific key
redis-cli OBJECT ENCODING key # current encoding
redis-cli DEBUG JMAP # memory breakdown
redis-cli SCAN 0 MATCH "user:*" COUNT 100 # iterate keys safely
```

Q27. What is Redis pipelining?**Difficulty: Medium**

Pipelining: send multiple commands in one network round-trip

Reduces RTT overhead dramatically
 Commands batched, sent together, responses received together

Without pipelining:

GET key1 → wait 1ms → GET key2 → wait 1ms → GET key3 → wait 1ms = 3ms total

With pipelining:

[GET key1, GET key2, GET key3] → send once → receive all → 1ms total

```
// Go with go-redis:
pipe := rdb.Pipeline()
get1 := pipe.Get(ctx, "key1")
get2 := pipe.Get(ctx, "key2")
pipe.Set(ctx, "key3", "value", time.Hour)
pipe.Exec(ctx)

val1, _ := get1.Result()
val2, _ := get2.Result()

// TxPipeline: pipeline wrapped in MULTI/EXEC (transaction)
txPipe := rdb.TxPipeline()
// ... commands ...
txPipe.Exec(ctx)
```

When to use:

Reading/writing many keys at once
 Batch operations (populate cache, read multiple users)

When NOT to use:

Commands depend on previous result (need scripting/transactions)
 Very low volume (overhead not worth it)

Q28. What is a Redis Lua script and when do you use it?

Difficulty: Hard

Lua scripts: execute multiple commands atomically on Redis
 Server-side execution: no network round-trips between commands
 Atomic: no other command runs between script commands
 Use: compound operations requiring atomicity

```
// Atomic check-and-set (compare-and-swap):
local current = redis.call('GET', KEYS[1])
if current == ARGV[1] then
    redis.call('SET', KEYS[1], ARGV[2])
    return 1
end
return 0

// Go:
script := redis.NewScript(`
    local current = redis.call('GET', KEYS[1])
    if current == ARGV[1] then
        redis.call('SET', KEYS[1], ARGV[2])
        return 1
    end
    return 0
`)
result, err := script.Run(ctx, rdb, []string{"mykey"}, "expected", "new_value").Int()
```



```
// Rate limiter (atomic):
local count = redis.call('INCR', KEYS[1])
if count == 1 then redis.call('EXPIRE', KEYS[1], ARGV[1]) end
return count

// Lock release (only if owner):
if redis.call('GET', KEYS[1]) == ARGV[1] then
    return redis.call('DEL', KEYS[1])
end
return 0
```

Caching scripts: EVALSHA (cache by SHA1 hash) vs EVAL
 SCRIPT LOAD script → returns SHA1 → use EVALSHA sha key args

Q29. What is Redis MULTI/EXEC (transactions)?

Difficulty:
Medium

MULTI/EXEC: queue commands, execute atomically
 Not a full ACID transaction (no rollback on error)
 Isolation: commands execute without interruption
 All-or-nothing execution: if EXEC called, all queued commands run

Behavior:

```
MULTI          → enter transaction mode (queue commands)
command1       → queued (not executed yet)
command2       → queued
EXEC           → execute all queued commands atomically
DISCARD        → cancel transaction
WATCH key      → optimistic locking (abort if key changed)
```

Redis transaction limitations:

No rollback on runtime errors (command already ran)
 If SET fails mid-transaction: previous SETs not undone
 Not like SQL transactions!

WATCH for optimistic locking:

```
WATCH account:42
value = GET account:42
MULTI
SET account:42 (value - 100)
result = EXEC
if result == nil: WATCH detected change, retry
```

// Go:

```
_, err = rdb.TxPipelined(ctx, func(pipe redis.Pipeliner) error {
    pipe.Set(ctx, "key1", "val1", 0)
    pipe.Set(ctx, "key2", "val2", 0)
    return nil
})
// Either both set or neither (on connection failure)
```

Q30. What is Redis single-threaded model?

Difficulty:
Medium

Redis: single-threaded command processing (mostly)
 One goroutine handles all commands sequentially
 No locking needed for data structures

No context switching overhead

Why fast despite single-thread:

- All data in RAM (no disk I/O in critical path)
- Simple data structures ($O(1)$ or $O(\log N)$)
- Minimal memory allocation (pre-allocated buffers)
- OS socket I/O using epoll/kqueue (non-blocking)

Multi-threading in Redis 6.0+:

- I/O threads: multiple threads for reading/writing sockets
- Command execution: still single-threaded
- Result: 2-3× throughput improvement on multi-core servers

Implications:

- Slow commands (KEYS *, SMEMBERS on huge set, SORT) block ALL clients
- Avoid $O(N)$ commands on large data sets
- Use SCAN instead of KEYS (non-blocking iteration)
- Use SSCAN/HSCAN/ZSCAN for large sets/hashtables/zsets

Benchmarks:

- Single Redis node: ~100K-1M simple ops/sec
- With Redis 6.0 I/O threads: up to 2M+ ops/sec

Q31. #### Q31-Q40: Additional Redis Topics

Difficulty:
Medium

Q	Topic
Q31	Redis memory maxmemory and its behavior when limit hit
Q32	Redis key expiry: lazy expiry vs active expiry
Q33	Redis OBJECT ENCODING: ziplist vs skiplist vs listpack
Q34	Redis replication: REPLICCONF, PSYNC, partial resync
Q35	Redis Sentinel: automatic failover configuration
Q36	Redis keyspace notifications for expired/modified keys
Q37	Redis SCAN vs KEYS: why KEYS is dangerous in production
Q38	Redis GEO commands: GEOADD, GEODIST, GEOSEARCH
Q39	Redis OBJECT FREQ for LFU eviction tracking
Q40	Redis ACL (Access Control Lists) for security

Eviction Policies & TTL

Q32. What are Redis eviction policies?

Difficulty:
Medium

When maxmemory is reached, Redis uses eviction policy:

noeviction (default):

- Return error on write commands when full
- Use for: queues, lists where data loss is unacceptable

allkeys-lru:

- Evict least recently used key from ALL keys
- Best for: pure cache (all keys eligible for eviction)

volatile-lru:

- Evict LRU key from keys WITH TTL set
- Use: mixed use case (cache + persistent data in same Redis)

allkeys-lfu (Redis 4.0+):

- Evict least frequently used from ALL keys
- Better than LRU when some keys accessed infrequently but recently

volatile-lfu:

- Evict LFU from keys with TTL

allkeys-random:

- Random eviction from all keys

volatile-random:

- Random from keys with TTL

volatile-ttl:

- Evict key with shortest TTL first
- Use: prioritize evicting keys about to expire anyway

Recommended for pure cache: allkeys-lru or allkeys-lfu

allkeys-lfu: better when some keys are accessed rarely but recently inserted

Q33. What is LRU and how is it implemented efficiently?

Difficulty: Hard

LRU (Least Recently Used): evict key not accessed for longest time

Naive LRU: $O(1)$ lookup (hash map) + $O(1)$ access order update (doubly-linked list)

- Hash map: key $\rightarrow$ node in linked list

- Linked list: MRU end ... LRU end

- On access: move node to MRU end

- On evict: remove from LRU end

Redis approximation (not exact LRU):

- Redis 3.0+: pool of eviction candidates

- Sample N random keys (maxmemory-samples, default 5)

- Evict least recently used among sample

- Approximation is 95%+ accurate vs true LRU

- Much faster and less memory than true LRU

LRU vs LFU:

- LRU: good for temporal locality (recently used = likely to be reused)

- LFU: good for frequency-based (popular keys always in cache)

- LRU problem: long-running low-frequency scans evict hot keys

- LFU problem: new popular key starts with frequency=1 (might be evicted!)

LFU solution: decay counter (frequency halved every time period)

```
// Go: implement O(1) LRU cache
import "github.com/hashicorp/golang-lru/v2"
cache, _ := lru.New[string, User](10000)
cache.Add("user:42", user)
if v, ok := cache.Get("user:42"); ok { ... }
```

Q34. What is TTL strategy and jitter?

Difficulty:
Medium

TTL (Time To Live): key automatically deleted after N seconds

Setting TTL:

```
SET key value EX 300          # 300 seconds
SET key value PX 300000       # 300000 milliseconds
SET key value EXAT 1700000000 # Unix timestamp
EXPIRE key 300                # set TTL on existing key
PERSIST key                   # remove TTL (make permanent)
TTL key                       # seconds remaining (-1=no TTL, -2=doesn't exist)
```

TTL strategy:

Short TTL (seconds): frequently changing data (stock prices, sessions)

Medium TTL (minutes): user profiles, product catalog

Long TTL (hours/days): rarely changing data, expensive to compute

Jitter: add randomness to prevent synchronized expiration

Without jitter: all keys set at 3:00 PM expire at 3:05 PM → stampede

With jitter: TTL = base + rand(0, jitter)

```
// Go:
baseTTL := 5 * time.Minute
jitter := time.Duration(rand.Intn(60)) * time.Second
rdb.Set(ctx, key, value, baseTTL+jitter)
```

Sliding TTL: reset TTL on access (session-style)

On each access: EXPIRE key 30min → extends by 30 min from now

Stays alive as long as accessed, expires when idle

Q35. What is the difference between EXPIRE and EXPIREAT?

Difficulty: Easy

```
EXPIRE key seconds      : expire in N seconds from now
PEXPIRE key milliseconds : expire in N milliseconds from now
EXPIREAT key timestamp  : expire at Unix timestamp (seconds)
PEXPIREAT key timestamp  : expire at Unix timestamp (milliseconds)
```

Use cases:

EXPIRE 3600: cache for 1 hour from now

EXPIREAT (midnight): expire at end of day

Atomicity with SET:

SET key value EX 3600 ← atomic set + expire

vs

SET key value ← not atomic (race condition if crash between SET and EXPIRE)

EXPIRE key 3600 ← use SET EX form instead

TTL introspection:

TTL key → remaining seconds (-1 = no TTL, -2 = doesn't exist)
 PTTL key → remaining milliseconds

Change expiry without changing value:

EXPIRE key 7200 → extend to 2 hours from now

EXPIRETIME key → returns absolute expiry timestamp (Redis 7.0+)

Q36. What is cache TTL best practice per data type?

Difficulty:
Medium

TTL guidelines by data volatility:

Immutable data (5 min - 24 hours or no TTL):

Product images, static config, country list

SET product:img:42 url TTL=24h

User-specific data (15 min - 4 hours):

User profile, preferences, permissions

SET user:42:profile json TTL=1h

Computed/aggregated data (5 min - 1 hour):

Leaderboard, stats, analytics summaries

ZADD leaderboard score user → rebuild every 5 min via cron

Real-time data (seconds to minutes):

Feed, notifications, online status

SET user:42:online 1 EX 30 # heartbeat every 15s

Session data (sliding, 30 min - 24 hours):

Reset TTL on each request

EXPIRE session:xyz 3600

Rate limits (per window):

TTL = window duration

SET ratelimit:ip:1.2.3.4:2024010112 0 EX 3600

Never cache:

Financial balances, active transactions, security tokens

Anything requiring strong consistency

Q37. ### Q46-Q50: More Eviction Topics

Difficulty:
Medium

Q	Topic
Q46	maxmemory-policy selection guide for different workloads
Q47	volatile-lru vs allkeys-lru: when to use each
Q48	Redis memory fragmentation and how to reduce it
Q49	TTL precision: EXPIRE vs PEXPIRE differences
Q50	Cache size estimation: how much memory do you need?

Cache Patterns & Strategies

Q38. What is the cache-aside pattern in detail?

Difficulty:
Medium

Most common caching pattern:

Read path:

1. App checks cache (GET key)
2. Hit → return value
3. Miss → fetch from DB → SET in cache → return value

Write path (option A: delete-on-write – safer):

1. Write to DB
2. DELETE cache key
3. Next read will repopulate from DB

Write path (option B: update-on-write – riskier):

1. Write to DB
 2. SET cache key with new value
- Race condition: concurrent write may set stale value

// Go implementation:

```
func (r *Repo) GetUser(ctx context.Context, id int64) (*User, error) {
    key := fmt.Sprintf("user:%d", id)

    // Check cache
    data, err := r.cache.Get(ctx, key).Bytes()
    if err == nil {
        var u User
        return &u, json.Unmarshal(data, &u)
    }

    // DB fallback
    u, err := r.db.GetUser(ctx, id)
    if err != nil { return nil, err }

    // Populate cache (ignore error – cache is best-effort)
    b, _ := json.Marshal(u)
    r.cache.Set(ctx, key, b, 5*time.Minute+jitter())
    return u, nil
}

func (r *Repo) UpdateUser(ctx context.Context, u *User) error {
    if err := r.db.UpdateUser(ctx, u); err != nil { return err }
    r.cache.Del(ctx, fmt.Sprintf("user:%d", u.ID)) // invalidate
    return nil
}
```

Q39. What is the write-through pattern?

Difficulty:
Medium

Write-through: every write goes to cache AND DB synchronously

1. App writes to cache → cache writes to DB (or app writes both)
2. Both succeed → return success

3. Read: always hits cache (never stale after write)

Pros:

- Cache always consistent with DB
- Read never misses for recently written data

Cons:

- Write latency: waits for both cache and DB
- Cache pollution: write data that's never read (wasted memory)
- Complex implementation in app or requires cache layer support

// Go:

```
func (r *Repo) CreateProduct(ctx context.Context, p *Product) error {
    // Write to DB first
    if err := r.db.InsertProduct(ctx, p); err != nil { return err }

    // Write to cache
    b, _ := json.Marshal(p)
    key := fmt.Sprintf("product:%d", p.ID)
    r.cache.Set(ctx, key, b, 1*time.Hour)
    return nil
}
```

When to use write-through:

- Read-after-write consistency critical (show user their own updates)
- Read-heavy with frequent re-reads of written data

When NOT to use:

- Write-heavy without read (cache fills with unread data)
- High write latency tolerance not available

Q40. What is the write-behind (write-back) pattern?

Difficulty: Hard

Write-behind: write to cache only; async flush to DB

- App writes to cache → returns immediately
- Background process: read from cache → write to DB

Write latency: near 0 (just in-memory write)

Risk: data loss if cache fails before flush

Implementation:

```
// Write: ultra fast (memory only)
rdb.Set(ctx, key, value, time.Hour)
rdb.LPush(ctx, "write-queue", key) // mark for flush

// Background flusher:
for {
    keys, _ := rdb.BRPop(ctx, time.Second, "write-queue").Result()
    for _, key := range keys[1:] {
        value, _ := rdb.Get(ctx, key).Result()
        db.Upsert(key, value) // flush to DB
    }
}
```

Safer implementation with Redis Streams:

- XADD write-log * key user:42 op SET value {...}
- Background consumer processes stream, writes to DB

Use cases:

- Shopping cart (acceptable to lose on crash)
- Game state (saved periodically)
- Analytics counters (approximate is OK)
- Session data (recreatable on next login)

Not for: financial transactions, inventory, anything critical

Q41. What is the refresh-ahead pattern?
Difficulty:
Medium

Refresh-ahead: proactively refresh cache before TTL expires
 Predict what will be needed next → warm up in background
 User never sees cache miss latency

Implementation:

- Option A: background goroutine refreshes popular keys
- Option B: refresh on near-expiry (probabilistic early expiration)

Probabilistic Early Expiration:

When TTL is close to 0, increase probability of background refresh

```
func shouldRefresh(ttl, totalTTL time.Duration, beta float64) bool {
    // beta: controls aggressiveness (typically 1.0)
    // Higher beta = refresh earlier
    elapsed := float64(totalTTL - ttl)
    threshold := -beta * math.Log(rand.Float64()) * float64(totalTTL)
    return elapsed >= float64(totalTTL) - threshold
}

// Simple refresh-ahead with goroutines:
func getWithRefreshAhead(ctx context.Context, key string) (string, error) {
    val, ttl, err := getWithTTL(ctx, key)
    if err != nil { return "", err }

    // If less than 20% TTL remains, refresh in background
    if ttl < originalTTL/5 {
        go func() {
            fresh := fetchFromDB(key)
            rdb.Set(context.Background(), key, fresh, originalTTL+jitter())
        }()
    }
    return val, nil
}
```

Q42. What is stale-while-revalidate?
Difficulty:
Medium

Stale-while-revalidate: serve stale data immediately, refresh async
 Client gets fast response (even if slightly stale)
 Background goroutine fetches fresh data
 Next request gets fresh data

HTTP Cache-Control header:

- Cache-Control: max-age=300, stale-while-revalidate=60
- Fresh for 300s
- Between 300s-360s: serve stale, trigger background revalidation

→ After 360s: must revalidate (blocking)

Implementation in Go:

```
type cachedValue struct {
    Data      []byte
    ExpiresAt time.Time
    RefreshAt time.Time // start background refresh at this time
}

func get(ctx context.Context, key string) ([]byte, error) {
    v := localCache.Get(key)
    if v == nil { return fetchAndCache(key) } // miss

    if time.Now().After(v.RefreshAt) && time.Now().Before(v.ExpiresAt) {
        go refreshAsync(key) // stale but still valid, refresh in background
    }

    if time.Now().After(v.ExpiresAt) {
        return fetchAndCache(key) // truly expired, blocking fetch
    }

    return v.Data, nil
}
```

Use case: news feed, product listings – slight staleness acceptable

Q43. What is a write-invalidate vs write-update strategy?

Difficulty:
Medium

Write-invalidate (delete on write):

On DB write: DELETE cache key

Next read: miss → fetch from DB → populate cache

Pros: simpler, no race conditions

Cons: first read after write is slower (cold miss)

Write-update (update on write):

On DB write: SET new value in cache

Next read: hit → return fresh cached value

Pros: next read is fast (no cold miss)

Cons: race conditions possible!

Race condition example:

T1: read user (v1) → start UPDATE query

T2: read user (v1) → start UPDATE query → finish → SET cache(v2)

T1: finish → SET cache(v1) ← STALE! T2's update overwritten

Why write-invalidate is safer:

No matter what order writes happen:

After all writes complete: cache is empty or has latest value

Race condition: two deletes for same key → both succeed, no stale data

Best practice: ALWAYS use write-invalidate (delete) not write-update

Exception: truly atomic compare-and-swap with version check

Q44. What is cache-aside with negative caching?**Difficulty:**
Medium

Negative caching: cache "not found" results to prevent DB hammering

Problem: request for non-existent key:

user:99999 → cache miss → DB query → not found → cache miss again → DB...

Solution: cache the "not found" result too

```
// Go implementation:
const notFoundMarker = "__NOT_FOUND__"

func getUser(ctx context.Context, id int64) (*User, error) {
    key := fmt.Sprintf("user:%d", id)

    data, err := rdb.Get(ctx, key).Bytes()
    if err == nil {
        if string(data) == notFoundMarker { return nil, ErrNotFound }
        var u User
        return &u, json.Unmarshal(data, &u)
    }

    u, err := db.GetUser(ctx, id)
    if errors.Is(err, ErrNotFound) {
        // Cache "not found" with short TTL (30s – user might be created soon)
        rdb.Set(ctx, key, notFoundMarker, 30*time.Second)
        return nil, ErrNotFound
    }
    if err != nil { return nil, err }

    b, _ := json.Marshal(u)
    rdb.Set(ctx, key, b, 5*time.Minute)
    return u, nil
}

// When user IS created: delete the negative cache entry
func createUser(ctx context.Context, u *User) error {
    if err := db.CreateUser(ctx, u); err != nil { return err }
    rdb.Del(ctx, fmt.Sprintf("user:%d", u.ID)) // clear negative cache
    return nil
}
```

Q45. What is cache stampede prevention with mutex?**Difficulty:** Hard

Mutex pattern: only ONE goroutine fetches from DB on miss; others wait

Distributed mutex with Redis:

Lock key: lock:user:42

First requester: acquires lock → fetches from DB → stores in cache → releases lock

Others: wait for lock → lock released → get value from cache

// Go with Redis lock:

```
func getWithMutex(ctx context.Context, key string) (string, error) {
    // Check cache first
    if val, err := rdb.Get(ctx, key).Result(); err == nil { return val, nil }
```

```

// Try to acquire lock
lockKey := "lock:" + key
locked, _ := rdb.SetNX(ctx, lockKey, "1", 5*time.Second).Result()

if !locked {
    // Another goroutine is fetching; wait and retry
    time.Sleep(50 * time.Millisecond)
    return getWithMutex(ctx, key) // retry (will likely hit cache now)
}
defer rdb.Del(ctx, lockKey)

// Double-check cache (another goroutine may have populated it)
if val, err := rdb.Get(ctx, key).Result(); err == nil { return val, nil }

// Fetch from DB
val := fetchFromDB(key)
rdb.Set(ctx, key, val, 5*time.Minute)
return val, nil
}

// Better: use golang.org/x/sync/singleflight for in-process dedup
var g singleflight.Group
val, _, _ := g.Do(key, func() (interface{}, error) {
    return fetchFromDB(key)
})

```

Q46. What is cache hierarchies (L1/L2)?

Difficulty:
Medium

Multi-level cache: L1 (in-process) + L2 (distributed Redis)

L1 (in-process):

- Location: application memory
- Latency: ~0ms (no network)
- Capacity: small (1K-100K items, limited by heap)
- Scope: per-instance only (not shared)
- TTL: very short (seconds to minutes)

L2 (distributed):

- Location: Redis cluster
- Latency: 1-5ms (network)
- Capacity: large (GB to TB)
- Scope: shared across all instances
- TTL: minutes to hours

Algorithm:

- Read: check L1 → hit return; miss → check L2 → hit: store in L1, return
→ L2 miss: fetch DB → store in L1 + L2 → return
- Write: invalidate L1 + delete L2 key

// Go L1 implementation:

```
import lru "github.com/hashicorp/golang-lru/v2"
```

```

type MultiLevelCache struct {
    l1 *lru.Cache[string, []byte] // in-process
    l2 *redis.Client             // Redis
}

```

```
func (c *MultiLevelCache) Get(ctx context.Context, key string) ([]byte, error) {
    if v, ok := c.l1.Get(key); ok { return v, nil } // L1 hit

    v, err := c.l2.Get(ctx, key).Bytes() // L2 hit
    if err == nil {
        c.l1.Add(key, v) // populate L1
        return v, nil
    }
    return nil, err // total miss
}
```

Q47. What is a CDN cache and how does it differ from Redis?

Difficulty:
Medium

CDN (Content Delivery Network) cache:

- Geographically distributed edge servers
- Cache HTTP responses at network edge
- Serves users from nearest PoP (Point of Presence)

Latency: 5-20ms (user to CDN PoP, not origin)

Capacity: petabytes across all PoPs

Use: static assets (images, JS, CSS), public API responses

Redis cache:

- Single or clustered in-datacenter
- Application-level caching
- Complex data structures

Latency: 1-5ms (within datacenter)

Capacity: TB in cluster

Use: session state, rate limiting, DB query results

Cache-Control headers for CDN:

public, max-age=3600	→ cache for 1 hour at CDN + browser
public, s-maxage=3600	→ cache 1h at CDN, browser uses max-age
no-cache	→ always revalidate (but can serve stale)
no-store	→ never cache (sensitive data)
stale-while-revalidate=60	→ serve stale, refresh async

CDN invalidation:

Purge by URL: `curl -X DELETE https://cdn/purge?url=...`

Purge by tag: all assets tagged "product:42"

Versioned URLs: /static/bundle.abc123.js (no purge needed)

Q48. ### Q61-Q65: More Cache Patterns

Difficulty:
Medium

Q	Topic
Q61	Caching SQL query results: when and how
Q62	Function-level memoization in Go
Q63	Cache-aside with circuit breaker for Redis failures
Q64	Idempotency key caching for API deduplication
Q65	Batch cache operations for N+1 prevention

Distributed Caching & Redis Cluster

Q49. What is Redis Cluster and how does it shard data?

Difficulty: Hard

Redis Cluster: horizontal sharding across multiple Redis nodes

Total hash slots: 16,384

Each key $\rightarrow$ $\text{CRC16}(\text{key}) \% 16384 \rightarrow$ hash slot $\rightarrow$ node

Node assignment:

3 primary nodes: Node A (slots 0-5460), B (5461-10922), C (10923-16383)

Each primary has 1-2 replicas for HA

Minimum: 3 primaries + 3 replicas = 6 nodes

Hash tags: {user_id} $\rightarrow$ hash only the part in {}

{user:42}:cart and {user:42}:profile $\rightarrow$ same slot

Enables MULTI/EXEC across multiple keys (same slot only)

Scaling:

Add nodes: Redis Cluster rebalances slots automatically

Remove nodes: drain slots first, then remove

Limitations:

Cross-slot operations: must use hash tags or pipeline separately

Lua scripts: all keys must be in same slot

Pub/Sub: works cluster-wide

Client routing:

Client receives MOVED error if queries wrong node $\rightarrow$ updates routing table

Smart clients (redis-go Cluster, ioredis): route automatically

When to use:

Dataset > 100GB OR ops/sec > 1M (single node limit)

Single node: ~100GB RAM, ~1M ops/sec

Cluster: horizontal scale to any size

Q50. What is Redis Sentinel?

Difficulty: Medium

Redis Sentinel: HA for single-master Redis setup (no sharding)

Monitors primary + replicas

Automatic failover if primary dies

Service discovery: clients ask Sentinel "who is primary?"

Architecture:

3 Sentinel processes (quorum=2 for decisions)

1 Primary + N replicas

Failover process:

1. Sentinels detect primary unreachable (3 $\times$ ping timeout)
2. Quorum (2/3) agrees primary is down (SDOWN $\rightarrow$ ODOWN)
3. Leader Sentinel elected (Raft-like)
4. Leader promotes best replica to primary
5. Reconfigures other replicas to follow new primary

6. Notifies clients via Sentinel API

Client connection:

Connect to Sentinel → ask for primary address → connect to primary
On failover: Sentinel notifies → client reconnects to new primary

```
// Go:
rdb := redis.NewFailoverClient(&redis.FailoverOptions{
    MasterName:      "mymaster",
    SentinelAddrs: []string{"sentinel1:26379", "sentinel2:26379", "sentinel3:26379"},
})
```

Sentinel vs Cluster:

Sentinel: HA only, single master, no sharding
Cluster: HA + sharding, multiple masters
Use Sentinel when: data fits single node, want simpler ops
Use Cluster when: need horizontal scale

Q51. What is consistent hashing for cache distribution?

Difficulty: Hard

Consistent hashing: distribute keys across N cache nodes
Minimizes remapping when nodes are added/removed

Problem with naive hashing:

key → hash(key) % N → server
Add 1 server (N=4→5): almost ALL keys remap → massive cache misses

Consistent hashing:

Hash space: 0 → 2³² (ring)
Servers placed at positions on ring (via hash(server_name))
Key maps to nearest server clockwise on ring

Add server: only keys between (new, predecessor) remapped
Remove server: only keys on that server remap to successor
Impact: K/N keys remapped (K=total keys, N=nodes)

Virtual nodes:

Each real node = 150 virtual nodes on ring
→ Better distribution, more even load
→ New server takes equal share from all existing servers

```
// Go implementation:
import "github.com/stathat/consistent"
c := consistent.New()
c.Add("server1")
c.Add("server2")
c.Add("server3")
server, _ := c.Get("user:42") // consistently routes to same server
```

Used by: Memcached clients, Redis clients (some), CDN, service discovery

Q52. What is a cache hot key and how do you solve it?

Difficulty: Hard

Hot key: one cache key receiving disproportionate traffic
Example: celebrity profile read by 1M followers/sec
Redis: single-threaded → one key becomes bottleneck

Solutions:**1. Local in-process cache (L1):**

Each app server caches the hot key locally

TTL = 5-10 seconds (short for freshness)

100 app servers × 10K req/sec → 1M req/sec served without Redis

```
var localCache = sync.Map{}

func getHotKey(ctx context.Context, key string) (string, error) {
    if v, ok := localCache.Load(key); ok {
        return v.(string), nil // 0ms hit
    }
    val, err := rdb.Get(ctx, key).Result() // Redis fallback
    if err == nil {
        localCache.Store(key, val)
        time.AfterFunc(5*time.Second, func() { localCache.Delete(key) })
    }
    return val, err
}
```

2. Key replication / scatter reads:

key:copy1, key:copy2, key:copy3 on different Redis nodes

Read: key:"copy" + strconv.Itoa(rand.Intn(3)+1)

Write: update all copies

3. Redis Cluster + hash tags:

Distribute natural sharding key better

4. Read replicas:

Redis replica nodes for read-only traffic

Q53. What is Redis cross-slot operations and hash tags?**Difficulty: Hard**

Redis Cluster limitation: multi-key operations only for keys in same slot

MGET key1 key2 → keys must be in same slot

MULTI/EXEC transactions → all keys same slot

Lua scripts → all keys same slot

SUNIONSTORE, ZINTERSTORE → same slot

Hash tags: {tag} forces key into hash slot of tag

CRC16 computed on tag only, not full key

{user:42}:cart → slot of "user:42"

{user:42}:profile → slot of "user:42" ← SAME slot!

{user:42}:sessions → slot of "user:42" ← SAME slot!

Now: MGET {user:42}:cart {user:42}:profile → works!

Now: MULTI/EXEC across these keys → works!

Downside: all keys with same tag → same slot → hot spot

Don't tag everything with same value

Tag by entity that needs atomic operations

// Go:

```
keys := []string{
    fmt.Sprintf("{user:%d}:cart", userID),
```

```
    fmt.Sprintf("{user:%d}:profile", userID),
}
vals, _ := rdb.MGet(ctx, keys...).Result()
```

Redis Cluster routing:

MOVED 3999 127.0.0.1:6381 → wrong slot, go to this node

ASK 3999 127.0.0.1:6381 → slot being migrated, try this node temporarily

Q54. ### Q71-Q80: More Distributed Cache Topics

Difficulty:
Medium

Q	Topic
Q71	Memcached vs Redis: when to use each
Q72	Redis read replicas for read scaling
Q73	Cache partition strategies: by user, by feature, by data type
Q74	Geo-distributed caching: CDN + regional Redis
Q75	Cache warm-up strategies for Blue-Green deployments
Q76	Redis connection pool tuning for high-concurrency Go apps
Q77	Cache key design: namespacing, versioning, structure
Q78	Cache serialization: JSON vs MessagePack vs Protobuf
Q79	Redis cluster topology: primary-only vs primary-replica per shard
Q80	Handling Redis cluster failover in Go applications

Advanced Topics & Go Integration

Q55. How do you integrate Redis with Go using go-redis?

Difficulty: Easy

```
import "github.com/redis/go-redis/v9"

// Single node
rdb := redis.NewClient(&redis.Options{
    Addr:      "localhost:6379",
    Password:  "",
    DB:        0,
    PoolSize:  10,
    MinIdleConns: 5,
    MaxRetries: 3,
    DialTimeout: 5 * time.Second,
    ReadTimeout: 3 * time.Second,
    WriteTimeout: 3 * time.Second,
})
```



```
// Sentinel (HA)
rdb = redis.NewFailoverClient(&redis.FailoverOptions{
    MasterName:      "mymaster",
    SentinelAddrs: []string{"s1:26379", "s2:26379", "s3:26379"},
    Password:        "password",
    PoolSize:        10,
})

// Cluster
rdb = redis.NewClusterClient(&redis.ClusterOptions{
    Addrs: []string{"node1:6379", "node2:6379", "node3:6379"},
    PoolSize: 5,
})

// Ping
if err := rdb.Ping(ctx).Err(); err != nil {
    log.Fatalf("redis connection failed: %v", err)
}
```

Q56. How do you implement a leaderboard in Go with Redis?

Difficulty:
Medium

```
type LeaderboardService struct { rdb *redis.Client }

const leaderboardKey = "leaderboard:global"

func (s *LeaderboardService) AddScore(ctx context.Context, playerID string, score float64) error {
    return s.rdb.ZAdd(ctx, leaderboardKey, redis.Z{
        Score: score,
        Member: playerID,
    }).Err()
}

func (s *LeaderboardService) IncrementScore(ctx context.Context, playerID string, delta float64) (float64, error) {
    return s.rdb.ZIncrBy(ctx, leaderboardKey, delta, playerID).Result()
}

func (s *LeaderboardService) GetTopN(ctx context.Context, n int) ([]redis.Z, error) {
    return s.rdb.ZRevRangeWithScores(ctx, leaderboardKey, 0, int64(n-1)).Result()
}

func (s *LeaderboardService) GetRank(ctx context.Context, playerID string) (int64, error) {
    rank, err := s.rdb.ZRevRank(ctx, leaderboardKey, playerID).Result()
    return rank + 1, err // 1-indexed
}

func (s *LeaderboardService) GetNeighbors(ctx context.Context, playerID string, spread int) ([]redis.Z, error) {
    rank, _ := s.rdb.ZRevRank(ctx, leaderboardKey, playerID).Result()
    start := rank - int64(spread)
    if start < 0 { start = 0 }
    return s.rdb.ZRevRangeWithScores(ctx, leaderboardKey, start, rank+int64(spread)).Result()
}
```

Q57. How do you implement Redis-based caching middleware in Go?

Difficulty:
Medium

```
// HTTP middleware for response caching
```

```

func CacheMiddleware(rdb *redis.Client, ttl time.Duration) func(http.Handler) http.Handler {
    return func(next http.Handler) http.Handler {
        return http.HandlerFunc(func(w http.ResponseWriter, r *http.Request) {
            // Only cache GET requests
            if r.Method != http.MethodGet {
                next.ServeHTTP(w, r)
                return
            }

            key := "http:" + r.URL.RequestURI()

            // Check cache
            if data, err := rdb.Get(r.Context(), key).Bytes(); err == nil {
                w.Header().Set("X-Cache", "HIT")
                w.Header().Set("Content-Type", "application/json")
                w.Write(data)
                return
            }

            // Capture response
            rec := &responseRecorder{ResponseWriter: w, body: &bytes.Buffer{}}
            next.ServeHTTP(rec, r)

            // Cache successful responses
            if rec.status == http.StatusOK {
                rdb.Set(r.Context(), key, rec.body.Bytes(),
                    ttl+time.Duration(rand.Intn(30))*time.Second)
                w.Header().Set("X-Cache", "MISS")
            }
        })
    }
}

type responseRecorder struct {
    http.ResponseWriter
    status int
    body   *bytes.Buffer
}

func (r *responseRecorder) WriteHeader(status int) {
    r.status = status; r.ResponseWriter.WriteHeader(status)
}

func (r *responseRecorder) Write(b []byte) (int, error) {
    r.body.Write(b); return r.ResponseWriter.Write(b)
}

```

Q58. How do you implement a Redis-based job queue in Go?

Difficulty: Hard

```

// Simple reliable job queue using Redis List + processing set
type Queue struct {
    rdb      *redis.Client
    queueKey string
    processingKey string
}

func (q *Queue) Enqueue(ctx context.Context, payload []byte) error {
    return q.rdb.LPush(ctx, q.queueKey, payload).Err()
}

```

```

func (q *Queue) Dequeue(ctx context.Context, timeout time.Duration) ([]byte, error) {
    // BRPOPLPUSH: atomically move from queue to processing set
    result, err := q.rdb.BRPopLPush(ctx, q.queueKey, q.processingKey, timeout).Result()
    if err == redis.Nil { return nil, nil } // timeout
    return []byte(result), err
}

func (q *Queue) Acknowledge(ctx context.Context, payload []byte) error {
    // Remove from processing set after successful handling
    return q.rdb.LRem(ctx, q.processingKey, 1, payload).Err()
}

func (q *Queue) RecoverStuck(ctx context.Context, timeout time.Duration) error {
    // Move stuck jobs back to main queue (ran by cron or background worker)
    items, _ := q.rdb.LRange(ctx, q.processingKey, 0, -1).Result()
    for _, item := range items {
        // In production: check timestamp embedded in job to determine age
        q.rdb.RPush(ctx, q.queueKey, item)
        q.rdb.LRem(ctx, q.processingKey, 1, item)
    }
    return nil
}

```

Q59. How do you implement Redis pub/sub for cache invalidation in Go?

Difficulty:
Medium

```

// Multi-server cache invalidation via Redis Pub/Sub

type CacheInvalidator struct {
    rdb      *redis.Client
    localCache sync.Map
}

// Subscriber: listen for invalidation events
func (c *CacheInvalidator) Subscribe(ctx context.Context) {
    pubsub := c.rdb.Subscribe(ctx, "cache:invalidate")
    defer pubsub.Close()

    for {
        select {
        case <-ctx.Done(): return
        case msg := <-pubsub.Channel():
            c.localCache.Delete(msg.Payload)
            log.Printf("invalidated local cache key: %s", msg.Payload)
        }
    }
}

// Publisher: broadcast invalidation when data changes
func (c *CacheInvalidator) InvalidateAll(ctx context.Context, key string) error {
    // 1. Delete from Redis
    c.rdb.Del(ctx, key)
    // 2. Invalidate local cache on this server
    c.localCache.Delete(key)
    // 3. Notify all other servers to invalidate their local caches
    return c.rdb.Publish(ctx, "cache:invalidate", key).Err()
}

```

```
// Usage:
inv := &CacheInvalidator{rdb: rdb}
go inv.Subscribe(ctx) // start listener in background

// On data update:
db.UpdateUser(user)
inv.InvalidateAll(ctx, fmt.Sprintf("user:%d", user.ID))
```

Q60. How do you monitor Redis health in Go?

Difficulty: Easy

```
type RedisHealthChecker struct { rdb *redis.Client }

func (h *RedisHealthChecker) Check(ctx context.Context) error {
    // Ping
    if err := h.rdb.Ping(ctx).Err(); err != nil {
        return fmt.Errorf("redis ping failed: %w", err)
    }
    return nil
}

// Readiness probe endpoint
http.HandleFunc("/ready", func(w http.ResponseWriter, r *http.Request) {
    ctx, cancel := context.WithTimeout(r.Context(), 2*time.Second)
    defer cancel()
    if err := checker.Check(ctx); err != nil {
        w.WriteHeader(http.StatusServiceUnavailable)
        json.NewEncoder(w).Encode(map[string]string{"redis": err.Error()})
        return
    }
    json.NewEncoder(w).Encode(map[string]string{"status": "ok"})
})

// Connection pool stats
stats := rdb.PoolStats()
log.Printf("redis pool: hits=%d misses=%d timeouts=%d totalConns=%d idleConns=%d",
    stats.Hits, stats.Misses, stats.Timeouts,
    stats.TotalConns, stats.IdleConns)

// Export to Prometheus:
redisPoolHits.Set(float64(stats.Hits))
redisPoolMisses.Set(float64(stats.Misses))
```

Q61. How do you implement rate limiting with Redis in Go?

Difficulty:
Medium

```
type RateLimiter struct {
    rdb      *redis.Client
    script   *redis.Script
}

func NewRateLimiter(rdb *redis.Client) *RateLimiter {
    // Atomic Lua: increment + set expiry if new key
    script := redis.NewScript(`
        local count = redis.call('INCR', KEYS[1])
        if count == 1 then
            redis.call('EXPIRE', KEYS[1], ARGV[1])
        end
    `)
```

```

        return count
    `)
    return &RateLimiter{rdb: rdb, script: script}
}

type RateLimitResult struct {
    Allowed    bool
    Count      int64
    Remaining  int64
    ResetAt    time.Time
}

func (r *RateLimiter) Allow(ctx context.Context, key string, limit int64, window time.Duration) (*RateLimitResult, error) {
    windowSecs := int(window.Seconds())
    now := time.Now()

    count, err := r.script.Run(ctx, r.rdb,
        []string{fmt.Sprintf("rl:%s:%d", key, now.Unix()/int64(windowSecs))},
        windowSecs,
    ).Int64()
    if err != nil { return nil, err }

    resetAt := time.Unix((now.Unix()/int64(windowSecs)+1)*int64(windowSecs), 0)
    return &RateLimitResult{
        Allowed:    count <= limit,
        Count:      count,
        Remaining:  max(0, limit-count),
        ResetAt:    resetAt,
    }, nil
}

// HTTP middleware:
func RateLimitMiddleware(limiter *RateLimiter) func(http.Handler) http.Handler {
    return func(next http.Handler) http.Handler {
        return http.HandlerFunc(func(w http.ResponseWriter, r *http.Request) {
            result, _ := limiter.Allow(r.Context(), r.RemoteAddr, 100, time.Minute)
            w.Header().Set("X-RateLimit-Remaining", strconv.FormatInt(result.Remaining, 10))
            if !result.Allowed {
                w.Header().Set("Retry-After", strconv.Itoa(int(time.Until(result.ResetAt).Seconds())))
                w.WriteHeader(http.StatusTooManyRequests)
                return
            }
            next.ServeHTTP(w, r)
        })
    }
}

```

Q62. How do you implement cache warming on startup in Go?

Difficulty:
Medium

```

type CacheWarmer struct {
    rdb *redis.Client
    db  *pgxpool.Pool
}

func (w *CacheWarmer) WarmProducts(ctx context.Context) error {
    log.Println("warming product cache...")
}

```

```

rows, err := w.db.Query(ctx,
    "SELECT id, data FROM products ORDER BY views DESC LIMIT 10000")
if err != nil { return err }
defer rows.Close()

// Use pipeline for batch SET
pipe := w.rdb.Pipeline()
count := 0
for rows.Next() {
    var id int64
    var data []byte
    rows.Scan(&id, &data)

    key := fmt.Sprintf("product:%d", id)
    ttl := 30*time.Minute + time.Duration(rand.Intn(300))*time.Second
    pipe.Set(ctx, key, data, ttl)
    count++

    // Execute in batches of 500
    if count%500 == 0 {
        pipe.Exec(ctx)
        pipe = w.rdb.Pipeline()
    }
}
if count%500 != 0 { pipe.Exec(ctx) }

log.Printf("warmed %d products", count)
return nil
}

// Call on startup before accepting traffic
func main() {
    warmer := &CacheWarmer{rdb: rdb, db: db}
    if err := warmer.WarmProducts(ctx); err != nil {
        log.Printf("cache warm failed (non-fatal): %v", err)
    }
    // Start server...
}

```

Q63. What is a circuit breaker for Redis in Go?

Difficulty: Hard

```

import "github.com/sony/gobreaker"

type ResilientCache struct {
    rdb      *redis.Client
    breaker *gobreaker.CircuitBreaker
}

func NewResilientCache(rdb *redis.Client) *ResilientCache {
    cb := gobreaker.NewCircuitBreaker(gobreaker.Settings{
        Name:      "redis",
        MaxRequests: 5,
        Interval:   30 * time.Second,
        Timeout:    60 * time.Second,
        ReadyToTrip: func(counts gobreaker.Counts) bool {
            return counts.ConsecutiveFailures > 10
        },
    },

```

```

        OnStateChange: func(name string, from, to gobreaker.State) {
            log.Printf("Redis circuit breaker: %s → %s", from, to)
        },
    })
    return &ResilientCache{rdb: rdb, breaker: cb}
}

func (c *ResilientCache) Get(ctx context.Context, key string) ([]byte, error) {
    result, err := c.breaker.Execute(func() (interface{}, error) {
        return c.rdb.Get(ctx, key).Bytes()
    })
    if err != nil {
        if errors.Is(err, gobreaker.ErrOpenState) {
            return nil, nil // cache unavailable, fall through to DB
        }
        return nil, err
    }
    return result.([]byte), nil
}

func (c *ResilientCache) Set(ctx context.Context, key string, value []byte, ttl time.Duration) {
    c.breaker.Execute(func() (interface{}, error) {
        return nil, c.rdb.Set(ctx, key, value, ttl).Err()
    })
    // Ignore cache set errors (best-effort)
}

```

Q64. How do you implement idempotency keys with Redis in Go?

Difficulty:
Medium

// Idempotency: same request with same key → same result (no duplicate processing)

```
type IdempotencyStore struct{ rdb *redis.Client }
```

```
const idempotencyTTL = 24 * time.Hour
```

```
type StoredResult struct {
    StatusCode int    `json:"status_code"`
    Body       []byte `json:"body"`
    CreatedAt  int64  `json:"created_at"`
}
```

```
func (s *IdempotencyStore) GetOrProcess(
    ctx context.Context,
    key string,
    handler func() (int, []byte, error),
) (int, []byte, error) {
    redisKey := "idempotency:" + key

    // Check if already processed
    data, err := s.rdb.Get(ctx, redisKey).Bytes()
    if err == nil {
        var result StoredResult
        json.Unmarshal(data, &result)
        return result.StatusCode, result.Body, nil // return cached result
    }

    // Not processed: use singleflight to prevent concurrent processing

```

```

var g singleflight.Group
v, err, _ := g.Do(key, func() (interface{}, error) {
    status, body, err := handler()
    if err != nil { return nil, err }

    result := StoredResult{StatusCode: status, Body: body, CreatedAt: time.Now().Unix()}
    d, _ := json.Marshal(result)
    s.rdb.Set(ctx, redisKey, d, idempotencyTTL)
    return &result, nil
})
if err != nil { return 0, nil, err }
r := v.(*StoredResult)
return r.StatusCode, r.Body, nil
}

// Usage in payment handler:
status, body, err := idempotency.GetOrProcess(ctx, idempotencyKey, func() (int, []byte, error) {
    result, err := paymentService.Charge(order)
    // ... process payment ...
})

```

Q65. ### Q91-Q100: Additional Advanced Cache Topics**Difficulty:**
Medium

Q	Topic
Q91	Redis WATCH for optimistic locking in Go
Q92	Implementing a distributed counter with Redis in Go
Q93	Cache key versioning for zero-downtime deploys
Q94	Redis keyspace notifications for event-driven cache expiry
Q95	Testing Go code that uses Redis (testcontainers, miniredis)
Q96	Redis memory profiling and optimization techniques
Q97	Implementing a feature flag service with Redis in Go
Q98	Cache stampede prevention: Go singleflight vs Redis lock
Q99	Redis Streams as a reliable job queue vs List-based queue
Q100	Cache design review: choosing TTL, eviction, and consistency

Master these 100 questions and you'll handle any caching interview at SDE2 level. Key areas: cache patterns (aside, through, behind), Redis data structures (Sorted Set for leaderboards, Streams for queues), stampede prevention (singleflight + jitter), distributed cache (Cluster, Sentinel, consistent hashing), and Go integration (go-redis, pipeline, Lua scripts). ■

Additional Questions (Q66–Q150)

Q66. What is Redis Cluster and how does it work?

Difficulty: Hard

Redis Cluster: horizontal scaling via sharding across 16384 hash slots.

16384 hash slots distributed across N master nodes

Each node handles a range of slots (e.g., node1: 0-5460)

Key $\rightarrow$ CRC16(key) % 16384 $\rightarrow$ slot $\rightarrow$ node

3-node cluster:

Node A: slots 0-5460 (+ replica A')

Node B: slots 5461-10922 (+ replica B')

Node C: slots 10923-16383 (+ replica C')

Replication: each master has 1+ replicas

Failover: replica promotes if master unreachable (majority vote)

Create cluster

```
redis-cli --cluster create \
  127.0.0.1:7000 127.0.0.1:7001 127.0.0.1:7002 \
  127.0.0.1:7003 127.0.0.1:7004 127.0.0.1:7005 \
  --cluster-replicas 1
```

Check cluster status

```
redis-cli -p 7000 cluster info
```

```
redis-cli -p 7000 cluster nodes
```

// Go client (redis/go-redis v9)

```
rdb := redis.NewClusterClient(&redis.ClusterOptions{
  Addrs: []string{"7000", "7001", "7002"},
  Password: "secret",
})
```

Interview tip: "Hash tags {user} force related keys to same slot. Use for multi-key operations (MGET, pipelining). Without same slot: CROSSSLOT error."

Q67. What are hash tags in Redis Cluster?

Difficulty: Medium

Hash tag: {key} – only the part inside {} is hashed for slot assignment

Used to colocate related keys on same node

Without hash tag:

user:1:profile $\rightarrow$ slot 123 (Node A)

user:1:orders $\rightarrow$ slot 456 (Node B)

$\rightarrow$ can't MGET both!

With hash tag:

{user:1}:profile $\rightarrow$ CRC16("user:1") % 16384 $\rightarrow$ same slot

{user:1}:orders $\rightarrow$ CRC16("user:1") % 16384 $\rightarrow$ same slot

$\rightarrow$ MGET works! Both on same node

// Ensure keys are on same slot for pipeline

```
pipe := rdb.Pipeline()
pipe.Get(ctx, "{user:1}:profile")
pipe.Get(ctx, "{user:1}:orders")
```

```
pipe.Get(ctx, "{user:1}:cart")
cmds, err := pipe.Exec(ctx)
```

Q68. What is Redis Sentinel?

Difficulty:
Medium

Redis Sentinel: high availability for single Redis master (not cluster)
Monitors master + replicas, auto-failover if master dies

Sentinel topology:

Sentinel1, Sentinel2, Sentinel3 (need odd number for quorum)
→ Monitor: Redis-Master + Redis-Replica1 + Redis-Replica2

Failover process:

1. Sentinels detect master unreachable (down-after-milliseconds)
2. Quorum reached (majority of sentinels agree)
3. One sentinel elected as leader
4. Leader selects best replica (most recent data)
5. Replica promoted to master
6. Other replicas reconfigured to follow new master
7. Clients notified (get new master address)

sentinel.conf:

```
sentinel monitor mymaster 127.0.0.1 6379 2 # quorum=2
sentinel down-after-milliseconds mymaster 5000
sentinel failover-timeout mymaster 60000
```

```
rdb := redis.NewFailoverClient(&redis.FailoverOptions{
    MasterName:      "mymaster",
    SentinelAddrs: []string{":26379", ":26380", ":26381"},
    Password:        "secret",
})
```

Q69. What is Redis persistence: RDB vs AOF?

Difficulty: Hard

RDB (Redis Database Snapshot):

Periodic point-in-time snapshot of all data
Fork process writes snapshot to disk
Fast restart (loads compact binary)
Data loss: up to last snapshot interval

```
CONFIG: save 900 1 # save if 1 key changed in 900s
        save 300 10 # save if 10 keys changed in 300s
        save 60 10000 # save if 10000 keys changed in 60s
FILE: dump.rdb
```

AOF (Append-Only File):

Logs every write command
More durable (configurable fsync policy)
Larger file, slower restart

```
appendonly yes
appendfsync everysec # fsync every second (good balance)
# appendfsync always # fsync every write (safest, slowest)
# appendfsync no      # OS decides (fastest, least safe)
FILE: appendonly.aof
```

Hybrid (Redis 4.0+):

```
RDB snapshot + AOF tail (since last snapshot)
aof-use-rdb-preamble yes
Best of both: fast load + durability
```

Q70. What is Redis pub/sub?

Difficulty:
Medium

```
// Redis Pub/Sub: simple publish-subscribe (no persistence, no ACK)
// Messages lost if no subscribers
// Not for reliable messaging (use Streams for that)

// Publisher
err := rdb.Publish(ctx, "notifications:user:42", payload).Err()

// Subscriber
pubsub := rdb.Subscribe(ctx, "notifications:user:42")
defer pubsub.Close()

ch := pubsub.Channel()
for msg := range ch {
    fmt.Printf("channel=%s payload=%s\n", msg.Channel, msg.Payload)
}

// Pattern subscribe (wildcard)
pubsub := rdb.PSubscribe(ctx, "notifications:*")
for msg := range pubsub.Channel() {
    fmt.Println(msg.Channel, msg.Payload)
}

// Pub/Sub vs Redis Streams:
// Pub/Sub: fire-and-forget, no persistence, fan-out
// Streams: persistent, consumer groups, replay, ACK
// Use Streams for: task queues, event sourcing, reliable messaging
```

Q71. What are Redis Streams?

Difficulty: Hard

```
// Redis Streams: persistent, ordered log (like Kafka-lite)
// Features: consumer groups, ACK, pending entries, trim

// Produce
id, err := rdb.XAdd(ctx, &redis.XAddArgs{
    Stream: "orders",
    Values: map[string]interface{}{
        "order_id": "42",
        "amount":   "99.99",
        "status":   "pending",
    },
    MaxLen: 10000, // keep last 10K messages
}).Result()
// id = "1234567890-0" (timestamp-sequence)

// Consume (consumer group for load balancing)
rdb.XGroupCreate(ctx, "orders", "order-processors", "0")

msgs, err := rdb.XReadGroup(ctx, &redis.XReadGroupArgs{
    Group:   "order-processors",
    Consumer: "worker-1",
})
```

```

Streams: []string{"orders", ">"}, // ">" = new messages only
Count:   10,
Block:   5 * time.Second,
}).Result()

for _, msg := range msgs[0].Messages {
    process(msg.Values)
    rdb.XAck(ctx, "orders", "order-processors", msg.ID) // acknowledge
}

// Check pending (unacked) messages
pending, _ := rdb.XPendingExt(ctx, &redis.XPendingExtArgs{
    Stream: "orders", Group: "order-processors",
    Start: "-", Stop: "+", Count: 100,
}).Result()

```

Q72. What are Redis Bloom Filters?

Difficulty: Hard

Bloom Filter: probabilistic data structure
 Test if element is in a set: "possibly yes" OR "definitely no"
 False positives possible, false negatives impossible
 Space-efficient: fixed size regardless of element count

Redis Bloom Filter (RedisBloom module):

```

BF.ADD key item          # add item
BF.EXISTS key item       # check membership
BF.MADD key i1 i2 i3    # batch add
BF.MEXISTS key i1 i2    # batch check

```

Use cases:

- Cache warming: check if item might be in cache before DB query
- Duplicate request detection (has this webhook been processed?)
- Spam email filter
- Username availability check

go-redis with Bloom:

```

rdb.BFAdd(ctx, "processed-events", eventID)
exists, err := rdb.BFExists(ctx, "processed-events", eventID)

```

False positive rate: $p = (1 - e^{(-kn/m)})^k$

k = number of hash functions
 n = number of elements
 m = bit array size
 Lower p = larger bit array needed

Q73. What is Redis HyperLogLog?

Difficulty:
Medium

```

// HyperLogLog: approximate count of unique elements
// Uses ~12KB regardless of cardinality
// Error rate: ~0.81%

// Count unique page views
rdb.PFAdd(ctx, "visitors:2024-01-15", "user1", "user2", "user3")
rdb.PFAdd(ctx, "visitors:2024-01-15", "user1", "user4") // user1 already counted

count, _ := rdb.PFCount(ctx, "visitors:2024-01-15").Result()

```

```
// count ≈ 4 (unique visitors)

// Merge multiple HLL (union)
rdb.PFMerge(ctx, "visitors:week",
    "visitors:2024-01-13",
    "visitors:2024-01-14",
    "visitors:2024-01-15",
)
weekCount, _ := rdb.PFCount(ctx, "visitors:week").Result()

// vs exact count with SET:
// SET: exact but O(n) memory (all IDs stored)
// HLL: ~0.81% error, O(1) memory (12KB always)

// Use for: analytics, cardinality estimation
// Not for: when you need exact count or to retrieve items
```

Q74. What is Redis sorted set use cases?

Difficulty:
Medium

```
// Sorted Set (ZSET): members with float scores, sorted by score
// O(log N) for most operations

// Leaderboard
rdb.ZAdd(ctx, "game:leaderboard",
    redis.Z{Score: 9500, Member: "alice"},
    redis.Z{Score: 8200, Member: "bob"},
    redis.Z{Score: 11000, Member: "charlie"},
)

// Get top 10 (highest score first)
leaders, _ := rdb.ZRevRangeWithScores(ctx, "game:leaderboard", 0, 9).Result()

// Get player rank (0-indexed)
rank, _ := rdb.ZRevRank(ctx, "game:leaderboard", "alice").Result()
// alice is rank 1 (charlie is 0)

// Rate limiting with sliding window
now := time.Now().UnixMilli()
window := int64(60000) // 60 seconds
rdb.ZAdd(ctx, "requests:user:42", redis.Z{Score: float64(now), Member: now})
rdb.ZRemRangeByScore(ctx, "requests:user:42", "0", strconv.FormatInt(now-window, 10))
count, _ := rdb.ZCard(ctx, "requests:user:42").Result()
if count > 100 { return ErrRateLimited }

// Scheduled tasks (score = execution time)
rdb.ZAdd(ctx, "scheduled_jobs", redis.Z{
    Score: float64(time.Now().Add(5*time.Minute).Unix()),
    Member: "job:send-email:42",
})
// Poll: get jobs due now
jobs, _ := rdb.ZRangeByScore(ctx, "scheduled_jobs", &redis.ZRangeBy{
    Min: "0",
    Max: strconv.FormatInt(time.Now().Unix(), 10),
    Count: 10,
}).Result()
```

Q75. What is Redis transaction (MULTI/EXEC)?**Difficulty:**
Medium

```
// MULTI/EXEC: execute commands atomically (no interleaving)
// NOT isolation like SQL transactions
// All commands queued, executed atomically

// Basic transaction
pipe := rdb.TxPipeline()
pipe.Decr(ctx, "inventory:product:42")
pipe.Incr(ctx, "orders:count")
pipe.Set(ctx, "order:99:status", "confirmed", 0)
_, err := pipe.Exec(ctx)

// WATCH + MULTI/EXEC (optimistic locking)
// Retry if watched key changed (CAS pattern)
const maxRetries = 5
for i := 0; i < maxRetries; i++ {
    err := rdb.Watch(ctx, func(tx *redis.Tx) error {
        // Get current value
        n, err := tx.Get(ctx, "inventory").Int()
        if err != nil { return err }
        if n <= 0 { return errors.New("out of stock") }

        // Atomic decrement only if inventory unchanged
        _, err = tx.TxPipelined(ctx, func(pipe redis.Pipeliner) error {
            pipe.Decr(ctx, "inventory")
            return nil
        })
        return err
    }, "inventory")

    if err == nil { break }
    if err != redis.TxFailedErr { return err } // retry only on conflict
    time.Sleep(time.Duration(i) * 10 * time.Millisecond)
}
```

Q76. What is Redis connection pooling in Go?**Difficulty:**
Medium

```
rdb := redis.NewClient(&redis.Options{
    Addr:      "localhost:6379",
    Password:  "secret",
    DB:        0,

    // Pool settings
    PoolSize: 10, // max connections (default: 10 per CPU)
    MinIdleConns: 3, // keep this many idle connections
    MaxIdleConns: 5, // max idle
    PoolTimeout: 4 * time.Second, // wait for connection from pool
    ConnMaxIdleTime: 5 * time.Minute, // close idle connections after
    ConnMaxLifetime: 30 * time.Minute,

    // Timeouts
    DialTimeout: 5 * time.Second,
    ReadTimeout: 3 * time.Second,
    WriteTimeout: 3 * time.Second,
```

```
})

// Monitor pool stats
stats := rdb.PoolStats()
fmt.Printf("hits=%d misses=%d timeouts=%d total=%d idle=%d\n",
    stats.Hits, stats.Misses, stats.Timeouts,
    stats.TotalConns, stats.IdleConns)

// Health check
if err := rdb.Ping(ctx).Err(); err != nil {
    log.Fatal("redis ping failed:", err)
}
```

Q77. What is Redis keyspace notifications?

Difficulty:
Medium

```
// Keyspace notifications: get notified when keys expire, are set, deleted

// Enable in redis.conf or via CONFIG SET:
// notify-keyspace-events "Ex" # E=keyevent events, x=expired events
// notify-keyspace-events "KEA" # K=keyspace, E=keyevent, A=all events

rdb.ConfigSet(ctx, "notify-keyspace-events", "Ex")

// Subscribe to expiry events
pubsub := rdb.Subscribe(ctx, "__keyevent@0__:expired")
ch := pubsub.Channel()

for msg := range ch {
    expiredKey := msg.Payload
    fmt.Printf("key expired: %s\n", expiredKey)
    // Handle expiry: refresh cache, notify user, etc.
}

// Use cases:
// - Cache expiry hooks (warm up cache before item expires)
// - Session expiry notification (logout user)
// - Distributed lock expiry detection
// - TTL-based job scheduling

// Note: notifications are best-effort (may be lost)
// Don't use for critical workflows
```

Q78. What are Redis pipelines and batching?

Difficulty:
Medium

```
// Pipeline: send multiple commands in single round-trip
// No interleaving guarantee (vs MULTI/EXEC)

pipe := rdb.Pipeline()
for i := 0; i < 100; i++ {
    pipe.Set(ctx, fmt.Sprintf("key:%d", i), i, 0)
}
_, err := pipe.Exec(ctx)

// Pipelined (functional style)
cmds, err := rdb.Pipelined(ctx, func(pipe redis.Pipeliner) error {
```

```

    for i := 0; i < 100; i++ {
        pipe.Get(ctx, fmt.Sprintf("key:%d", i))
    }
    return nil
})
for _, cmd := range cmds {
    fmt.Println(cmd.(*redis.StringCmd).Val())
}

// Performance comparison:
// 100 individual GETs: 100 round trips × 0.1ms = 10ms
// 100 pipelined GETs: 1 round trip × 0.1ms + processing = ~1ms

// Cluster pipeline: routes to correct node automatically
pipe := rdb.(*redis.ClusterClient).Pipeline()

```

Q79. What is Redis memory management?

Difficulty: Hard

```

# Memory info
redis-cli INFO memory
# used_memory: current memory
# used_memory_peak: peak memory
# mem_fragmentation_ratio: >1.5 = fragmented (defrag needed)

# Memory limit and eviction
CONFIG SET maxmemory 2gb
CONFIG SET maxmemory-policy allkeys-lru

# Eviction policies:
# noeviction:      error when memory full (default)
# allkeys-lru:     evict any key by LRU
# volatile-lru:    evict keys with TTL by LRU
# allkeys-random:  evict random key
# volatile-random: evict random key with TTL
# volatile-ttl:    evict key with shortest TTL
# allkeys-lfu:     evict by LFU (least frequently used, Redis 4+)
# volatile-lfu:    evict keys with TTL by LFU

# For cache: allkeys-lru or allkeys-lfu
# For session store: volatile-lru (keep keys without TTL)

# Check memory per key type
redis-cli --bigkeys # find largest keys
redis-cli MEMORY USAGE mykey # bytes used by specific key

// In Go: monitor memory
info, _ := rdb.Info(ctx, "memory").Result()
// Parse used_memory field

```

Q80. What is Redis Lua scripting?

Difficulty: Hard

```

// Lua scripts: atomic execution of multiple commands
// No interruption between commands (like transaction but with logic)

// Example: atomic increment with cap
luaScript := redis.NewScript(`
local current = tonumber(redis.call('GET', KEYS[1])) or 0

```



```

local max = tonumber(ARGV[1])
if current < max then
    return redis.call('INCR', KEYS[1])
end
return -1 -- at cap
`)

result, err := luaScript.Run(ctx, rdb, []string{"counter:api-calls"}, 100).Int()
if result == -1 {
    return ErrRateLimitExceeded
}

// Rate limiter using Lua (atomic check-and-increment)
rateLimitScript := `
local count = redis.call('INCR', KEYS[1])
if count == 1 then
    redis.call('EXPIRE', KEYS[1], ARGV[1])
end
return count
`

count, _ := rdb.Eval(ctx, rateLimitScript, []string{"rate:user:42"}, 60).Int()
if count > 100 { return ErrRateLimited }

```

Q81. ### Q81–Q100: Advanced Cache Patterns

Difficulty:
Medium

Q82. What is a multi-level cache architecture?

Difficulty:
Medium

L1: In-process cache (e.g., ristretto, bigcache)
Access: ~1ns, size limited by service memory
Scope: single pod only

L2: Redis (shared cache)
Access: ~0.5ms, larger capacity
Scope: all pods share same data

L3: Database
Access: ~5ms (indexed), ~50ms (complex queries)
Source of truth

Lookup order:

1. Check L1 (in-process)
2. Cache miss → check L2 (Redis)
3. Cache miss → query L3 (DB)
4. Populate L2, then L1

Invalidation:

On write → DELETE from L2 (all pods' L1 will expire naturally)
Or: short TTL in L1 (accept some staleness)

```

func (c *MultiLevelCache) Get(ctx context.Context, key string) (string, error) {
    // L1 check
    if v, ok := c.l1.Get(key); ok { return v.(string), nil }

    // L2 check
    v, err := c.redis.Get(ctx, key).Result()

```

```

    if err == nil {
        c.l1.Set(key, v, 30*time.Second) // populate L1
        return v, nil
    }

    // DB
    v, err = c.db.Get(ctx, key)
    if err != nil { return "", err }
    c.redis.Set(ctx, key, v, 5*time.Minute) // populate L2
    c.l1.Set(key, v, 30*time.Second)        // populate L1
    return v, nil
}

```

Q83. What is consistent hashing for cache distribution?

Difficulty:
Medium

Consistent hashing: distribute keys across nodes with minimal redistribution

Problem with simple $\text{hash}(\text{key}) \% N$:

Add/remove node $\rightarrow$ almost ALL keys remap $\rightarrow$ cache miss storm

Consistent hashing:

Virtual ring of 2^{32} slots

Nodes placed at multiple points (virtual nodes) on ring

Key $\rightarrow$ hash $\rightarrow$ clockwise to next node

Add node: only keys between new node and its predecessor remap

Remove node: only keys from removed node remap to next node

$\sim 1/N$ keys remap (vs $\sim N-1/N$ for simple hashing)

Virtual nodes: each physical node has 150+ virtual nodes

$\rightarrow$ more even distribution

$\rightarrow$ fewer hot spots

Libraries:

github.com/stathat/consistent (Go)

hashring pattern

Q84. What is Redis Geo for location-based queries?

Difficulty:
Medium

// Redis GEO: store/query geographic coordinates

// Internally uses sorted set with geohash score

// Add locations

```

rdb.GeoAdd(ctx, "restaurants",
    &redis.GeoLocation{Name: "pizza-palace", Longitude: -122.4194, Latitude: 37.7749},
    &redis.GeoLocation{Name: "sushi-bar", Longitude: -122.4089, Latitude: 37.7851},
)

```

// Find nearby (within 5km)

```

results, _ := rdb.GeoRadius(ctx, "restaurants", -122.4194, 37.7749,
    &redis.GeoRadiusQuery{
        Radius:    5,
        Unit:      "km",
        WithDist:  true,
        WithCoord: true,
    })

```

```

        Count:    10,
        Sort:     "ASC", // nearest first
    },
).Result()

for _, r := range results {
    fmt.Printf("%s: %.2f km away\n", r.Name, r.Dist)
}

// Distance between two points
dist, _ := rdb.GeoDist(ctx, "restaurants", "pizza-palace", "sushi-bar", "km").Result()

```

Q85. What is the Read-Aside vs Write-Through pattern?

Difficulty:
Medium

Cache-Aside (Lazy Loading):

Read: check cache → miss → read DB → populate cache

Write: write to DB → invalidate/update cache

Pros: only cache what's needed, resilient to cache failure

Cons: cold start, thundering herd on first access

When: read-heavy, unpredictable access patterns

Write-Through:

Write: write to cache AND DB atomically (or DB first, then cache)

Read: always hits cache (freshest data)

Pros: cache always up-to-date, no stale reads

Cons: write latency (wait for both), cache bloat (write everything)

When: read-after-write consistency required, data always reused

Write-Behind (Write-Back):

Write: write to cache immediately, async flush to DB

Pros: ultra-fast writes

Cons: data loss if cache crashes before flush, complexity

When: high-write scenarios, eventual consistency acceptable

Q86. What is the cache warming strategy?

Difficulty:
Medium

// Cache warming: pre-populate cache before going live

// Prevents cold start / thundering herd

// Strategy 1: Eager warming on startup

```

func warmCache(ctx context.Context, db *sql.DB, rdb *redis.Client) error {
    // Load top 1000 products (most accessed)
    products, err := db.QueryContext(ctx,
        "SELECT id,data FROM products ORDER BY view_count DESC LIMIT 1000")
    if err != nil { return err }

```

```

    pipe := rdb.Pipeline()
    for products.Next() {
        var id int64; var data []byte

```

```

        products.Scan(&id, &data)
        pipe.Set(ctx, fmt.Sprintf("product:%d", id), data, time.Hour)
    }
    _, err = pipe.Exec(ctx)
    return err
}

// Strategy 2: Background warming (don't block startup)
go func() {
    time.Sleep(5 * time.Second) // let service start first
    warmCache(ctx, db, rdb)
}()

// Strategy 3: Predictive warming (next likely-needed keys)
// After serving product:42, pre-warm related products

```

Q87. What is Redis TTL management best practices?

Difficulty:
Medium

```

// TTL strategies:
// 1. Absolute TTL: key expires at fixed time
rdb.Set(ctx, "session:abc", data, 30*time.Minute)

// 2. Sliding TTL: reset on each access
func getWithSlidingTTL(ctx context.Context, key string) (string, error) {
    pipe := rdb.Pipeline()
    get := pipe.Get(ctx, key)
    pipe.Expire(ctx, key, 30*time.Minute) // reset TTL on access
    pipe.Exec(ctx)
    return get.Val(), get.Err()
}

// 3. TTL jitter: prevent thundering herd (all keys expiring together)
func setWithJitter(ctx context.Context, key string, val interface{}, base time.Duration) error {
    jitter := time.Duration(rand.Intn(int(base / 10))) // ±10%
    return rdb.Set(ctx, key, val, base+jitter).Err()
}

// Check remaining TTL
ttl, _ := rdb.TTL(ctx, "session:abc").Result()
if ttl < 5*time.Minute {
    // Proactively refresh before expiry
    refreshSession(ctx, key)
}

```

Q88. What is the session store pattern with Redis?

Difficulty:
Medium

```

// Redis as centralized session store

type Session struct {
    UserID    int64
    ExpiresAt time.Time
    Data      map[string]interface{}
}

func createSession(ctx context.Context, rdb *redis.Client, userID int64) (string, error) {

```

```

    sessionID := generateSecureToken()
    session := Session{
        UserID:    userID,
        ExpiresAt: time.Now().Add(24 * time.Hour),
        Data:       make(map[string]interface{}),
    }

    data, _ := json.Marshal(session)
    err := rdb.Set(ctx, "session:"+sessionID, data, 24*time.Hour).Err()
    return sessionID, err
}

func getSession(ctx context.Context, rdb *redis.Client, sessionID string) (*Session, error) {
    data, err := rdb.Get(ctx, "session:"+sessionID).Bytes()
    if errors.Is(err, redis.Nil) { return nil, ErrSessionNotFound }
    if err != nil { return nil, err }

    var session Session
    json.Unmarshal(data, &session)
    return &session, nil
}

// Sticky sessions vs centralized:
// Sticky: user always hits same server (simpler but hard to scale)
// Centralized Redis: any server handles any user (horizontal scale)

```

Q89. What is distributed rate limiting with Redis?

Difficulty:
Medium

```

// Sliding window rate limiter (most accurate)
func isRateLimited(ctx context.Context, rdb *redis.Client, key string, limit int, window time.Duration) (bool, error) {
    now := time.Now().UnixMilli()
    windowMs := window.Milliseconds()

    pipe := rdb.TxPipeline()
    pipe.ZRemRangeByScore(ctx, key, "0", strconv.FormatInt(now-windowMs, 10)) // remove old
    pipe.ZAdd(ctx, key, redis.Z{Score: float64(now), Member: now})              // add current
    countCmd := pipe.ZCard(ctx, key)                                           // count in window
    pipe.Expire(ctx, key, window)                                              // set TTL
    _, err := pipe.Exec(ctx)
    if err != nil { return false, err }

    count, _ := countCmd.Result()
    return count > int64(limit), nil
}

// Token bucket (bursting allowed)
// Fixed window (simpler, less accurate)
// Leaky bucket (smooth output rate)

// Choose based on requirements:
// API rate limiting: sliding window (most fair)
// Burst tolerance: token bucket
// Simple protection: fixed window

```

Q90. What is Redis ACL (Access Control Lists)?**Difficulty:**
Medium

```
# Redis 6.0+ ACL: fine-grained access control

# Create user
ACL SETUSER alice on >password ~cache:* +get +set +del
# alice: enabled, password "password", can only access keys matching "cache:*"
# can only use GET, SET, DEL commands

ACL SETUSER readonly on >pass ~* +@read # read-only on all keys
ACL SETUSER admin on >pass ~* +@all # full access

# In Go
rdb := redis.NewClient(&redis.Options{
    Addr:      "localhost:6379",
    Username:  "alice",
    Password:  "password",
})

# List users
ACL LIST
ACL WHOAMI
ACL LOG # log auth failures
```

Q91. What is Redis keyspace design best practices?**Difficulty:**
Medium

Key naming conventions:

```
<object>:<id>           user:42
<object>:<id>:<field>     user:42:profile
<namespace>:<object>:<id> myapp:prod:user:42
```

Avoid:

- Very long keys (memory overhead)
- JSON as key (use hash tags for clustering)
- Spaces in keys (use _ or :)

Key size:

- Keys: keep < 1KB (ideally < 128 bytes)
- Values: depends on use case
- Large values: consider compression (msgpack vs JSON)

Naming for different data types:

```
String:  user:42           (serialized JSON)
Hash:    user:42:fields    (individual fields)
Set:     user:42:friends   (friend IDs)
List:    user:42:feed      (recent activity)
ZSet:    leaderboard       (score-based)
Stream:  events:orders     (event log)
```

Key expiry:

- Always set TTL for cache keys
- No TTL for persistent data (sessions with explicit logout)
- Monitor keys without TTL: `redis-cli --scan --pattern "cache:*`

Q92. What are Redis data structure memory optimization?**Difficulty:**
Medium

Encoding optimization:

Small hashes use ziplist (< hash-max-ziplist-entries=128):

hash-max-ziplist-entries 128

hash-max-ziplist-value 64

→ Saves 10x memory vs hash table

Small lists use ziplist:

list-max-ziplist-size -2 (-2 = 8KB max)

Small sets of integers use intset:

set-max-intset-entries 512

Small sorted sets use ziplist:

zset-max-ziplist-entries 128

zset-max-ziplist-value 64

Object encoding check:

DEBUG OBJECT mykey → encoding: ziplist | hashtable | intset | etc.

OBJECT ENCODING mykey

Compression:

Compress values with gzip/snappy before storing

Trade CPU for memory

Worth it for values > 1KB

redis-cli --memkeys: analyze memory usage by key pattern

Q93. What is cache consistency in distributed systems?**Difficulty:**
Medium

Consistency problems:

Write races: two writers update cache and DB in different orders

Stale reads: cache has old value after DB update

Solutions:

1. Cache-aside with TTL (accept eventual consistency):

Read from cache if fresh, else DB

Stale window = TTL

Simple, works for most cases

2. Write-invalidate (most reliable):

Write to DB → DELETE cache key

Next read fetches fresh from DB

Problem: delete-before-write race (two processes)

3. Write-invalidate with version:

DB row has version column

Cache stores: {"version": 5, "data": {...}}

On cache hit: compare version with DB (occasional check)

4. CDC-based invalidation:

Debezium reads DB WAL → invalidates Redis keys

Eventually consistent, no application-level sync

5. Read-your-writes consistency:

After write: route reads to primary DB (bypass cache)

After TTL or delay: switch back to cache reads

Use sticky routing or version token

Q94. What is Redis keyspace memory analysis?

Difficulty:
Medium

```
# Memory analysis tools

# 1. redis-cli --bigkeys (find largest keys)
redis-cli --bigkeys
# Samples: 100 keys, finds largest per type

# 2. redis-cli --memkeys (Go 1.x, newer redis-cli)
redis-cli -u redis://localhost:6379 --memkeys

# 3. MEMORY USAGE (individual key)
redis-cli MEMORY USAGE user:42 # bytes used by key + value

# 4. SCAN + MEMORY USAGE (script all keys in pattern)
redis-cli --scan --pattern "session:*" | xargs -L 1 redis-cli MEMORY USAGE

# 5. RDB analysis (rdbtools)
rdb --command memory dump.rdb | sort -t, -k4 -rn | head -20

# Common memory hogs:
# Large string values (JSON blobs)
# Unbounded lists/sets
# Keys without TTL accumulating
# Hash fields explosion (1M fields in one hash)
```

Q95. What is Redis eviction and OOM handling?

Difficulty:
Medium

```
// Handle Redis OOM in application

func setCache(ctx context.Context, key string, val interface{}, ttl time.Duration) error {
    data, _ := json.Marshal(val)
    err := rdb.Set(ctx, key, data, ttl).Err()
    if err != nil {
        if strings.Contains(err.Error(), "OOM") {
            // Redis is out of memory
            // Log but don't fail the request (cache is optional)
            log.Printf("cache OOM, skipping cache write for key %s", key)
            return nil // graceful degradation
        }
        return err
    }
    return nil
}

// Monitor eviction rate
// redis-cli INFO stats | grep evicted_keys
// Alert if evicted_keys rate > 0 (means maxmemory hit)
```



```
// Set appropriate maxmemory:
// total_memory × 0.75 (leave headroom for fragmentation)
// Example: 8GB Redis → maxmemory 6gb

// Monitor with Prometheus:
// redis_evicted_keys_total (counter, alert on rate > 0)
// redis_memory_used_bytes / redis_memory_max_bytes (usage ratio, alert >80%)
```

Q96. What is Redis replication lag?

Difficulty:
Medium

```
# Check replication status
redis-cli INFO replication
# role: master
# connected_slaves: 2
# slave0: ip=...,port=...,state=online,offset=123456,lag=0
# slave1: ip=...,port=...,state=online,offset=123450,lag=2

# lag=0: fully caught up
# lag=2: 2 seconds behind (normal under load)

# master_repl_offset: current master offset
# slave0.offset: replica's confirmed offset
# Lag bytes = master_repl_offset - slave0.offset

# Monitoring:
# redis_connected_slaves (alert if drops below expected)
# redis_replication_backlog_size
# slave lag (from INFO replication parsing)

# Replication buffer
config set repl-backlog-size 1mb # buffer for disconnected replicas

# Partial resync: replica reconnects, catches up from backlog
# Full resync: if backlog too small, full RDB copy needed (expensive)
```

Q97. What is in-memory cache libraries in Go?

Difficulty:
Medium

```
// Options:

// 1. ristretto (Cloudflare): most production-grade
import "github.com/dgraph-io/ristretto"

cache, _ := ristretto.NewCache(&ristretto.Config{
    NumCounters: 1e7,      // track 10M items for admission
    MaxCost:     1 << 30,  // 1GB max
    BufferItems:  64,      // performance tuning
})
cache.Set("key", value, cost) // cost = estimated memory
val, found := cache.Get("key")

// 2. bigcache (Allegro): for millions of small items
import "github.com/allegro/bigcache/v3"

cache, _ := bigcache.New(context.Background(), bigcache.DefaultConfig(10*time.Minute))
cache.Set("key", []byte(value))
```

```
entry, _ := cache.Get("key")

// 3. groupcache: distributed cache without Redis
// Perfect for read-heavy, shared computation

// 4. freecache: zero GC overhead
import "github.com/coocood/freecache"
cache := freecache.NewCache(100 * 1024 * 1024) // 100MB
cache.Set([]byte("key"), []byte("val"), 300)
```

Q98. What is the cache stampede (thundering herd) and prevention?

Difficulty:
Medium

```
// Cache stampede: many requests hit DB simultaneously when cache expires
```

```
// Solution 1: singleflight (coalesce identical requests)
import "golang.org/x/sync/singleflight"
```

```
var sfGroup singleflight.Group
```

```
func getUser(ctx context.Context, id int64) (*User, error) {
    key := fmt.Sprintf("user:%d", id)

    result, err, _ := sfGroup.Do(key, func() (interface{}, error) {
        // Only ONE goroutine runs this for the same key
        // Others wait and share the result
        user, err := db.GetUser(ctx, id)
        if err != nil { return nil, err }
        redis.Set(ctx, key, user, time.Hour)
        return user, nil
    })
    if err != nil { return nil, err }
    return result.(*User), nil
}
```

```
// Solution 2: mutex per key (Redis SETNX lock)
```

```
func getWithLock(ctx context.Context, key string) (string, error) {
    val, err := rdb.Get(ctx, key).Result()
    if err == nil { return val, nil } // cache hit

    // Try to acquire lock
    locked, _ := rdb.SetNX(ctx, key+":lock", "1", 10*time.Second).Result()
    if !locked {
        time.Sleep(50 * time.Millisecond)
        return rdb.Get(ctx, key).Result() // wait and retry
    }
    defer rdb.Del(ctx, key+":lock")

    // We have the lock: compute value
    val = fetchFromDB()
    rdb.Set(ctx, key, val, time.Hour)
    return val, nil
}
```

```
// Solution 3: probabilistic early expiration (PER)
// Refresh cache slightly before expiry with probability
```

Q99. What is Redis TLS configuration?**Difficulty:**
Medium

```
import "crypto/tls"

rdb := redis.NewClient(&redis.Options{
    Addr: "redis.prod.example.com:6380",
    TLSConfig: &tls.Config{
        MinVersion: tls.VersionTLS12,
        // For custom CA:
        RootCAs: loadCACert("redis-ca.pem"),
        // For mTLS (client cert):
        Certificates: []tls.Certificate{loadClientCert()},
    },
})

// Redis Cluster with TLS
rdb := redis.NewClusterClient(&redis.ClusterOptions{
    Addrs: []string{":7000", ":7001", ":7002"},
    TLSConfig: &tls.Config{MinVersion: tls.VersionTLS12},
})

// redis.conf:
// tls-port 6380
// tls-cert-file /etc/ssl/redis.crt
// tls-key-file /etc/ssl/redis.key
// tls-ca-cert-file /etc/ssl/ca.crt
// tls-auth-clients no # optional client certs
```

Q100. What are Redis performance benchmarks?**Difficulty:**
Medium

```
# Redis benchmark tool
redis-benchmark -q -n 100000
# -n: number of requests
# -c: concurrent connections (default 50)
# -q: quiet (only summary)

redis-benchmark -q -n 100000 -c 100 -t get,set,incr,lpush

# Typical results (local, single thread):
# GET: ~110,000 ops/sec
# SET: ~110,000 ops/sec
# INCR: ~110,000 ops/sec
# LPUSH: ~100,000 ops/sec

# With pipeline (batch)
redis-benchmark -n 100000 -P 16 # pipeline 16 cmds
# GET: ~1,000,000 ops/sec (10x improvement)

# Latency percentiles
redis-benchmark --latency-history -t get -i 1
redis-cli --latency-history # live latency monitoring

# Redis is single-threaded for commands:
# I/O multiplexing handles many connections
# Commands serialize on main thread
```

Cluster: scale horizontally per shard

Q101. What is the cache invalidation strategy for microservices?

Difficulty:
Medium

Cache invalidation is hard ("one of two hard problems in CS")

Strategies for microservices:

1. TTL-based (simplest):
Set short TTL, accept staleness window
No invalidation needed – just wait for expiry
Good for: product catalog, configs, non-critical data
2. Event-driven invalidation:
Service A updates DB → publishes "user:updated" event
Service B (cache owner) subscribes → deletes/refreshes cache
Requires: message broker (Kafka/RMQ)
Good for: cross-service cache dependencies
3. Write-through invalidation:
Update DB → immediately invalidate/update cache
In same transaction or 2-phase with retry
Good for: single-service, same-DB ownership
4. CDC-based (most reliable):
Debezium reads DB WAL → publishes changes → cache invalidator subscribes
Decoupled from application code
Near-real-time
Good for: large-scale, multi-consumer invalidation
5. Version/ETag based:
Cache stores version alongside data
Check version on read, invalidate if stale
Good for: API caching, HTTP caching

Rule: cache owner is responsible for invalidation
Never let service B invalidate service A's cache

Extended Questions (Q102–Q150)

Q102. What is ristretto cache in Go?

Difficulty:
Medium

```
import "github.com/dgraph-io/ristretto"

// Ristretto: high-performance in-process cache for Go
// Features: admission policy (TinyLFU), per-item cost, metrics

cache, err := ristretto.NewCache(&ristretto.Config{
    NumCounters: 1e7,      // 10M counters for admission tracking
    MaxCost:     1 << 30,  // 1GB max total cost
    BufferItems:  64,       // async operations per goroutine
    Metrics:     true,      // enable metrics
})
```

```

}))
if err != nil { panic(err) }
defer cache.Close()

// Set with cost (estimated memory bytes)
cost := int64(len(key) + len(value))
cache.Set(key, value, cost)
cache.SetWithTTL(key, value, cost, time.Hour)

// Get
val, found := cache.Get(key)
if found {
    use(val.(string))
}

// Wait for async operations to complete (in tests)
cache.Wait()

// Metrics
m := cache.Metrics
fmt.Printf("hits=%d misses=%d ratio=%.2f\n",
    m.Hits(), m.Misses(), m.Ratio())

// Eviction callback
config.OnEvict = func(item *ristretto.Item) {
    log.Printf("evicted: key=%v", item.Key)
}

```

Q103. What is bigcache in Go?

Difficulty:
Medium

```

import "github.com/allegro/bigcache/v3"

// BigCache: fast, GC-friendly in-process cache
// Key feature: stores values as []byte → no GC pressure from interface{}
// Suitable for: millions of small items

config := bigcache.DefaultConfig(10 * time.Minute)
config.MaxEntriesInWindow = 1000 * 10 * 60
config.MaxEntrySize = 500           // bytes
config.HardMaxCacheSize = 256       // MB
config.OnRemove = func(key string, entry []byte) {
    log.Printf("removed: %s", key)
}

cache, err := bigcache.New(context.Background(), config)
if err != nil { panic(err) }
defer cache.Close()

// Store bytes
cache.Set("user:42", jsonBytes)

// Get bytes
entry, err := cache.Get("user:42")
if errors.Is(err, bigcache.ErrEntryNotFound) {
    // cache miss
}

```

```
// Delete
cache.Delete("user:42")

// Stats
stats := cache.Stats()
fmt.Printf("hits=%d misses=%d\n", stats.Hits, stats.Misses)

// Limitations:
// No per-item TTL (global TTL only)
// No cost-based eviction
// LRU-ish eviction (hash segments)
```

Q104. What is cache-aside pattern in Go?

Difficulty:
Medium

```
type UserCache struct {
    redis *redis.Client
    db     *pgxpool.Pool
    ttl    time.Duration
}

func (c *UserCache) GetUser(ctx context.Context, id int64) (*User, error) {
    key := fmt.Sprintf("user:%d", id)

    // 1. Try cache
    data, err := c.redis.Get(ctx, key).Bytes()
    if err == nil {
        var user User
        if err := json.Unmarshal(data, &user); err == nil {
            return &user, nil
        }
    }

    // 2. Cache miss → DB
    var user User
    err = c.db.QueryRow(ctx, "SELECT id,name,email FROM users WHERE id=$1", id).
        Scan(&user.ID, &user.Name, &user.Email)
    if errors.Is(err, pgx.ErrNoRows) {
        // Cache negative result to prevent DB hammering
        c.redis.Set(ctx, key+":notfound", "1", time.Minute)
        return nil, ErrNotFound
    }
    if err != nil { return nil, err }

    // 3. Populate cache
    data, _ = json.Marshal(user)
    c.redis.Set(ctx, key, data, c.ttl)

    return &user, nil
}

func (c *UserCache) UpdateUser(ctx context.Context, user *User) error {
    _, err := c.db.Exec(ctx, "UPDATE users SET name=$1,email=$2 WHERE id=$3",
        user.Name, user.Email, user.ID)
    if err != nil { return err }

    // Invalidate cache
```

```

    c.redis.Del(ctx, fmt.Sprintf("user:%d", user.ID))
    return nil
}

```

Q105. What is negative caching?

Difficulty:
Medium

```

// Negative caching: cache "not found" results to prevent DB hammering
// Problem: DoS via non-existent keys → every request hits DB

func (c *Cache) GetProduct(ctx context.Context, id int64) (*Product, error) {
    key := fmt.Sprintf("product:%d", id)
    negKey := fmt.Sprintf("product:%d:notfound", id)

    // Check negative cache first
    if c.redis.Exists(ctx, negKey).Val() > 0 {
        return nil, ErrNotFound // fast return, no DB hit
    }

    // Check positive cache
    data, err := c.redis.Get(ctx, key).Bytes()
    if err == nil {
        var p Product
        json.Unmarshal(data, &p)
        return &p, nil
    }

    // DB lookup
    product, err := c.db.GetProduct(ctx, id)
    if errors.Is(err, pgx.ErrNoRows) {
        // Cache the "not found" with shorter TTL
        c.redis.Set(ctx, negKey, "1", 30*time.Second) // short TTL
        return nil, ErrNotFound
    }
    if err != nil { return nil, err }

    // Cache positive result
    data, _ = json.Marshal(product)
    c.redis.Set(ctx, key, data, time.Hour)
    return product, nil
}

```

Q106. What is cache penetration vs avalanche vs breakdown?

Difficulty: Hard

Cache Penetration:

Definition: query for non-existent key → always hits DB

Cause: malicious requests with fake IDs, deleted data

Solution:

1. Negative caching: cache "not found" with short TTL
2. Bloom filter: fast check if key exists before hitting cache/DB
3. Input validation: reject obviously invalid IDs

Cache Avalanche:

Definition: many keys expire at same time → traffic floods DB

Cause: mass deployment (all keys set with same TTL), power outage

Solution:

1. TTL jitter: add random offset (TTL + rand(0, 10%*TTL))

2. Gradual expiry: different TTLs for different key groups
3. Circuit breaker: throttle DB when overloaded
4. Multi-level cache: L1 in-process + L2 Redis

Cache Breakdown (hot key expiry):

Definition: one hot key expires → thundering herd on that key

Cause: viral post, trending product, live event

Solution:

1. Mutex/singleflight: only one goroutine rebuilds
2. Never expire hot keys (or very long TTL)
3. Pre-warm before expiry (probabilistic early refresh)
4. Background async refresh (serve stale while refreshing)

Q107. What is Redis cluster hash slot migration?

Difficulty: Hard

Resharding: move hash slots between nodes (for rebalancing)
redis-cli --cluster reshard 127.0.0.1:7000

Interactive: how many slots? from which node? to which node?

Or automated:

```
redis-cli --cluster reshard 127.0.0.1:7000 \
  --cluster-from source-node-id \
  --cluster-to target-node-id \
  --cluster-slots 1000 \
  --cluster-yes
```

Add new node (scale out)

```
redis-cli --cluster add-node 127.0.0.1:7006 127.0.0.1:7000
```

Then reshard some slots to new node

Remove node (scale in)

First: move all its slots to other nodes

```
redis-cli --cluster reshard ... # move slots away
```

Then remove:

```
redis-cli --cluster del-node 127.0.0.1:7000 node-id
```

During migration:

MOVED error: key moved to different node (client redirects)

ASK error: key being migrated (temporary redirect)

go-redis handles both automatically (smart client)

Check cluster balance

```
redis-cli --cluster check 127.0.0.1:7000
```

```
redis-cli --cluster info 127.0.0.1:7000
```

Q108. What is Redis RESP protocol?

Difficulty: Hard

RESP (Redis Serialization Protocol): simple, binary-safe protocol

Types:

- + Simple string: +OK\r\n
- Error: -ERR unknown command\r\n
- : Integer: :1000\r\n
- \$ Bulk string: \$6\r\nfoobar\r\n
- \$ Bulk string: \$-1\r\n (null bulk string)
- * Array: *3\r\n:1\r\n:2\r\n:3\r\n ([1,2,3])

Example: SET key value

Client sends: *3\r\n\$3\r\nSET\r\n\$3\r\nkey\r\n\$5\r\nvalue\r\n

Server responds: +OK\r\n

Example: GET key

Client: *2\r\n\$3\r\nGET\r\n\$3\r\nkey\r\n

Server: \$5\r\nvalue\r\n (or \$-1\r\n for nil)

RESP3 (Redis 6.0+):

New types: Map, Set, Double, Boolean, BigNumber, VerbatimString

Push messages (for pub/sub, keyspace notifications)

go-redis v9: supports RESP3 automatically

Pipelining efficiency:

Send N commands without waiting for responses

Server processes in order, sends all responses

1 RTT instead of N RTTs

Q109. What is Redis Object Encoding internals?

Difficulty: Hard

Redis uses different internal encodings based on value characteristics

Check encoding:

OBJECT ENCODING mykey

String encodings:

int: integer $\leq 2^{63}-1$ (e.g., INCR counter)

embstr: string ≤ 44 bytes (single allocation, immutable)

raw: string > 44 bytes (two allocations, mutable)

Hash encodings:

ziplist $\rightarrow$ listpack: $\leq$ hash-max-listpack-entries (128) and value ≤ 64 bytes

hashtable: when exceeds thresholds (10-15x larger!)

List encodings:

listpack: $\leq$ list-max-listpack-size

quicklist: when exceeds (doubly-linked list of listpacks)

Set encodings:

listpack: $\leq$ set-max-listpack-entries (128) and values ≤ 64 bytes

intset: all integers $\leq$ set-max-intset-entries (512)

hashtable: otherwise

Sorted Set encodings:

listpack: $\leq$ zset-max-listpack-entries (128)

skiplist + hashtable: when exceeds

Performance impact:

ziplist/listpack: compact, cache-friendly, $O(n)$ operations

hashtable: $O(1)$ lookup, more memory

Threshold tuning: larger thresholds $\rightarrow$ more memory savings but slower operations

Q110. What is Redis pub/sub vs Streams comparison?

Difficulty:
Medium

Redis Pub/Sub:

Pattern: fire-and-forget publish/subscribe

Persistence: NONE (if no subscribers $\rightarrow$ message lost)

Consumer groups: NO (all subscribers get all messages)
 Replay: IMPOSSIBLE (no history)
 Delivery: at-most-once
 Use: real-time notifications, live chat, metrics broadcasting

Redis Streams:

Pattern: persistent, ordered log
 Persistence: YES (stored until XDEL or XTRIM)
 Consumer groups: YES (like Kafka groups, each group independent)
 Replay: YES (read from any position)
 Delivery: at-least-once (with ACK)
 Pending entries: track unacked messages per consumer
 Use: reliable event log, task queues, audit trail

When to use Pub/Sub:

Live dashboard (stale data ok if missed)
 Cache invalidation notifications (TTL handles stale anyway)
 Chat messages (missed = user reconnects and loads from DB)

When to use Streams:

Task queue (can't lose tasks)
 Order events (need replay for new consumers)
 Audit logging
 Replacing Kafka for small-scale event streaming

Code comparison:

Pub/Sub: SUBSCRIBE channel, PUBLISH channel message
 Streams: XADD stream * field1 val1, XREADGROUP GROUP g consumer STREAMS stream >

Q111. What is cache serialization formats?

Difficulty:
Medium

```
// JSON: human readable, universal, slower
data, _ := json.Marshal(user)           // ~200 ns
json.Unmarshal(data, &user)             // ~300 ns
// Size: ~100 bytes for simple struct

// MessagePack: binary JSON, 20-30% smaller, 2-3x faster
import "github.com/vmihailov/msgpack"
data, _ := msgpack.Marshal(user)         // ~80 ns
msgpack.Unmarshal(data, &user)          // ~120 ns
// Size: ~70 bytes

// Protocol Buffers: fastest, smallest, requires schema
import "google.golang.org/protobuf/proto"
data, _ := proto.Marshal(userProto)      // ~50 ns
proto.Unmarshal(data, userProto)         // ~80 ns
// Size: ~50 bytes

// GOB: Go-specific, no schema needed
var buf bytes.Buffer
gob.NewEncoder(&buf).Encode(user)
gob.NewDecoder(&buf).Decode(&user)
// Comparable to JSON but Go-only

// Recommendation for Redis cache:
// High throughput, simple types: MessagePack
```

```
// Cross-language compatibility: JSON or Protobuf
// Go-only microservice: MessagePack or GOB
// Complex schemas, many services: Protobuf
```

Q112. What is Redis geo commands for location services?

Difficulty:
Medium

```
// GEO: location-based features (restaurant finder, delivery tracking)

// Add locations (lon, lat order!)
rdb.GeoAdd(ctx, "drivers", &redis.GeoLocation{
    Name:      "driver:42",
    Longitude: 77.5946, // Delhi longitude
    Latitude:  12.9716, // Bangalore latitude
})

// Update driver location
rdb.GeoAdd(ctx, "drivers", &redis.GeoLocation{
    Name:      "driver:42",
    Longitude: 77.5946 + 0.01,
    Latitude:  12.9716 + 0.01,
})

// Find drivers within 5km of customer
results, _ := rdb.GeoSearch(ctx, "drivers", &redis.GeoSearchQuery{
    Member:      "customer:99", // search around this member
    // OR: Longitude + Latitude for a point
    RadiusMiles: 5,
    Sort:        "ASC",
    Count:       10,
}).Result()

// GeoSearchStore: store results in another sorted set
rdb.GeoSearchStore(ctx, "drivers", "nearby_drivers:99", &redis.GeoSearchStoreQuery{
    GeoSearchQuery: redis.GeoSearchQuery{Longitude: 77.59, Latitude: 12.97, RadiusKM: 5, Sort: "ASC"},
    StoreDist: true,
})

// Distance
dist, _ := rdb.GeoDist(ctx, "drivers", "driver:42", "driver:43", "km").Result()
fmt.Printf("distance: %.2f km\n", dist)

// Get coordinates
pos, _ := rdb.GeoPos(ctx, "drivers", "driver:42").Result()
fmt.Printf("lat=%.4f lon=%.4f\n", pos[0].Latitude, pos[0].Longitude)
```

Q113. What is Redis WAIT command for replication sync?

Difficulty: Hard

```
// WAIT: block until N replicas acknowledged writes
// Use: read-after-write consistency in Redis with replicas

// After critical write: ensure N replicas have it
numReplicas, err := rdb.Wait(ctx, 1, 1000).Result()
// Wait for 1 replica to ack, timeout 1000ms
// Returns: number of replicas that acknowledged (may be 0 if timeout)

if numReplicas < 1 {
```

```

    // Replica didn't ack in time – decide: proceed anyway or error?
    log.Printf("replica sync timeout, proceeding with %d replicas", numReplicas)
}

// Use case: payment service
func processPayment(ctx context.Context, rdb *redis.Client, payment Payment) error {
    // Write payment to Redis
    rdb.HSet(ctx, "payment:"+payment.ID, payment.ToMap())

    // Ensure at least 1 replica has it before returning to client
    n, _ := rdb.Wait(ctx, 1, 500).Result()
    if n == 0 {
        log.Warn("payment not yet replicated, continuing")
    }

    return nil
}

// WAIT overhead: adds latency equal to replication lag
// Typical: 1-5ms in same datacenter
// Trade-off: stronger durability vs higher latency

```

Q114. What is Redis memory optimization with encoding thresholds?

Difficulty: Hard

```

# Encoding thresholds control memory vs performance trade-off

# Hash: use listpack (compact) for small hashes
hash-max-listpack-entries 128 # (default, up to 512 saves significant memory)
hash-max-listpack-value 64    # max value size in listpack

# Experiment: storing user sessions as hashes
# user:42 = {session_id, token, user_id, created_at, ...}
# With listpack: ~200 bytes per user session
# With hashtable: ~800 bytes per user session
# Savings: 75% memory at 128 entries threshold

# Sorted set: use listpack for leaderboards with few entries
zset-max-listpack-entries 128
zset-max-listpack-value 64

# Set: intset for integer sets
set-max-intset-entries 512 # pure integer sets up to 512 elements

# List: quicklist node size
list-max-listpack-size -2 # -2 = 8KB per node

# Test memory usage
redis-cli --bigkeys # find keys using most memory
redis-cli MEMORY USAGE mykey # bytes for specific key
redis-cli OBJECT ENCODING mykey # current encoding

# Memory fragmentation
INFO memory | grep mem_fragmentation_ratio
# > 1.5: high fragmentation → MEMORY PURGE or restart
redis-cli MEMORY PURGE # release fragmented memory (Redis 4.0+)

```

Q115. What is multi-region cache architecture?**Difficulty: Hard**

Problem: single Redis region → high latency for distant users

Architecture options:

1. Read replicas in each region:
 Master (us-east-1) → Replica (eu-west-1) → Replica (ap-southeast-1)
 Writes: always go to master (cross-region latency)
 Reads: from local replica (low latency)
 Consistency: eventual (replica lag 5-50ms)
2. Region-local caches:
 Each region has independent Redis
 Cache miss: read from local DB replica
 Invalidation: cross-region event (Kafka, SNS)
 Consistency: eventual (event propagation lag)
3. CRDTs (Conflict-free Replicated Data Types):
 Redis Enterprise: active-active geo-distribution
 Each region accepts writes, CRDTs merge conflicts
 Suitable: counters, sets, sorted sets

Cache warmup on region startup:

New region: cold cache → thundering herd
 Solution: warm from S3 snapshot or pre-fetch top N keys

Cross-region latency costs:

us-east → eu-west: ~80ms
 us-east → ap-southeast: ~200ms
 Same region: <1ms
 → Never do cross-region cache reads in hot path!

Q116. What are cache metrics and alerting?**Difficulty: Medium**

```
// Key cache metrics to track:

// 1. Hit rate (most important)
hitRate := float64(cacheHits) / float64(cacheHits+cacheMisses)
// Alert: hit rate < 80% (depends on use case)

// 2. Latency percentiles
// p50 < 1ms, p99 < 5ms for Redis

// 3. Memory usage
// Alert: > 80% of maxmemory

// 4. Eviction rate
// redis-cli INFO stats | grep evicted_keys
// Alert: evicted_keys > 0 (means maxmemory reached)

// 5. Connection pool utilization
// Alert: > 80% of pool size used consistently

// Prometheus metrics (go-redis built-in)
rdb := redis.NewClient(&redis.Options{...})
```

```

collector := redisprometheus.NewCollector("namespace", "subsystem", rdb)
prometheus.MustRegister(collector)

// Grafana dashboard metrics:
// redis_connected_clients
// redis_memory_used_bytes / redis_memory_max_bytes
// redis_keyspace_hits_total / (redis_keyspace_hits_total + redis_keyspace_misses_total)
// redis_evicted_keys_total
// redis_commands_duration_seconds_bucket (latency histogram)

// Application-level metrics
var (
    cacheHits = promauto.NewCounter(prometheus.CounterOpts{Name: "cache_hits_total"})
    cacheMisses = promauto.NewCounter(prometheus.CounterOpts{Name: "cache_misses_total"})
    cacheLatency = promauto.NewHistogram(prometheus.HistogramOpts{
        Name:    "cache_operation_duration_seconds",
        Buckets: []float64{.0005, .001, .005, .01, .05},
    })
)

```

Q117. What is cache testing strategies?

Difficulty:
Medium

```

// 1. Unit test with mock cache
type MockCache struct {
    data map[string]interface{}
    mu    sync.Mutex
}

func (m *MockCache) Get(key string) (interface{}, bool) {
    m.mu.Lock(); defer m.mu.Unlock()
    v, ok := m.data[key]
    return v, ok
}

func (m *MockCache) Set(key string, val interface{}, ttl time.Duration) {
    m.mu.Lock(); defer m.mu.Unlock()
    m.data[key] = val
}

// 2. Integration test with miniredis
import "github.com/alicebob/miniredis/v2"

func TestWithRedis(t *testing.T) {
    mr := miniredis.RunT(t) // in-memory Redis server for tests
    rdb := redis.NewClient(&redis.Options{Addr: mr.Addr()})

    svc := NewService(rdb)

    // Test cache hit
    mr.Set("user:42", `{"id":42,"name":"Alice"}`)
    user, err := svc.GetUser(ctx, 42)
    assert.NoError(t, err)
    assert.Equal(t, "Alice", user.Name)

    // Test TTL
    mr.FastForward(time.Hour + time.Second)
    assert.False(t, mr.Exists("user:42"))

    // Test cache miss → DB fallback

```

```
    user, err = svc.GetUser(ctx, 99)
    assert.Equal(t, ErrNotFound, err)
}
```

Q118. ### Q118–Q150: Final Cache Questions

Difficulty:
Medium

Q119. What is distributed lock with Redlock?

Difficulty:
Medium

```
// Redlock: distributed lock using N Redis nodes (usually 5)
// Majority of nodes must grant lock
import "github.com/go-redis/redis/v4"

pool := goredis.NewPool(rdb)
rs := redsync.New(pool)

mutex := rs.NewMutex("critical-section",
    redsync.WithExpiry(15*time.Second),
    redsync.WithRetryDelay(100*time.Millisecond),
    redsync.WithTries(32),
)

if err := mutex.Lock(); err != nil {
    return ErrCouldNotLock
}
defer mutex.Unlock()
// Critical section here
```

Q120. What is Redis backup and disaster recovery?

Difficulty:
Medium

```
# RDB backup (point-in-time snapshot)
redis-cli BGSAVE          # async background save
redis-cli LASTSAVE        # timestamp of last save
# File: /var/lib/redis/dump.rdb

# AOF backup (append-only log)
redis-cli BGREWRITEAOF    # compact AOF file

# Automated backup with cron
0 * * * * redis-cli BGSAVE && cp /var/redis/dump.rdb /backup/redis-$(date +%Y%m%d%H).rdb

# Cloud: Redis ElastiCache automated backups
# Backup window: set in cluster config
# Retention: 1-35 days
# Restore: creates new cluster from backup

# DR strategy:
# Cross-region replica: ElastiCache global datastore
# Manual: backup to S3, restore in target region
# RTO: 30 min (create new cluster + restore)
# RPO: 1 hour (hourly backups) or minutes (with replica)
```

Q121. What is Redis slow log?**Difficulty:**
Medium

```
# Slow log: commands that exceeded threshold
slowlog-log-slower-than 10000 # microseconds (10ms)
slowlog-max-len 128           # keep last 128 entries

# View slow log
SLOWLOG GET 10    # last 10 slow commands
SLOWLOG LEN       # number in slow log
SLOWLOG RESET     # clear slow log

# Entry format:
# 1) ID
# 2) timestamp
# 3) duration (microseconds)
# 4) command + args

# Common causes:
# KEYS * : O(N) scan of all keys – never use in production
# SMEMBERS on large set: O(N)
# SORT without LIMIT: O(N+M*log(M))
# LRANGE 0 -1 on large list: O(N)

# Use instead:
# SCAN cursor MATCH pattern COUNT 100 # iterative, non-blocking
# SSCAN for sets
# HSCAN for hashes
# ZSCAN for sorted sets
```

Q122. ### Q121–Q150: Remaining Cache Questions**Difficulty:**
Medium

Q	Topic
Q121	Redis OBJECT FREQ and LFU eviction
Q122	Cache versioning strategies
Q123	Write coalescing (batch cache writes)
Q124	Redis Sentinel automatic failover
Q125	Cache compression trade-offs
Q126	Redis key expiry notifications in Go
Q127	In-memory cache vs Redis: when to use each
Q128	Cache coherence in microservices
Q129	Redis OBJECT IDLETIME for access patterns
Q130	Distributed caching patterns for APIs
Q131	Cache hit ratio optimization
Q132	Redis Modules: RedisSearch, RedisJSON
Q133	Redis replication buffer tuning
Q134	Cache warming from database dump
Q135	Session caching best practices

Q136	Redis INFO command interpretation
Q137	Cache tier design for e-commerce
Q138	Redis key pattern analysis with redis-cli
Q139	Cache consistency models (strong vs eventual)
Q140	Redis cluster cross-slot operations
Q141	Cache monitoring with Datadog/New Relic
Q142	Redis configuration for low-latency
Q143	Multi-tenant caching (namespace isolation)
Q144	Cache debugging techniques
Q145	Redis Enterprise vs open source Redis
Q146	Cache testing in CI/CD pipelines
Q147	Redis bulk operations (MSET, MGET, pipelines)
Q148	Cache observability (OpenTelemetry)
Q149	Valkey vs Redis (post-license fork)
Q150	Cache production readiness checklist

Q123. What is write-through vs write-behind caching?

Difficulty:
Medium

Write-Through:

- Write to cache AND DB synchronously
- Cache always consistent with DB
- Higher write latency (both writes must complete)
- Use: read-heavy, consistency required

Write-Behind (Write-Back):

- Write to cache only, async write to DB
- Lower write latency (only cache write blocks)
- Risk: data loss if cache fails before DB write
- Use: write-heavy, eventual consistency ok
- Mitigation: WAL (write-ahead log) in cache

Write-Around:

- Write directly to DB, bypass cache
- Cache populated on next read (cache miss)
- Use: infrequently read data, large objects
- Prevents cache pollution from one-time writes

Cache-Aside (Lazy Loading):

- App manages cache explicitly
- Read: cache → DB on miss → populate cache
- Write: update DB + invalidate cache
- Use: most common pattern, flexible

Read-Through:

- Cache sits in front of DB transparently
- Cache fetches from DB on miss automatically
- App only talks to cache
- Use: simplifies app code, uniform caching layer

Q124. What is Redis keyspace notifications?**Difficulty:**
Medium

```
// Keyspace notifications: get events when keys expire/modified
// Config (redis.conf or SET):
// notify-keyspace-events "Ex"
// E = keyspace events, x = expired events

rdb.Do(ctx, "CONFIG", "SET", "notify-keyspace-events", "Ex")

// Subscribe to expiry notifications
pubsub := rdb.Subscribe(ctx, "__keyevent@0__:expired")
defer pubsub.Close()

ch := pubsub.Channel()
for msg := range ch {
    expiredKey := msg.Payload
    log.Printf("key expired: %s", expiredKey)
    // Trigger: cleanup, re-fetch, notify
}

// Other event types:
// K = keyspace events (channel: __keyspace@<db>__:<key>)
// E = keyevent events (channel: __keyevent@<db>__:<event>)
// g = generic commands (del, expire, rename)
// s = set commands
// l = list commands
// z = sorted set commands
// x = expired events
// d = module key events
// t = stream commands

// notify-keyspace-events "KEA" = all events
// Use carefully: high-volume events can overwhelm Redis
```

Q125. What is Redis memory usage analysis?**Difficulty:**
Medium

```
# Memory info
redis-cli INFO memory

# Key fields:
# used_memory:          total allocated by Redis
# used_memory_rss:      RSS from OS (includes fragmentation)
# mem_fragmentation_ratio: RSS/used (> 1.5 = fragmented, < 1 = swapping)
# maxmemory:            limit (0 = no limit)
# mem_allocator:        jemalloc (default, good for fragmentation)

# Find memory hogs
redis-cli --bigkeys          # top 1 key per type by size
redis-cli MEMORY USAGE mykey # bytes for one key
redis-cli MEMORY DOCTOR      # diagnostic suggestions

# Scan all keys and estimate memory (don't use KEYS * in production!)
redis-cli --scan --pattern "user:*" | xargs -I{} redis-cli MEMORY USAGE {} | awk '{sum+=$1} END{print "Total"}'

# Key count by pattern (approximate)
```

```
redis-cli DBSIZE          # total keys in current db

# Memory breakdown:
# MEMORY MALLOC-STATS    # jemalloc stats
# MEMORY PURGE           # return fragmented memory to OS (Redis 4.0+)

# Per-database key counts (Redis has 16 DBs by default)
redis-cli INFO keyspace
```

Q126. What is Redis OBJECT ENCODING for memory optimization?

Difficulty:
Medium

```
# Encoding determines memory usage per data type

# String
SET counter 100
OBJECT ENCODING counter    # "int" (very compact!)
SET name "Alice"
OBJECT ENCODING name       # "embstr" (<= 44 bytes)
SET longtext "...long text..."
OBJECT ENCODING longtext   # "raw"

# Hash
HSET user:1 name Alice age 30
OBJECT ENCODING user:1     # "listpack" (<= 128 fields and <= 64 bytes each)
# Add many fields:
OBJECT ENCODING user:1     # "hashtable" (after threshold exceeded)

# Memory savings (listpack vs hashtable):
# 10-field hash: listpack ~200 bytes, hashtable ~800 bytes

# Tune thresholds for your data:
CONFIG SET hash-max-listpack-entries 256 # allow larger listpack hashes
CONFIG SET hash-max-listpack-value 128   # larger values in listpack

# Sorted Set
ZADD leaderboard 100 "user1"
OBJECT ENCODING leaderboard # "listpack" or "skiplist"

# Always check OBJECT ENCODING after tuning to verify effect
```

Q127. What is Redis vs Memcached comparison?

Difficulty:
Medium

Redis:

- Data structures: string, hash, list, set, zset, stream, geo, bitmap
- Persistence: RDB snapshots + AOF log
- Replication: async master-replica
- Clustering: built-in cluster mode (hash slots)
- Pub/Sub: yes
- Transactions: MULTI/EXEC (optimistic)
- Lua scripting: yes
- Max value size: 512MB
- Eviction: multiple policies (LRU, LFU, TTL-based)
- Use: caching + data structures + sessions + queues + pub/sub

Memcached:

Data structures: string only (key-value)
 Persistence: NO (in-memory only)
 Replication: NO (client-side sharding only)
 Clustering: client-side (consistent hashing)
 Pub/Sub: NO
 Transactions: NO
 Lua scripting: NO
 Max value size: 1MB
 Eviction: LRU only
 Use: pure caching (high performance, simple)

Memcached advantages:

Slightly faster for simple string ops (less overhead)
 Multi-threading (Redis: single-threaded command loop until 6.0)
 Simpler: nothing to misconfigure

Recommendation (2025):

Redis for virtually all new projects
 Memcached: only for very specific high-throughput pure cache needs

Q128. What is Redis connection pooling in Go?

Difficulty:
Medium

```

import "github.com/redis/go-redis/v9"

// go-redis: connection pool built-in
rdb := redis.NewClient(&redis.Options{
    Addr:      "localhost:6379",
    Password:  "",
    DB:        0,

    // Pool settings
    PoolSize:  10,                // connections per CPU (default: 10 * runtime.NumCPU())
    MinIdleConns: 5,            // keep N idle connections ready
    MaxIdleConns: 10,           // max idle connections

    // Timeouts
    DialTimeout: 5 * time.Second,
    ReadTimeout: 3 * time.Second,
    WriteTimeout: 3 * time.Second,
    PoolTimeout: 4 * time.Second, // wait for free connection
    ConnMaxIdleTime: 30 * time.Minute, // close idle after
    ConnMaxLifetime: time.Hour,        // max connection age
})

// Health check
ctx := context.Background()
if err := rdb.Ping(ctx).Err(); err != nil {
    log.Fatalf("redis ping failed: %v", err)
}

// Monitor pool stats
stats := rdb.PoolStats()
fmt.Printf("hits=%d misses=%d timeouts=%d totalConns=%d idleConns=%d\n",
    stats.Hits, stats.Misses, stats.Timeouts, stats.TotalConns, stats.IdleConns)

// Alert: Timeouts > 0 sustained → pool exhausted → increase PoolSize

```

Q129. What is Redis WAIT for read-after-write consistency?**Difficulty:**
Medium

```
// Problem: write to master → immediately read from replica → stale data!
// Read-after-write: user writes then immediately reads own data

// Solution 1: always read from master (simplest, correct)
masterRDB := redis.NewClient(&redis.Options{Addr: "master:6379"})

// Solution 2: WAIT for replica sync
func updateAndRead(ctx context.Context, key, value string) (string, error) {
    if err := masterRDB.Set(ctx, key, value, time.Hour).Err(); err != nil {
        return "", err
    }

    // Wait for at least 1 replica to acknowledge the write
    numReplicas, err := masterRDB.Wait(ctx, 1, 100*time.Millisecond).Result()
    if err != nil || numReplicas < 1 {
        // Replica sync timeout - read from master to be safe
        return masterRDB.Get(ctx, key).Result()
    }

    // Safe to read from replica now
    return replicaRDB.Get(ctx, key).Result()
}

// Solution 3: sticky session routing
// User X always reads from same replica after write
// Route by user_id hash: user writes/reads go to same replica

// Solution 4: version tokens
// Return version on write, pass version on read, replica confirms it has version
```

Q130. What is Redis ACL (Access Control Lists)?**Difficulty:**
Medium

```
# Redis 6.0+: fine-grained access control
# Default user: "default" (backward compat, can restrict)

# Create read-only user for application
redis-cli ACL SETUSER api-reader on \ # enabled
>securepassword \ # password
~data:* \ # can access keys matching data:*
&* \ # can subscribe to all channels
+GET +MGET +HGET +HGETALL \ # allowed commands
+INFO +PING \ # monitoring
-DEBUG # explicit deny

# Create write user for specific namespace
redis-cli ACL SETUSER api-writer on >password ~session:* +SET +GET +DEL +EXPIRE

# Admin user (full access)
redis-cli ACL SETUSER admin on >adminpass ~* &* +@all

# List users
redis-cli ACL LIST
```

```
redis-cli ACL WHOAMI # current user
redis-cli ACL CAT    # list command categories

# Disable default user (security hardening)
redis-cli ACL SETUSER default off

# In go-redis:
redis.NewClient(&redis.Options{
    Username: "api-reader",
    Password: "securepassword",
})
```

Q131. What is Redis sentinel vs cluster for HA?

Difficulty:
Medium

Redis Sentinel:

- Purpose: HA for single master + replicas
- Components: N sentinel processes (recommend: 3)
- Failure detection: majority of sentinels agree (quorum=2)
- Failover: sentinel promotes replica → updates clients
- Scale: vertical only (one master handles writes)
- Use: single-region, < 100GB data, moderate throughput

Redis Cluster:

- Purpose: HA + horizontal scaling
- Data split: 16384 hash slots across N masters
- Each master: has 1-N replicas
- Failover: replica promoted automatically
- Scale: add shards (horizontal)
- Client: must be cluster-aware (MOVED/ASK handling)
- Limitation: multi-key ops only within same slot/hash tag
- Use: > 100GB data, high throughput, multi-region

Choosing:

- < 10GB, simple HA → Sentinel
- > 10GB or > 100K ops/sec → Cluster
- Need cross-slot transactions → Sentinel (single master)
- Massive scale → Cluster

Managed options:

- AWS ElastiCache: both Sentinel and Cluster modes
- Redis Cloud: Redis Enterprise (better than open-source cluster)
- Upstash: serverless Redis (pay per request)

Q132. What is Redis rate limiting with Lua?

Difficulty:
Medium

```
// Lua script: atomic check-and-increment
// Prevents race conditions in distributed rate limiting

const rateLimitScript = `
local key = KEYS[1]
local limit = tonumber(ARGV[1])
local window = tonumber(ARGV[2])
local now = tonumber(ARGV[3])

-- Remove old entries outside window
```

```

redis.call('ZREMRANGEBYSCORE', key, 0, now - window)

-- Count current requests
local count = redis.call('ZCARD', key)

if count < limit then
    -- Add this request
    redis.call('ZADD', key, now, now .. math.random())
    redis.call('EXPIRE', key, window / 1000 + 1)
    return 1 -- allowed
end
return 0 -- denied
`

func isAllowed(ctx context.Context, rdb *redis.Client, userID string, limit int, windowMs int64) bool {
    now := time.Now().UnixMilli()
    key := fmt.Sprintf("ratelimit:%s", userID)

    result, err := rdb.Eval(ctx, rateLimitScript,
        []string{key},
        limit, windowMs, now,
    ).Int()

    if err != nil { return true } // fail open
    return result == 1
}

```

Q133. What is Redis caching for database queries?

Difficulty:
Medium

```

// Pattern: cache DB query results by normalized query hash

type QueryCache struct {
    rdb *redis.Client
    db  *pgxpool.Pool
    ttl time.Duration
}

func (qc *QueryCache) Query(ctx context.Context, sql string, args ...interface{}) ([]map[string]interface{},
    // Create cache key from query + args
    h := sha256.New()
    fmt.Fprintf(h, "%s:%v", sql, args)
    cacheKey := fmt.Sprintf("query:%x", h.Sum(nil))

    // Check cache
    data, err := qc.rdb.Get(ctx, cacheKey).Bytes()
    if err == nil {
        var rows []map[string]interface{}
        if json.Unmarshal(data, &rows) == nil {
            return rows, nil
        }
    }

    // Execute query
    rows, err := qc.executeQuery(ctx, sql, args...)
    if err != nil { return nil, err }

    // Cache result

```

```

    if data, err := json.Marshal(rows); err == nil {
        qc.rdb.Set(ctx, cacheKey, data, qc.ttl)
    }

    return rows, nil
}

// Invalidation: use tags (store set of keys per tag)
func (qc *QueryCache) Invalidate(ctx context.Context, tags ...string) {
    for _, tag := range tags {
        keys, _ := qc.rdb.SMembers(ctx, "tag:"+tag).Result()
        if len(keys) > 0 {
            qc.rdb.Del(ctx, keys...)
        }
        qc.rdb.Del(ctx, "tag:"+tag)
    }
}

```

Q134. What is Redis sorted set for leaderboard?

Difficulty:
Medium

```

// Leaderboard: sorted by score, rank queries efficient

// Add/update score
rdb.ZAdd(ctx, "leaderboard:global", redis.Z{Score: 15420, Member: "user:42"})
rdb.ZIncrBy(ctx, "leaderboard:global", 100, "user:42") // add 100 points

// Get top 10 (highest scores first)
top10, _ := rdb.ZRevRangeWithScores(ctx, "leaderboard:global", 0, 9).Result()
for i, z := range top10 {
    fmt.Printf("#%d: %s → %.0f points\n", i+1, z.Member, z.Score)
}

// Get user's rank (0-indexed)
rank, _ := rdb.ZRevRank(ctx, "leaderboard:global", "user:42").Result()
score, _ := rdb.ZScore(ctx, "leaderboard:global", "user:42").Result()
fmt.Printf("user:42 rank=%d score=%.0f\n", rank+1, score)

// Get neighborhood (user's rank ± 2)
rdb.ZRevRangeWithScores(ctx, "leaderboard:global", rank-2, rank+2)

// Weekly leaderboard (expire after 1 week)
weekKey := fmt.Sprintf("leaderboard:week:%s", time.Now().Format("2006-W01"))
rdb.ZAdd(ctx, weekKey, redis.Z{Score: points, Member: userID})
rdb.Expire(ctx, weekKey, 8*24*time.Hour)

// Count users above threshold
rdb.ZCount(ctx, "leaderboard:global", "1000", "+inf")

```

Q135. What is Redis HyperLogLog for cardinality?

Difficulty:
Medium

```

// HyperLogLog: probabilistic cardinality estimation
// Estimate unique count with ~0.81% error, uses only 12KB per HLL
// Count unique visitors to a page

```



```

rdb.PFAdd(ctx, "visitors:2024-01-15", "user:123", "user:456", "user:789")
rdb.PFAdd(ctx, "visitors:2024-01-15", "user:123") // duplicate, not counted

// Get unique count
count, _ := rdb.PFCount(ctx, "visitors:2024-01-15").Result()
fmt.Printf("unique visitors: %d\n", count) // approximately 3

// Merge multiple HLLs (e.g., weekly from daily)
rdb.PFMerge(ctx, "visitors:2024-W03",
    "visitors:2024-01-15",
    "visitors:2024-01-16",
    "visitors:2024-01-17",
)

weeklyUnique, _ := rdb.PFCount(ctx, "visitors:2024-W03").Result()

// Comparison vs SADD (exact):
// HLL: 12KB for any cardinality, ~1% error
// SET: N bytes per element, exact count
// Use HLL for: daily active users, search queries, A/B test reach
// Use SET for: must-be-exact, small cardinality (< 100K elements)

```

Q136. What is Redis pipeline vs transaction?

Difficulty:
Medium

```

// Pipeline: batch commands, reduce RTT (NOT atomic)
pipe := rdb.Pipeline()
pipe.Set(ctx, "key1", "val1", time.Hour)
pipe.Set(ctx, "key2", "val2", time.Hour)
pipe.Get(ctx, "key1")
cmds, err := pipe.Exec(ctx)
// All 3 commands sent in 1 RTT, but if key1 SET fails, key2 SET still runs

// TxPipeline: pipeline + MULTI/EXEC (atomic)
// If any command fails during EXEC, all commands rolled back
txPipe := rdb.TxPipeline()
txPipe.Set(ctx, "key1", "val1", time.Hour)
txPipe.Set(ctx, "key2", "val2", time.Hour)
_, err = txPipe.Exec(ctx)
// Atomic: either both set or neither

// WATCH + MULTI/EXEC (optimistic locking):
err = rdb.Watch(ctx, func(tx *redis.Tx) error {
    val, err := tx.Get(ctx, "balance").Int()
    if err != nil { return err }

    if val < 100 { return errors.New("insufficient funds") }

    _, err = tx.TxPipelined(ctx, func(pipe redis.Pipeliner) error {
        pipe.Set(ctx, "balance", val-100, 0)
        return nil
    })
    return err
}, "balance")
// If "balance" changed between WATCH and EXEC → retried (optimistic)

```

Q137. What is Redis Bloom filter?**Difficulty:**
Medium

```
# Bloom filter: probabilistic set membership test
# False positives possible, false negatives impossible
# Extremely memory efficient

# Requires RedisBloom module (or Redis Stack)
BF.RESERVE emails 0.01 1000000 # 1% false positive rate, 1M elements
BF.ADD emails "user@example.com"
BF.EXISTS emails "user@example.com" # 1 = likely member, 0 = definitely not

# Use case: prevent duplicate email registration
func canRegister(ctx context.Context, email string) bool {
    exists, _ := rdb.Do(ctx, "BF.EXISTS", "registered-emails", email).Int()
    return exists == 0 // definitely not registered (bloom filter says so)
}

func registerUser(ctx context.Context, email string) error {
    // ... create user in DB ...
    rdb.Do(ctx, "BF.ADD", "registered-emails", email)
    return nil
}

# Cuckoo filter (better: supports deletion)
CF.RESERVE urls 1000000
CF.ADD urls "https://example.com"
CF.DEL urls "https://example.com" # deletion supported!
CF.EXISTS urls "https://example.com"
```

Q138. What is Redis Streams consumer groups in Go?**Difficulty:**
Medium

```
// Create consumer group
rdb.XGroupCreate(ctx, "events", "processors", "0") // "0" = start from beginning
// "$" = only new messages

// Producer: add to stream
rdb.XAdd(ctx, &redis.XAddArgs{
    Stream: "events",
    Values: map[string]interface{}{"type": "order.placed", "id": "42"},
})

// Consumer: read from group
result, _ := rdb.XReadGroup(ctx, &redis.XReadGroupArgs{
    Group:    "processors",
    Consumer: "worker-1",
    Streams:  []string{"events", ">"}, // ">" = undelivered to this group
    Count:    10,
    Block:    2 * time.Second,
}).Result()

for _, stream := range result {
    for _, msg := range stream.Messages {
        if err := process(msg); err == nil {
            rdb.XAck(ctx, "events", "processors", msg.ID) // acknowledge
        }
        // No ACK: message stays in PEL (pending entry list)
    }
}
```

```

    }
}

// Claim stuck messages (consumer crashed without ACK)
pending, _ := rdb.XPendingExt(ctx, &redis.XPendingExtArgs{
    Stream: "events", Group: "processors",
    Idle: 5 * time.Minute, Start: "-", End: "+", Count: 10,
}).Result()
// Claim: rdb.XClaim → reassign to current consumer

```

Q139. What is multi-level caching architecture?

Difficulty:
Medium

L1: In-process cache (ristretto/bigcache)

Latency: ~100ns

Size: 100MB-1GB (bounded by process memory)

Consistency: per-instance (different servers may have different values)

TTL: 30s-5min (short, accept stale)

L2: Redis (shared cache)

Latency: <1ms (same datacenter)

Size: GBs (Redis Cluster for TBs)

Consistency: shared across all instances

TTL: minutes-hours

L3: Database (source of truth)

Latency: 5-50ms

Consistency: always correct

Capacity: unlimited (with sharding)

Request flow:

L1 hit → return (fastest)

L1 miss → L2 hit → populate L1 → return

L2 miss → DB query → populate L2 → populate L1 → return

Invalidation complexity:

L1: invalidate per-process (hard without messaging)

Strategy: short L1 TTL (accept brief stale) + invalidate on write

Redis Pub/Sub: broadcast invalidation to all processes

When to use multi-level:

Very high read throughput (> 100K req/s)

Same data accessed repeatedly in request burst

L2 latency still too high for hot path

Q140. What is Redis for session storage?

Difficulty:
Medium

// Redis: ideal for distributed session storage

// Shared across all app instances, TTL auto-cleanup

```

type SessionStore struct {
    rdb *redis.Client
    ttl time.Duration
}

```

```

func (s *SessionStore) Set(ctx context.Context, sessionID string, data interface{}) error {

```

```

    b, err := json.Marshal(data)
    if err != nil { return err }
    return s.rdb.Set(ctx, "session:"+sessionID, b, s.ttl).Err()
}

func (s *SessionStore) Get(ctx context.Context, sessionID string, out interface{}) error {
    b, err := s.rdb.Get(ctx, "session:"+sessionID).Bytes()
    if errors.Is(err, redis.Nil) { return ErrSessionNotFound }
    if err != nil { return err }
    return json.Unmarshal(b, out)
}

// Extend session TTL on activity
func (s *SessionStore) Touch(ctx context.Context, sessionID string) error {
    return s.rdb.Expire(ctx, "session:"+sessionID, s.ttl).Err()
}

func (s *SessionStore) Delete(ctx context.Context, sessionID string) error {
    return s.rdb.Del(ctx, "session:"+sessionID).Err()
}

// Session ID: generate with crypto/rand
sessionID := make([]byte, 32)
rand.Read(sessionID)
id := base64.URLEncoding.EncodeToString(sessionID)

```

Q141. ### Q141–Q150: Final Cache Questions

Difficulty:
Medium

Q142. What is Redis info command key metrics?

Difficulty:
Medium

```

redis-cli INFO all # all sections

# Server: version, uptime, OS, tcp port
# Clients: connected_clients, blocked_clients
# Memory: used_memory, maxmemory, mem_fragmentation_ratio
# Persistence: rdb_changes_since_last_save, aof_enabled
# Stats: keyspace_hits, keyspace_misses, evicted_keys
# Replication: role, connected_slaves, master_replid, master_repl_offset
# CPU: used_cpu_sys, used_cpu_user
# Keyspace: db0:keys=100000,expires=50000,avg_ttl=3600000

hit_rate = keyspace_hits / (keyspace_hits + keyspace_misses)

```

Q143. What is Valkey vs Redis?

Difficulty:
Medium

Redis relicensed (March 2024): SSPL + RSALv2 (not OSI-approved open source)
Valkey: Linux Foundation fork of Redis 7.2 (truly open source, BSD-3)

Valkey 7.2: feature-compatible with Redis 7.2

Valkey 8.0: new features

- Multi-threaded I/O (significant performance improvement)
- Performance improvements
- Active community (AWS, Google, Snap, Ericsson)

go-redis: supports both Redis and Valkey (same protocol)
 AWS ElastiCache: supports Valkey as drop-in Redis replacement
 DigitalOcean, Aiven: Valkey support added

Migration:

Valkey is wire-compatible with Redis
 Same commands, same protocol (RESP2/RESP3)
 go-redis client works with both
 Drop-in replacement for most use cases

Which to use (2025):

New projects: Valkey (truly OSS, growing community)
 Existing Redis: evaluate based on licensing requirements
 Redis Cloud/Enterprise: Redis (proprietary features like active-active)

Q144. What is cache stampede prevention with probabilistic early expiration?

Difficulty:
Medium

```
// XFetch algorithm: probabilistically refresh before expiry
// Avoids thundering herd on cache expiry

func fetchWithPER(ctx context.Context, key string, delta float64, fetchFn func() ([]byte, error)) ([]byte, error) {
    // Try cache first
    pipe := rdb.Pipeline()
    getCmd := pipe.Get(ctx, key)
    ttlCmd := pipe.TTL(ctx, key)
    pipe.Exec(ctx)

    data, err := getCmd.Bytes()
    ttl, _ := ttlCmd.Result()

    if err == nil && ttl > 0 {
        // Probabilistic early recomputation:
        // Refresh if: -delta * ln(random) > remaining_ttl
        remaining := ttl.Seconds()
        if -delta * math.Log(rand.Float64()) > remaining {
            // Eagerly refresh (beat the expiry)
            goto refresh
        }
        return data, nil
    }

refresh:
    fresh, err := fetchFn()
    if err != nil {
        if data != nil { return data, nil } // serve stale on error
        return nil, err
    }
    rdb.Set(ctx, key, fresh, time.Hour)
    return fresh, nil
}

// delta: typical recomputation time in seconds
// Higher delta = earlier probabilistic refresh
```

Q145. What is cache tagging for bulk invalidation?

Difficulty:
Medium

```
// Problem: how to invalidate all caches related to "user:42"?
// Solution: tag keys and invalidate by tag

func setCacheWithTags(ctx context.Context, key string, data []byte, ttl time.Duration, tags ...string) error {
    pipe := rdb.Pipeline()
    pipe.Set(ctx, key, data, ttl)
    for _, tag := range tags {
        pipe.SAdd(ctx, "tag:"+tag, key)
        pipe.Expire(ctx, "tag:"+tag, ttl+time.Minute)
    }
    _, err := pipe.Exec(ctx)
    return err
}

func invalidateByTag(ctx context.Context, tag string) error {
    keys, err := rdb.SMembers(ctx, "tag:"+tag).Result()
    if err != nil { return err }
    if len(keys) > 0 {
        keys = append(keys, "tag:"+tag)
        return rdb.Del(ctx, keys...).Err()
    }
    return rdb.Del(ctx, "tag:"+tag).Err()
}

// Usage
setCacheWithTags(ctx, "user:42:profile", data, time.Hour, "user:42", "profiles")
setCacheWithTags(ctx, "user:42:orders", orderData, time.Hour, "user:42", "orders")

// On user update: invalidate everything for user:42
invalidateByTag(ctx, "user:42") // deletes both profile and orders cache
```

Q146. What is Redis for distributed locks in Go?

Difficulty:
Medium

```
// Simple distributed lock (single Redis instance)
func acquireLock(ctx context.Context, rdb *redis.Client, key, token string, ttl time.Duration) (bool, error) {
    // SET key token NX PX ttl_ms (atomic: only set if not exists)
    ok, err := rdb.SetNX(ctx, "lock:"+key, token, ttl).Result()
    return ok, err
}

func releaseLock(ctx context.Context, rdb *redis.Client, key, token string) error {
    // Lua: only delete if we own the lock (check token)
    script := `
        if redis.call("get", KEYS[1]) == ARGV[1] then
            return redis.call("del", KEYS[1])
        else
            return 0
        end
    `
    return rdb.Eval(ctx, script, []string{"lock:" + key}, token).Err()
}

// Usage
token := uuid.New().String()
acquired, err := acquireLock(ctx, rdb, "resource-123", token, 30*time.Second)
if !acquired { return ErrLockNotAcquired }
```

```
defer releaseLock(ctx, rdb, "resource-123", token)

// Process critical section
doWork()

// For multi-node: use Redlock (redsync library)
```

Q147. What is cache warm-up strategies?

Difficulty:
Medium

```
// Problem: cold start → all requests miss cache → overload DB

// Strategy 1: Pre-warm from DB on startup
func warmCache(ctx context.Context, rdb *redis.Client, db *pgxpool.Pool) error {
    rows, err := db.Query(ctx, "SELECT id, data FROM products WHERE active=true ORDER BY view_count DESC LIMIT 1000")
    if err != nil { return err }
    defer rows.Close()

    pipe := rdb.Pipeline()
    for rows.Next() {
        var id int64
        var data []byte
        rows.Scan(&id, &data)
        pipe.Set(ctx, fmt.Sprintf("product:%d", id), data, 24*time.Hour)
    }
    _, err = pipe.Exec(ctx)
    return err
}

// Strategy 2: Shadow traffic (replay prod traffic to new instance)
// Strategy 3: Cache seeding from S3 snapshot
// Strategy 4: Gradual rollout (small % traffic first)
// Strategy 5: Request coalescing (singleflight during warm-up)

var sfGroup singleflight.Group
func getProduct(ctx context.Context, id int64) (*Product, error) {
    result, err, _ := sfGroup.Do(fmt.Sprintf("product:%d", id), func() (interface{}, error) {
        return fetchFromDB(ctx, id)
    })
    return result.(*Product), err
}
```

Q148. What is Redis for feature flags?

Difficulty:
Medium

```
// Feature flags stored in Redis: instant rollout without deploy

type FeatureFlags struct {
    rdb *redis.Client
}

// Set feature flag
func (ff *FeatureFlags) Enable(ctx context.Context, flag string) error {
    return ff.rdb.Set(ctx, "feature:"+flag, "true", 0).Err()
}

func (ff *FeatureFlags) Disable(ctx context.Context, flag string) error {
    return ff.rdb.Del(ctx, "feature:"+flag).Err()
}
```

```

}
func (ff *FeatureFlags) IsEnabled(ctx context.Context, flag string) bool {
    val, _ := ff.rdb.Get(ctx, "feature:"+flag).Result()
    return val == "true"
}

// Percentage rollout
func (ff *FeatureFlags) IsEnabledForUser(ctx context.Context, flag string, userID int64) bool {
    // Get rollout percentage
    pct, err := ff.rdb.Get(ctx, "feature:"+flag+":pct").Int()
    if err != nil { return false }
    // Deterministic: same user always gets same result
    return (userID % 100) < int64(pct)
}

// Listen for changes (keyspace notifications)
func (ff *FeatureFlags) WatchChanges(ctx context.Context, onChange func(flag string)) {
    pubsub := ff.rdb.Subscribe(ctx, "__keyevent@0__:set", "__keyevent@0__:del")
    for msg := range pubsub.Channel() {
        if strings.HasPrefix(msg.Payload, "feature:") {
            onChange(strings.TrimPrefix(msg.Payload, "feature:"))
        }
    }
}

```

Q149. What is Redis for search autocomplete?

Difficulty:
Medium

```

// Autocomplete using sorted set with score=0 trick (lexicographic range)

// Index: ZADD index 0 "apple" 0 "application" 0 "apply"
func indexWord(ctx context.Context, rdb *redis.Client, word string) error {
    return rdb.ZAdd(ctx, "autocomplete", redis.Z{Score: 0, Member: word}).Err()
}

// Search prefix
func autocomplete(ctx context.Context, rdb *redis.Client, prefix string, limit int) ([]string, error) {
    min := "[" + prefix // inclusive lower bound
    max := "[" + prefix + "ÿ" // prefix + max char (all strings starting with prefix)

    results, err := rdb.ZRangeByLex(ctx, "autocomplete", &redis.ZRangeBy{
        Min:    min,
        Max:    max,
        Count:  int64(limit),
    }).Result()
    return results, err
}

// With frequency scoring (better UX)
func recordSearch(ctx context.Context, rdb *redis.Client, query string) error {
    return rdb.ZIncrBy(ctx, "search:frequency", 1, query).Err()
}

func autocompleteByFrequency(ctx context.Context, rdb *redis.Client, prefix string) ([]string, error) {
    // Use trie or scan-based approach for prefix + score sort
    // Or: maintain separate sorted set per first character
    return rdb.ZRevRangeByScore(ctx, "search:frequency:"+prefix[0:1],

```



```

    &redis.ZRangeBy{Min: "-inf", Max: "+inf", Count: 10}).Result()
}

```

Q150. What is Redis for job queues?

Difficulty:
Medium

```

// Redis as job queue (simple, without Kafka/RMQ overhead)

// Enqueue job
func enqueue(ctx context.Context, rdb *redis.Client, queue, jobJSON string) error {
    return rdb.LPush(ctx, "queue:"+queue, jobJSON).Err()
}

// Dequeue (blocking pop, waits for jobs)
func dequeue(ctx context.Context, rdb *redis.Client, queue string) (string, error) {
    result, err := rdb.BRPop(ctx, 10*time.Second, "queue:"+queue).Result()
    if err != nil { return "", err }
    return result[1], nil // result[0] is queue name, result[1] is value
}

// Reliable queue (BRPOPLUSH: atomic move to processing list)
func dequeueReliable(ctx context.Context, rdb *redis.Client, queue string) (string, error) {
    return rdb.BRPopLPush(ctx, "queue:"+queue, "queue:"+queue+":processing", 10*time.Second).Result()
}

// Acknowledge (remove from processing list)
func ack(ctx context.Context, rdb *redis.Client, queue, job string) error {
    return rdb.LRem(ctx, "queue:"+queue+":processing", 1, job).Err()
}

// Recover stale jobs (crashed workers)
func recoverStale(ctx context.Context, rdb *redis.Client, queue string) error {
    jobs, _ := rdb.LRange(ctx, "queue:"+queue+":processing", 0, -1).Result()
    for _, job := range jobs {
        rdb.RPush(ctx, "queue:"+queue, job)
    }
    return rdb.Del(ctx, "queue:"+queue+":processing").Err()
}

```

Q151. What is cache production readiness checklist?

Difficulty:
Medium

Configuration:

- maxmemory set (never let Redis use unlimited memory)
- maxmemory-policy = allkeys-lru (or allkeys-lfu)
- Persistence configured (RDB for snapshots, AOF for durability)
- bind configured (not exposed to public internet)
- requirepass set (authentication)
- ACL configured (least privilege)
- notify-keyspace-events disabled (unless needed, adds overhead)

HA:

- Sentinel (3 nodes) or Cluster mode for HA
- Replica in separate AZ
- Automatic failover tested

Monitoring:

- Prometheus exporter (redis-exporter)
- Alert: `used_memory > 80% maxmemory`
- Alert: `evicted_keys > 0` (cache under pressure)
- Alert: `connected_clients` approaching max
- Alert: `replication_lag > 10s`
- Hit rate dashboard: `keyspace_hits / (hits + misses)`

Application:

- Connection pool properly sized
- Timeouts configured (read/write/pool)
- Retry logic with exponential backoff
- Circuit breaker (fail gracefully without cache)
- Cache key namespacing (environment prefix)
- TTL on all keys (no immortal keys)
- Singleflight for thundering herd prevention

Security:

- TLS enabled (Redis 6.0+)
- Redis not exposed to internet
- No sensitive data in keys (PII in values only, encrypted if needed)