

Go (Golang)

SDE2 Interview Guide

100 Questions & Answers with Code Examples

Level: Mid (2–4 yrs)	Stack: Go / Golang	Roles: SDE2 / Backend
----------------------	--------------------	-----------------------

#	Topic	Questions
1	Goroutines & Concurrency	Q1–Q15
2	Channels & Select	Q16–Q28
3	Runtime & Memory	Q29–Q40
4	Interfaces & Types	Q41–Q52
5	Error Handling	Q53–Q60
6	Context	Q61–Q67
7	Generics	Q68–Q74
8	Testing & Benchmarking	Q75–Q82
9	Go Patterns	Q83–Q92
10	Miscellaneous & Advanced	Q93–Q100

> **Focus:** SQS Fundamentals, Standard vs FIFO, Visibility Timeout, DLQ, SNS+SQS, SQS vs Kafka, Lambda Triggers, Go AWS SDK | **Level:** SDE2

Table of Contents

[Standard vs FIFO Queues](#2-standard-vs-fifo-queues) — Q21–Q35
[Visibility Timeout & Message Lifecycle](#3-visibility-timeout--message-lifecycle) — Q36–Q50
[Dead Letter Queues & Error Handling](#4-dead-letter-queues--error-handling) — Q51–Q65
[SNS + SQS Fan-Out Patterns](#5-sns--sqs-fan-out-patterns) — Q66–Q75
[SQS vs Kafka & Architecture Decisions](#6-sqs-vs-kafka--architecture-decisions) — Q76–Q85
[Lambda Triggers, Go SDK & Production Patterns](#7-lambda-triggers-go-sdk--production-patterns) — Q86–Q100

SQS Fundamentals

Q1. What is Amazon SQS?

Difficulty: Easy

Amazon SQS (Simple Queue Service) is a fully managed message queuing service. It decouples producers and consumers, enabling asynchronous processing without managing queue infrastructure.

Producer → [SQS Queue] → Consumer(s)

Key properties:

- Fully managed: no servers to provision, patch, or scale
- Highly available: 3 AZ redundancy within a region
- Scalable: handles any message volume automatically
- At-least-once delivery: messages delivered at least once (Standard queue)
- Pull-based: consumers poll the queue (not push)
- Pay-per-use: charged per API call and data transfer

Message limits:

- Max message size: 256 KB
- Retention: 1 minute to 14 days (default 4 days)
- Max in-flight (Standard): 120,000
- Max in-flight (FIFO): 20,000

Pricing (us-east-1):

- First 1M requests/month: free
- Standard: \$0.40 per million requests
- FIFO: \$0.50 per million requests

Q2. What are the core SQS operations?

Difficulty: Easy

SendMessage: publish message to queue
ReceiveMessage: poll queue for up to 10 messages
DeleteMessage: acknowledge and remove message after processing
ChangeMessageVisibility: extend or reduce visibility timeout
GetQueueAttributes: queue metadata (depth, ARN, etc.)

Flow:

1. Producer: SendMessage (message enters queue)
 2. Consumer: ReceiveMessage (message becomes invisible)
 3. Consumer: process message
 4. Consumer: DeleteMessage (remove permanently)
- OR
- 4b: Visibility timeout expires → message becomes visible again → retry

Important: `ReceiveMessage` does NOT delete the message!

You must explicitly call `DeleteMessage` after successful processing

This enables at-least-once delivery (safe retry on failure)

Batch operations (up to 10 messages per call):

`SendMessageBatch`: send up to 10 messages in one API call

`DeleteMessageBatch`: delete up to 10 messages in one call

`ChangeMessageVisibilityBatch`: change timeout for up to 10 messages

→ Reduces API calls, lowers cost, improves throughput

Q3. What is the SQS message structure?

Difficulty: Easy

SQS Message fields:

`MessageId`: unique identifier assigned by SQS

`ReceiptHandle`: token for `DeleteMessage`/`ChangeVisibility` (changes each receive)

`Body`: message content (string, max 256KB)

`MD5OfBody`: checksum for data integrity verification

`Attributes`: system attributes (`ApproximateReceiveCount`, `SentTimestamp`, etc.)

`MessageAttributes`: custom metadata (type-safe key-value pairs)

`MessageAttributes` (custom metadata):

`Name`: string key (max 256 chars)

`Type`: String | Number | Binary

`Value`: the data

// Send with attributes:

```
{
  "Body": "{\"orderId\": 42, \"amount\": 99.99}",
  "MessageAttributes": {
    "EventType": {
      "DataType": "String",
      "StringValue": "order.created"
    },
    "TraceId": {
      "DataType": "String",
      "StringValue": "abc-123-xyz"
    }
  }
}
```

// System attributes on received message:

`ApproximateReceiveCount`: how many times received (retry count)

`ApproximateFirstReceiveTimestamp`: first time message was received

`SentTimestamp`: when message was sent

`SenderId`: IAM user/role that sent it

Q4. What is long polling vs short polling?

Difficulty: Medium

Short polling (`WaitTimeSeconds=0`, default):

`ReceiveMessage` returns IMMEDIATELY

Samples only a subset of servers

May return empty even when messages exist

Higher cost: many API calls return nothing

Long polling (`WaitTimeSeconds=1-20`):

`ReceiveMessage` WAITS up to 20 seconds for messages to arrive

- Queries ALL servers (not just a subset)
- Returns empty only if no messages arrive during wait period
- Lower cost: fewer empty API calls
- Recommended for all production use

Enable long polling:

- Queue-level: `QueueAttribute ReceiveMessageWaitTimeSeconds=20`
- Per-request: `WaitTimeSeconds=20` in `ReceiveMessage` call

Cost comparison example (100 messages/hour, empty otherwise):

- Short polling: $3600 \times \$0.40/M$ = minimal but wasteful
- Long polling: $300 \times \$0.40/M$ = 12× fewer API calls, same messages

Always use long polling (`WaitTimeSeconds=20`) in production:

- Reduces cost
- Reduces CPU on consumer
- Better latency (detects message arrival instantly vs polling delay)

Q5. What is the SQS visibility timeout?

Difficulty:
Medium

Visibility timeout: period during which a received message is invisible

- After `ReceiveMessage`: message hidden from other consumers

- Consumer processes: if success → `DeleteMessage` (permanent removal)

- If timeout expires without `DeleteMessage`: message becomes visible again → retry

Default: 30 seconds

Range: 0 seconds to 12 hours

Queue-level default + per-message override possible

Timeline:

- t=0: Message arrives in queue
- t=10: Consumer A receives message (invisible for 30s)
- t=25: Consumer A still processing...
- t=35: Timeout expires! Message visible again
- t=36: Consumer B receives the same message (duplicate!)
- t=40: Consumer A finishes and calls `DeleteMessage` (409 error - receipt handle invalid)
- t=41: Consumer B finishes and calls `DeleteMessage` (success)

Setting correctly:

- `visibility_timeout >= max_processing_time × 6`
- If processing takes 2 minutes: set to 12+ minutes

Extending timeout dynamically:

- Before timeout expires: call `ChangeMessageVisibility`
- Add more time when processing will take longer than expected
- Max: total visibility cannot exceed 12 hours

Q6. What are SQS queue attributes?

Difficulty:
Medium

Key queue attributes:

`VisibilityTimeout` (default 30s, max 12h):

- How long a received message is invisible

`MessageRetentionPeriod` (default 4 days, 1min-14days):

How long SQS keeps undeleted messages

MaximumMessageSize (default 256KB, 1KB-256KB):
Max message body size

DelaySeconds (default 0, max 15min):
Delivery delay for new messages (initially invisible)

ReceiveMessageWaitTimeSeconds (default 0, max 20s):
Long polling wait time at queue level

RedrivePolicy:
maxReceiveCount: failures before → DLQ
deadLetterTargetArn: DLQ queue ARN

Policy:
Resource-based policy (who can send/receive)

ApproximateNumberOfMessages: messages in queue
ApproximateNumberOfMessagesNotVisible: in-flight (being processed)
ApproximateNumberOfMessagesDelayed: delayed messages

Monitor these metrics in CloudWatch:
NumberOfMessagesSent
NumberOfMessagesDeleted
NumberOfMessagesReceived
ApproximateNumberOfMessagesVisible → alert if growing (consumers lagging)
ApproximateAgeOfOldestMessage → alert if too high

Q7. What is message deduplication in SQS?

Difficulty:
Medium

Standard queues: NO deduplication (at-least-once delivery)
Duplicate messages possible (rare but can happen)
Consumer MUST be idempotent

FIFO queues: deduplication within 5-minute window
ContentBasedDeduplication: SQS computes SHA-256 of body → dedup key
MessageDeduplicationId: custom dedup key (override)
If same dedup ID received within 5 min → silently discarded

Content-based deduplication:
Enable: ContentBasedDeduplication=true on FIFO queue
SQS hashes message body → rejects same body within 5 minutes
Limitation: same content = rejected even if genuinely different message
Better: set unique MessageDeduplicationId per logical message

Custom deduplication ID:
SendMessage with MessageDeduplicationId: "order-42-created-2024-01-15T10:30:00Z"
Same ID in 5 min window → second message silently ignored
First message already delivered → idempotent

Application-level idempotency (for Standard queues):
Check if event already processed before doing work
Store processed message IDs in DynamoDB/Redis with TTL
Conditional DB update: only apply if state matches expected state

Q8. What is the SQS Extended Client Library?**Difficulty:**
Medium

Problem: SQS max message size = 256 KB
 Large payloads (images, large JSON, PDFs) don't fit

Solution: SQS Extended Client Library
 Stores large payload in S3
 Sends SQS message with S3 reference (pointer)
 Consumer: reads SQS message → fetches from S3 → processes

Flow:

Producer:
 payload > 256KB → upload to S3 → send SQS message with S3 key
 payload ≤ 256KB → send directly in SQS message

Consumer:
 Receive SQS message
 If S3 pointer: download from S3 → get full payload
 Process payload
 Delete SQS message + optionally delete S3 object

Message format (with S3 reference):

```
[
  "software.amazon.payloadoffloading.PayloadS3Pointer",
  {"s3BucketName": "my-bucket", "s3Key": "messages/uuid"}
]
```

Manual implementation in Go (without official library):

```
const maxSQSSize = 256 * 1024
if len(payload) > maxSQSSize {
    key := uploadToS3(payload)
    sqsBody = `{"s3_key":"` + key + `"}`
} else {
    sqsBody = string(payload)
}
```

Q9. What is the ApproximateNumberOfMessages metric?**Difficulty: Easy**

CloudWatch metrics for SQS monitoring:

ApproximateNumberOfMessagesVisible:
 Messages available for retrieval (queue depth)
 Growing → consumers can't keep up → scale consumers
 Alert: > threshold for > 5 minutes

ApproximateNumberOfMessagesNotVisible:
 In-flight: received but not yet deleted
 Max: 120,000 (Standard), 20,000 (FIFO)
 High value → slow processing or processing hanging

ApproximateNumberOfMessagesDelayed:
 Waiting due to delay policy
 Expected for delayed queues, alert if unexpected

ApproximateAgeOfOldestMessage:
 Seconds since oldest unprocessed message was sent

High → messages building up (backlog)
 Alert: $> \text{retention_period} \times 0.8$ (risk of message expiry)

NumberOfMessagesSent: producer throughput
 NumberOfMessagesDeleted: consumer throughput
 NumberOfMessagesReceived: receive API calls

Auto-scaling with KEDA (Kubernetes):
 Scale consumer pods based on ApproximateNumberOfMessagesVisible
 lagThreshold: 100 → target 1 message per pod per second

Auto-scaling with AWS Application Auto Scaling:
 Scaling policy: scale ECS tasks when queue depth > 100

Q10. What are SQS message timers (delay queues)?

Difficulty:
Medium

Delay Queues: messages initially invisible for N seconds after send
 All messages delayed → use Queue DelaySeconds attribute
 Queue-level: 0-900 seconds (15 minutes) default delay

Per-message delay:
 SendMessage with DelaySeconds parameter
 Individual message delayed, not entire queue
 Max: 900 seconds (15 minutes)
 Standard queues: override queue default
 FIFO queues: not supported (use queue-level only)

Use cases:
 Retry with backoff: delay retried messages
 Scheduled tasks: "send email 5 min after registration"
 Webhook retries: delay before retrying failed webhook
 Order pipeline: "charge card 10 min after order placed" (fraud window)

Limitation: max 15 minutes
 For longer delays: use EventBridge Scheduler or Step Functions Wait state

```
// Go: send with 5 minute delay
svc.SendMessage(&sqs.SendMessageInput{
    QueueUrl:      &queueURL,
    MessageBody:   aws.String(body),
    DelaySeconds:  aws.Int64(300), // 5 minutes
})
```

Pattern: exponential backoff retries via delay
 Attempt 1: delay 0s
 Attempt 2: delay 30s
 Attempt 3: delay 120s
 Attempt 4: delay 300s (5min max for SQS)
 Attempt 5+: → Dead Letter Queue

Q11. What is the SQS message receipt handle?

Difficulty:
Medium

ReceiptHandle: token returned with ReceiveMessage
 Required for: DeleteMessage, ChangeMessageVisibility
 NOT the MessageId (different field!)

Changes each time message is received (new receipt handle)

Why it changes:

```
Prevents "stale" operations from old consumers
Consumer A receives (handle-A), times out
Consumer B receives (handle-B)
Consumer A tries DeleteMessage with handle-A → error (stale handle)
Consumer B deletes with handle-B → success
```

ReceiptHandle is 1000+ characters (internal SQS metadata)

```
// Go: store receipt handle from receive, use for delete
result, _ := svc.ReceiveMessage(&sqs.ReceiveMessageInput{...})
for _, msg := range result.Messages {
    handle := msg.ReceiptHandle // save this!

    if err := process(*msg.Body); err == nil {
        svc.DeleteMessage(&sqs.DeleteMessageInput{
            QueueUrl:      &queueURL,
            ReceiptHandle: handle, // use saved handle
        })
    }
}
```

Important: receipt handle is per-consumer, per-receive
 Multiple consumers: each gets different handle for same message
 Don't share receipt handles across goroutines

Q12. What is SQS access control?

Difficulty:
Medium

Two mechanisms for SQS access control:

1. IAM Policies (identity-based):
 Attach to IAM user/role/group


```
{
  "Effect": "Allow",
  "Action": ["sqs:SendMessage", "sqs:ReceiveMessage", "sqs:DeleteMessage"],
  "Resource": "arn:aws:sqs:us-east-1:123456789:my-queue"
}
```
2. Queue Policy (resource-based):
 Attached to the queue itself
 Controls who can access this queue
 Required for: cross-account access, SNS publish

Example: allow SNS to publish to SQS queue:

```
{
  "Effect": "Allow",
  "Principal": {"Service": "sns.amazonaws.com"},
  "Action": "sqs:SendMessage",
  "Resource": "arn:aws:sqs:...:my-queue",
  "Condition": {
    "ArnEquals": {"aws:SourceArn": "arn:aws:sns:...:my-topic"}
  }
}
```

VPC Endpoint for SQS:

Access SQS without internet (private network)
com.amazonaws.region.sqs endpoint
Policy restricts to specific queues

Least privilege:

Producer: sqs:SendMessage only
Consumer: sqs:ReceiveMessage + sqs:DeleteMessage + sqs:ChangeMessageVisibility
Admin: sqs:* (only for management tooling)

Q13. What is SQS server-side encryption?

Difficulty: Easy

SQS SSE: encrypt messages at rest using KMS
All messages encrypted before storing
Decrypted automatically on ReceiveMessage
Transparent to application code

KMS key options:

AWS managed key (aws/sqs): free, simpler
Customer managed key (CMK): control over key, rotation, audit

Enable encryption:

Queue attribute: SqsManagedSseEnabled=true (AWS managed)
Or: KmsMasterKeyId=<key-id> (CMK)

Encryption scope:

At-rest: stored in SQS infrastructure
In-transit: always TLS (HTTPS endpoints)

KMS key policy required:

Key policy must allow SQS to use it

Performance:

KMS calls: one call per unique (queue, key) pair cached for 5 minutes
Not one KMS call per message (cached)
High-throughput: minimal KMS latency impact

Compliance use cases:

PCI DSS: encrypt cardholder data in transit through SQS
HIPAA: encrypt PHI in message queues
GDPR: encrypt personal data at rest

Q14. What are SQS limits you must know?

Difficulty: Easy

Message limits:

Max size: 256 KB (use S3 Extended Client for larger)
Retention: 1 min to 14 days (default 4 days)
Delay: 0 to 900 seconds (15 minutes)
Visibility timeout: 0 to 12 hours (default 30 seconds)

Throughput limits:

Standard queue: unlimited messages/second
FIFO queue: 300 msg/sec (basic)
3,000 msg/sec (with batching, 10 msg/batch)

In-flight limits:

Standard: 120,000 messages in-flight simultaneously

FIFO: 20,000 messages in-flight simultaneously

Batch limits:

SendMessageBatch: 10 messages, total 256 KB

DeleteMessageBatch: 10 messages

ReceiveMessage: up to 10 messages per call

Per-account limits:

Default: 1,000 queues per region per account (can request increase)

Active message quota: up to 120,000 inflight Standard

Long polling:

Max WaitTimeSeconds: 20 seconds

Message deduplication (FIFO):

5-minute deduplication window

Q15. ### Q15-Q20: More SQS Fundamentals

Difficulty:
Medium

Q	Topic
Q15	SQS pricing model and cost optimization strategies
Q16	SQS message filtering vs routing
Q17	SQS purge queue: when and how to use it
Q18	SQS queue naming conventions and tagging
Q19	SQS monitoring with CloudWatch dashboards
Q20	SQS vs SNS vs EventBridge: which to use

Standard vs FIFO Queues

Q16. What is the difference between Standard and FIFO queues?

Difficulty:
Medium

	Standard	FIFO
Ordering	Best-effort	Strict (FIFO within group)
Delivery	At-least-once	Exactly-once
Deduplication	None	5-min window
Throughput	Unlimited	300/sec (3000 with batching)
Cost	\$0.40/M	\$0.50/M (25% more)
Use case	High throughput	Ordering + no duplicates
Name suffix	any name	must end with .fifo

Standard queue:

May deliver out of order (rare but possible)

May deliver same message more than once (consumer must be idempotent)

Unlimited throughput: burst millions of messages

Most use cases

FIFO queue:

Strict ordering: first sent = first received (within message group)
 Exactly-once: deduplication by MessageDeduplicationId
 Lower throughput: 300 TPS (3,000 TPS with batching)
 Required: financial transactions, ordered event streams, idempotent workflows

Decision rule:

"Do I need strict ordering OR guaranteed no duplicates?"

YES → FIFO

NO → Standard (simpler, cheaper, higher throughput)

Q17. What is a FIFO queue message group?

Difficulty: Hard

MessageGroupId: logical partition within FIFO queue
 Messages with same GroupId: strictly ordered (FIFO)
 Messages with different GroupId: can be processed in parallel

Without MessageGroupId:

All messages in one implicit group → sequential processing
 Only 1 consumer active at a time → low throughput

With MessageGroupId:

N groups → N consumers can process in parallel
 Ordering preserved per group, not globally

Example: order processing

MessageGroupId = customer_id
 All orders for customer-42 are processed in order
 Orders for customer-99 processed simultaneously by another consumer

Customer 42: create → add_item → checkout (in order)
 Customer 99: create → checkout (parallel with customer 42's orders)

Partitioning strategy:

High cardinality: each entity gets own group → max parallelism
 Low cardinality: few groups → serialized within group → less parallelism

Too few groups: bottleneck
 Too many groups: fine, just overhead of group tracking

FIFO In-flight limit:

20,000 messages in-flight per queue
 Per group: unlimited in-flight (but ordered delivery within group)

Q18. When should you use FIFO over Standard?

Difficulty: Medium

Use FIFO when:

- Order matters absolutely:
 - Bank account transactions: deposit before withdrawal
 - Order state machine: created → paid → shipped → delivered
 - Event sourcing: events must be applied in sequence
- Exactly-once processing required:
 - Payment processing: charge card exactly once per order
 - Inventory decrement: can't decrement twice for same order
 - Email sending: don't send duplicate confirmation emails

3. Idempotency is hard to implement:
Some operations are inherently non-idempotent
FIFO's deduplication handles it at infrastructure level

Use Standard when:

1. High throughput needed (>30000 TPS):
Log ingestion, metrics collection, analytics events
2. Order doesn't matter:
Email notifications (order doesn't affect outcome)
Image processing (each image independent)
Background jobs (unrelated tasks)
3. Consumer is naturally idempotent:
"Set user status to active" → safe to run twice
4. Cost-sensitive:
Standard is 20% cheaper

Rule of thumb:

Most backend jobs → Standard
Financial/transactional workflows → FIFO
Event sourcing → FIFO
Parallel processing of independent tasks → Standard

Q19. What are FIFO queue best practices?

Difficulty:
Medium

1. Set MessageDeduplicationId explicitly (don't rely on content-based):
UUID per logical operation: "order-42-created-" + uuid.New().String()
Predictable deduplication: same logical event = same dedup ID
Content-based dedup: risk of false positives (same body, different event)
 2. Choose MessageGroupId cardinality carefully:
Per-entity (user/order ID): good parallelism, ordered per entity
Global (all messages same group): sequential → use only if truly needed
 3. Handle FIFO throughput limits:
3,000 TPS with batching (10 msg/batch)
If more needed: multiple FIFO queues + routing layer
Or: reconsider if Standard queue fits (most cases it does)
 4. Consumer concurrency for FIFO:
One consumer per active message group (no parallel within group)
N consumers for N groups = N parallel streams
Don't consume same group ID in parallel (breaks ordering)
 5. Monitoring FIFO-specific:
NumberOfMessageGroupsWithInflightMessages
NumberOfDeduplicatedMessages (are duplicates being filtered?)
 6. FIFO queue naming:
Must end with .fifo: my-orders-queue.fifo
Important: include environment: orders-prod.fifo
- // Go: send to FIFO queue

```
svc.SendMessage(&sqs.SendMessageInput{
    QueueUrl:      &queueURL,
    MessageBody:    aws.String(body),
    MessageGroupId: aws.String(customerID),
    MessageDeduplicationId: aws.String(orderID + "-created"),
})
```

Q20. ### Q25-Q35: More Standard vs FIFO Topics

Difficulty:
Medium

Q	Topic
Q25	FIFO queue consumer group design patterns
Q26	Standard queue ordering: why it's not guaranteed
Q27	Migrating from Standard to FIFO queue
Q28	FIFO high-throughput mode (3000 TPS) configuration
Q29	SQS FIFO with Lambda: concurrency considerations
Q30	Duplicate handling strategies for Standard queues
Q31	FIFO deduplication window: what happens after 5 minutes
Q32	SQS FIFO vs Kinesis: ordering comparison
Q33	Multiple FIFO queues for higher throughput patterns
Q34	FIFO queue with DLQ configuration
Q35	Exactly-once semantics: SQS FIFO limitations

Visibility Timeout & Message Lifecycle

Q21. How do you set the right visibility timeout?

Difficulty:
Medium

Rule: $\text{visibility_timeout} \geq \text{max_processing_time} \times 6$

Multiplier provides buffer for:

- Slow processing (occasional)
- Network latency
- Retry overhead

Examples:

Processing takes ~10s → timeout = 60s

Processing takes ~1min → timeout = 6 minutes

Processing takes ~5min → timeout = 30 minutes

Processing takes ~1h → timeout = 6 hours (close to max 12h)

Too short:

Message redelivered while still being processed

Duplicate processing

Extra load on consumers (processing same message multiple times)

Too long:

- Failure detection delayed (consumer crashes → message invisible for hours)
- Queue appears empty but has in-flight messages

Dynamic visibility timeout (best practice):

- Extend before expiry if still processing
- Heartbeat pattern: goroutine extends timeout every 30s during processing

```
func processWithHeartbeat(msg *sqs.Message) error {
    done := make(chan struct{})
    go func() { // heartbeat goroutine
        ticker := time.NewTicker(30 * time.Second)
        defer ticker.Stop()
        for {
            select {
            case <-done: return
            case <-ticker.C:
                svc.ChangeMessageVisibility(&sqs.ChangeMessageVisibilityInput{
                    QueueUrl:      &queueURL,
                    ReceiptHandle:   msg.ReceiptHandle,
                    VisibilityTimeout: aws.Int64(60),
                })
            }
        }
    }()
    defer close(done)
    return doWork(msg)
}
```

Q22. What is the message lifecycle in SQS?

Difficulty:
Medium

Message lifecycle states:

Available (visible):

- Message sits in queue waiting for consumer
- ReceiveMessage will return this message

In-flight (invisible):

- Message received but not deleted
- Other consumers cannot receive it
- Visibility timeout counting down

Delayed:

- Message not yet available due to delay policy
- Timer running; becomes Available when timer expires

Dead (in DLQ):

- After maxReceiveCount failures → moved to DLQ
- Stays in DLQ until manually processed or expired

Deleted:

- Consumer called DeleteMessage
- Message permanently removed from queue

Timing:

- t=0: SendMessage → Available

```

t=5:  ReceiveMessage → In-flight (invisible for timeout period)
t=35:  VisibilityTimeout=30 expires without DeleteMessage → Available again
t=36:  Different consumer receives → In-flight
t=40:  DeleteMessage → Deleted (permanently gone)

```

If message received `maxReceiveCount` times without delete → DLQ
`ApproximateReceiveCount` attribute tracks retry count

Q23. What happens when a consumer crashes?

Difficulty:
Medium

Consumer crashes during processing:

Scenario:

```

t=0:  Consumer receives message (visibility=60s)
t=10: Consumer crashes (process killed, OOM, network loss)
t=60: VisibilityTimeout expires → message becomes visible again
t=61: Another consumer picks up message
t=75: Second consumer processes successfully → DeleteMessage

```

Result: at-least-once delivery

Message processed at least once (by the second consumer)
 If first consumer had partially applied changes: need idempotency

Without `DeleteMessage`: safe to retry

If consumer crashes after work but before `DeleteMessage`:
 → Message retried → must be idempotent

Consumer restart strategy:

On startup: check in-progress work table → clean up partial work
 Then: start consuming from SQS

Kubernetes pod crash:

Pod restarts automatically (`restartPolicy: Always`)
 SQS message timeout ensures retry even if pod takes time to restart
 Set `visibility_timeout > pod_restart_time + processing_time`

Monitoring crashes:

`ApproximateReceiveCount > 1` → message being retried
 Alert when specific messages hit `maxReceiveCount` → DLQ
 CloudWatch: `NumberOfMessagesReceived` without corresponding `DeleteMessage`

Q24. What is `ChangeMessageVisibility`?

Difficulty:
Medium

`ChangeMessageVisibility`: modify visibility timeout for an in-flight message

Extend: give more time to process large/complex message
 Reduce (set to 0): make message immediately visible (abandon processing)

Parameters:

`ReceiptHandle`: token from `ReceiveMessage`
`VisibilityTimeout`: new timeout (0 to 12 hours)
 Setting to 0: returns message to queue immediately

Use cases:

Extend (more time needed):

```
svc.ChangeMessageVisibility(&sqs.ChangeMessageVisibilityInput{
    QueueUrl:      &queueURL,
    ReceiptHandle:  msg.ReceiptHandle,
    VisibilityTimeout: aws.Int64(300), // 5 more minutes
})
```

Abandon (unrecoverable error, don't wait for timeout):

```
svc.ChangeMessageVisibility(&sqs.ChangeMessageVisibilityInput{
    QueueUrl:      &queueURL,
    ReceiptHandle:  msg.ReceiptHandle,
    VisibilityTimeout: aws.Int64(0), // return to queue now
})
```

Heartbeat pattern (recommended for long-running tasks):

Every 30s: extend timeout by 60s

Ensures timeout doesn't expire during legitimate processing

Stop heartbeat when done → delete or abandon

Note: cannot reduce timeout to < time already elapsed

If 10s of original 30s remain: can extend, cannot set to < 0

Q25. ### Q40-Q50: More Visibility Timeout Topics

Difficulty:
Medium

Q	Topic
Q40	Batch ChangeMessageVisibility for multiple messages
Q41	Visibility timeout vs message retention period
Q42	In-flight message limit: what happens when reached
Q43	Consumer scaling and visibility timeout interaction
Q44	Optimizing visibility timeout for different workload types
Q45	Message age monitoring: ApproximateAgeOfOldestMessage
Q46	Poison message detection via ApproximateReceiveCount
Q47	Message timer vs queue delay: which to use
Q48	Concurrent consumers and visibility timeout coordination
Q49	SQS backoff strategy with visibility timeout
Q50	Consumer group patterns for SQS

Dead Letter Queues & Error Handling

Q26. What is a Dead Letter Queue (DLQ)?**Difficulty:**
Medium

DLQ: queue for messages that fail processing N times
After maxReceiveCount failures → automatically moved to DLQ
Main queue: processed normally; DLQ: problem messages

Configuration:

```
RedrivePolicy on main queue:
{
  "maxReceiveCount": 5,           // failures before DLQ
  "deadLetterTargetArn": "arn:aws:sqs:...:orders-dlq"
}
```

DLQ requirements:

Same type as source (Standard DLQ for Standard, FIFO DLQ for FIFO)
Same region and account as source queue
DLQ should have LONGER retention than source queue

Why messages end up in DLQ:

Processing exceptions (bug, invalid data)
Downstream service down (DB unavailable)
Malformed message (schema changed)
Timeout exceeded (slow processing)
Business rule violations

DLQ best practices:

Always configure DLQ on every queue (don't let messages expire silently)
DLQ retention = 14 days (max) → time to investigate
Alert on DLQ depth > 0 (any message in DLQ = bug/issue)
Process DLQ: fix bug → redrive (replay) to main queue

Monitoring:

CloudWatch alarm: ApproximateNumberOfMessagesVisible > 0 on DLQ → SNS alert

Q27. How do you redrive messages from DLQ?**Difficulty:**
Medium

Redrive: replay DLQ messages back to the source queue
After fixing the bug that caused failures

Methods:

1. AWS Console: "Start DLQ redrive" button
Redrive all or subset of messages
Rate: up to 500 messages/second
2. SQS Redrive API (AWS SDK):
StartMessageMoveTask: move messages from DLQ to source
ListMessageMoveTasks: check redrive progress
CancelMessageMoveTask: stop if needed
3. Manual redrive (custom code):
Poll DLQ → process differently (fix applied) → delete from DLQ
More control: can filter which messages to redrive

// Go: manual DLQ redrive

```

for {
    result, _ := svc.ReceiveMessage(&sqs.ReceiveMessageInput{
        QueueUrl:      &dlqURL,
        MaxNumberOfMessages: aws.Int64(10),
        WaitTimeSeconds:    aws.Int64(20),
    })
    for _, msg := range result.Messages {
        // Re-send to main queue with fix applied
        svc.SendMessage(&sqs.SendMessageInput{
            QueueUrl:    &mainQueueURL,
            MessageBody: msg.Body,
        })
        svc.DeleteMessage(&sqs.DeleteMessageInput{
            QueueUrl:      &dlqURL,
            ReceiptHandle: msg.ReceiptHandle,
        })
    }
    if len(result.Messages) == 0 { break }
}

```

Redrive strategy:

Don't just blindly redrive → messages will fail again!
 Fix the root cause first → verify fix → then redrive

Q28. What is the maxReceiveCount and how do you set it?

Difficulty:
Medium

maxReceiveCount: how many times a message can be received before going to DLQ
 Counts: how many times ReceiveMessage returned this message
 After count exceeded: moved to DLQ automatically

Setting the right value:

Transient failures (DB restart, network blip):
 maxReceiveCount = 3-5 (allow a few retries)

External API rate limits (429 errors):
 maxReceiveCount = 10 + use delay between retries

Never-succeed cases (malformed messages):
 maxReceiveCount = 3 (fail fast → DLQ quickly)

Interaction with visibility timeout:

Each ReceiveMessage (and subsequent timeout) = +1 to ApproximateReceiveCount
 If consumer receives message but crashes before delete:
 Timeout expires → message visible → next receive = count+1

maxReceiveCount = 5, visibility = 60s:
 Message retried 5 times over 5+ minutes → DLQ

```

// Check receive count in your consumer:
for _, msg := range result.Messages {
    receiveCount, _ := strconv.Atoi(
        aws.StringValue(msg.Attributes["ApproximateReceiveCount"]),
    )
    if receiveCount > 3 {
        // Close to DLQ, log extra details for debugging
        log.Warnf("message %s received %d times", *msg.MessageId, receiveCount)
    }
}

```

```

    }
    processMessage(msg)
}

```

Q29. How do you handle different types of failures?

Difficulty: Hard

Failure categories and handling:

1. Transient failures (retry-able):

Network timeout, DB connection failure, temporary unavailability
 → Don't delete, let visibility timeout expire → auto-retry
 → Or: `ChangeMessageVisibility(0)` to retry immediately
 → `maxReceiveCount = 5`

2. Business logic errors (non-retryable):

Validation failure, missing required data, invariant violation
 → Move to DLQ immediately (don't waste retries)
 → Or: send to separate error queue for human review

```

if isValidatationError(err) {
    moveToErrorQueue(msg)
    svc.DeleteMessage(...) // remove from main queue
    return
}

```

3. Poison pills (malformed messages):

Can never be processed regardless of retries
 → `maxReceiveCount = 1` or `2` → fast path to DLQ
 → DLQ message with original payload for debugging

4. External API rate limits:

429 Too Many Requests
 → `ChangeMessageVisibility` to delay retry
 → Increase visibility timeout to implement backoff

5. Idempotency violations:

Message already processed (`ApproximateReceiveCount > 1`)
 Check idempotency key → if already done → delete + continue
 No failure, just duplicate detection

// Categorized error handling:

```

func handleSQSMessage(msg *sqs.Message) {
    err := processMessage(msg)
    if err == nil { deleteMessage(msg); return }

    if isNonRetryable(err) {
        sendToErrorQueue(msg, err)
        deleteMessage(msg)
    }
    // Otherwise: don't delete → let SQS retry (visibility timeout)
}

```

Q30. ### Q55-Q65: More DLQ Topics

Difficulty:
Medium

Q	Topic
---	-------

Q55	DLQ alarm configuration in CloudWatch
Q56	Multiple DLQs: per-error-type routing
Q57	DLQ message metadata for debugging (original timestamp, error)
Q58	DLQ retention period best practices
Q59	Redrive policy limits and considerations
Q60	DLQ with Lambda: partial batch failure handling
Q61	Circuit breaker pattern with SQS and DLQ
Q62	Error enrichment: adding context before DLQ
Q63	DLQ consumer for alerts and automated recovery
Q64	Message expiry: what happens when retention exceeded
Q65	Testing DLQ behavior in development

SNS + SQS Fan-Out Patterns

Q31. What is the SNS + SQS fan-out pattern?

Difficulty:
Medium

Fan-out: one event → multiple independent consumers

SNS Topic: broadcast to all subscribers

SQS Queues: one per consumer service (buffers for each)

Architecture:

OrderService → SNS:order-created → [SQS:email-queue, SQS:inventory-queue, SQS:analytics-queue]

EmailService → polls SQS:email-queue → sends confirmation email

InventoryService → polls SQS:inventory-queue → reserves stock

AnalyticsService → polls SQS:analytics-queue → updates dashboards

Benefits:

Decoupled: services don't know about each other

Independent scaling: each service scales based on its own queue depth

Independent failures: email service down → inventory unaffected

Independent retry: each service retries independently

Add new subscriber: create new SQS queue + SNS subscription (no code change)

vs Direct SNS → Lambda:

With SQS buffer: handles bursts (SQS absorbs spike, Lambda processes at own rate)

Without SQS: Lambda throttled if SNS bursts faster than Lambda concurrency

Setup:

1. Create SNS topic
2. Create SQS queue per consumer
3. Subscribe each SQS queue to SNS topic
4. SQS queue policy: allow SNS to send
5. SNS subscription: filter policy (optional, per-service filtering)

Q32. What is SNS message filtering?**Difficulty:**
Medium

Subscription Filter Policy: each SQS subscriber receives only matching messages

Without filtering: every subscriber receives every SNS message

With filtering: subscribers receive only relevant messages

Filter policy (JSON on the subscription):

```
{
  "eventType": ["order.created", "order.cancelled"],
  "region": ["us-east-1", "eu-west-1"]
}
```

Message attributes must match filter for delivery:

SNS message: {MessageAttributes: {eventType: "order.shipped"}}

Filter: eventType: ["order.created"] → NOT delivered to this subscriber

Filter: eventType: ["order.created", "order.shipped"] → DELIVERED

Filter operators:

```
Exact match: {"eventType": ["order.created"]}
Prefix:      {"eventType": [{"prefix": "order."}]}
Anything-but: {"eventType": [{"anything-but": ["payment.failed"]}]}
Numeric:     {"amount": [{"numeric": [">=", 100]}]}
Exists:      {"eventType": [{"exists": true}]}
```

Use cases:

EmailService: subscribe to order.created only

FraudService: subscribe to payment.* events

HighValueAlerts: subscribe to orders with amount >= \$1000

RegionalService: subscribe to events for specific region

Benefits:

Reduced queue depth (services only process relevant events)

Lower cost (fewer messages delivered to non-interested services)

Simplified consumer logic (no need to filter in code)

Q33. What is EventBridge vs SNS+SQS?**Difficulty: Hard**

	SNS+SQS	EventBridge
Routing	Topic-based (all subs)	Content-based (rules)
Filtering	Attribute filtering	Rich JSON pattern matching
Sources	Your apps + AWS services	100+ AWS services natively
Targets	SQS, Lambda, HTTP, email	20+ targets (SQS, Lambda, Step Functions, ...)
Schema	No schema enforcement	Schema registry available
Replay	No	Event archive + replay
Cost	Per message	Per event + rules
Setup	Lower	Higher (more powerful)

EventBridge advantages:

Richer event routing (any JSON field, nested paths)

Native AWS service events (EC2 start, S3 put, CodePipeline, etc.)

Event archive + replay for debugging/reprocessing

Schema registry for event discovery

Cross-account event routing

SNS+SQS advantages:

Simpler for basic fan-out
 Lower latency (SNS → SQS is faster than EventBridge)
 Message buffering per subscriber (SQS)
 Lower cost for simple routing

Decision:

EventBridge: complex routing, AWS service events, event sourcing, replay
 SNS+SQS: simple fan-out to services, buffered delivery, basic filtering
 Both: can coexist (EventBridge → SNS → SQS for complex + buffered)

Q34. ### Q69-Q75: More SNS+SQS Topics

Difficulty:
Medium

Q	Topic
Q69	SNS FIFO topic + SQS FIFO queue for ordered fan-out
Q70	Cross-account SNS → SQS: IAM and queue policy setup
Q71	SNS message structure and how it wraps SQS messages
Q72	HTTP/HTTPS SNS subscription with signature verification
Q73	SNS delivery retry policy and DLQ for SNS
Q74	SQS + EventBridge Pipes for transformation
Q75	Outbox pattern: DB → SQS vs DB → SNS → SQS

SQS vs Kafka & Architecture Decisions

Q35. When do you choose SQS over Kafka?

Difficulty:
Medium

Choose SQS when:

1. Simplicity + managed service:
 No Kafka cluster to manage, patch, or scale
 AWS handles all operations
2. Simple task queue pattern:
 Produce task → single consumer processes it → delete
 No need for replay or multiple consumers
3. AWS ecosystem integration:
 Lambda triggers, EventBridge, SNS — native integration
 IAM-based auth (no Kafka credential management)
4. Variable/unpredictable workload:
 SQS scales to zero (no idle broker cost)
 Pay-per-message (not per-hour for cluster)
5. Small/medium volume (< 1M msg/day):

SQS cost: nearly free at low volume
Kafka on EC2: \$50-200/month minimum

Choose Kafka when:

1. Message replay needed:
Audit, reprocessing, event sourcing
SQS: messages deleted after consumption
2. Multiple independent consumers:
Both OrderService AND AnalyticsService read same event
SQS: needs separate queue per consumer (duplication)
Kafka: multiple consumer groups, one topic
3. High throughput (>100K msg/sec):
SQS Standard: ~100K TPS per queue (works)
But Kafka: lower cost and better performance at this scale
4. Strict ordering with high throughput:
SQS FIFO: 3,000 TPS max
Kafka: millions/sec with ordering per partition
5. Stream processing:
Kafka Streams, Flink, Spark Streaming native Kafka integration
SQS: no native stream processing

Q36. What is the SQS vs RabbitMQ comparison?

Difficulty:
Medium

	SQS	RabbitMQ
Management	Fully managed (AWS)	Self-hosted or CloudAMQP
Cost	Per-message pricing	Infrastructure cost
Ops burden	Zero	High (clusters, upgrades)
Message routing	Simple (one queue)	Rich (exchanges, bindings)
Message TTL	Per-queue	Per-message
Priority	No	Yes (x-max-priority)
Replay	No	No (same as SQS)
Protocol	HTTP/AWS SDK	AMQP (richer)
Throughput	Unlimited (Standard)	~100K/sec (single cluster)
Max message	256KB	Default unlimited
Ordering	FIFO queue	Single consumer per queue

Choose SQS over RabbitMQ:

AWS-native deployment (no extra infra)
Don't want to manage message broker
Simple queue patterns without complex routing
Variable workload (pay only for what you use)

Choose RabbitMQ over SQS:

Complex routing: topic exchange, dead letter, priority queues
AMQP protocol required
Non-AWS environment
Per-message TTL needed
Already have RabbitMQ expertise

Q37. What are SQS architecture anti-patterns?

Difficulty: Hard

Anti-pattern 1: Using SQS as a database

Storing state in SQS messages (polling for state)

SQS is a queue, not storage

Fix: use DynamoDB, RDS, or Redis for state

Anti-pattern 2: Single queue for all message types

All events in one queue → consumers must filter in code

Fix: separate queues by type OR SNS → multiple SQS queues

Anti-pattern 3: No DLQ configured

Failed messages silently expire after retention period

Fix: always configure DLQ + CloudWatch alarm on DLQ depth

Anti-pattern 4: Tiny visibility timeout

Consumer needs 2 min, timeout is 30s

→ Duplicate processing on every message

Fix: set visibility timeout >> max processing time

Anti-pattern 5: No idempotency with Standard queue

Processing payment twice = overcharge

Fix: idempotency check before processing (DynamoDB conditional write)

Anti-pattern 6: Processing in ReceiveMessage loop without batching

1 message per poll → 1 API call per message → high cost

Fix: MaxNumberOfMessages=10, batch process

Anti-pattern 7: Synchronous caller waiting on SQS

Client sends message, then polls for result

Fix: for sync: use API call directly; for async results: use callback/webhook

Or: use SQS Request-Reply with correlation ID + separate response queue

Q38. ### Q79-Q85: More Architecture Topics**Difficulty:**
Medium

Q	Topic
Q79	SQS as a buffer for downstream rate limiting
Q80	SQS in event-driven microservices architecture
Q81	Saga pattern with SQS for distributed transactions
Q82	SQS fan-out vs Kinesis fan-out: when to choose each
Q83	SQS with Step Functions for workflow orchestration
Q84	SQS for scheduled task execution with EventBridge
Q85	SQS reliability: what AWS SLA guarantees

Lambda Triggers, Go SDK & Production Patterns

Q39. How does Lambda + SQS trigger work?**Difficulty:**
Medium

Lambda Event Source Mapping (ESM):

- Lambda polls SQS on your behalf
- Invokes Lambda function with batch of messages
- Lambda handles deletion on success

Configuration:

- BatchSize: 1-10,000 messages per invocation
- MaximumBatchingWindowInSeconds: wait up to N seconds to fill batch
- FunctionResponseType: ReportBatchItemFailures (partial success)

Lambda invocation:

- ESM polls queue → batch ready → invoke Lambda with event

```
{
  "Records": [
    {
      "messageId": "abc-123",
      "receiptHandle": "...",
      "body": "{\"orderId\": 42}",
      "attributes": {"ApproximateReceiveCount": "1"},
      "messageAttributes": {"eventType": {"stringValue": "order.created"}}
    }
  ]
}
```

Success: Lambda returns (no error) → ESM deletes ALL messages in batch

Failure: Lambda throws error → ESM makes messages visible → retry
After maxReceiveCount: → DLQ

Partial batch failure (best practice):

- Return {batchItemFailures: [{itemIdentifier: messageId}]}
- Only failed messages are retried; successful ones deleted
- Without this: ALL messages retried even if only 1 failed

FIFO + Lambda:

- One concurrent invocation per message group
- Within group: strict ordering maintained
- Across groups: parallel Lambda invocations

Q40. How do you implement SQS consumer in Go?**Difficulty:**
Medium

```
package main

import (
    "context"
    "encoding/json"
    "log"
    "sync"

    "github.com/aws/aws-sdk-go-v2/aws"
    "github.com/aws/aws-sdk-go-v2/config"
    "github.com/aws/aws-sdk-go-v2/service/sqs"
    "github.com/aws/aws-sdk-go-v2/service/sqs/types"
}
```

```

)

type Consumer struct {
    client    *sqs.Client
    queueURL string
    workers  int
}

func (c *Consumer) Run(ctx context.Context) {
    jobs := make(chan types.Message, c.workers*2)
    var wg sync.WaitGroup

    // Start worker pool
    for i := 0; i < c.workers; i++ {
        wg.Add(1)
        go func() {
            defer wg.Done()
            for msg := range jobs {
                c.process(ctx, msg)
            }
        }()
    }

    // Poll loop
    for {
        select {
        case <-ctx.Done():
            close(jobs)
            wg.Wait()
            return
        default:
            result, err := c.client.ReceiveMessage(ctx, &sqs.ReceiveMessageInput{
                QueueUrl:      aws.String(c.queueURL),
                MaxNumberOfMessages: 10,
                WaitTimeSeconds: 20, // long polling
                AttributeNames: []types.QueueAttributeName{"ApproximateReceiveCount"},
                MessageAttributeNames: []string{"All"},
            })
            if err != nil { log.Printf("receive error: %v", err); continue }
            for _, msg := range result.Messages {
                jobs <- msg
            }
        }
    }
}

func (c *Consumer) process(ctx context.Context, msg types.Message) {
    var payload OrderPayload
    if err := json.Unmarshal([]byte(aws.ToString(msg.Body)), &payload); err != nil {
        log.Printf("unmarshal error for %s: %v", aws.ToString(msg.MessageId), err)
        return // don't delete → goes to DLQ after maxReceiveCount
    }
    if err := processOrder(ctx, payload); err != nil {
        log.Printf("process error for %s: %v", aws.ToString(msg.MessageId), err)
        return // don't delete → retry
    }
    c.client.DeleteMessage(ctx, &sqs.DeleteMessageInput{

```

```

        QueueUrl:      aws.String(c.queueURL),
        ReceiptHandle: msg.ReceiptHandle,
    })
}

```

Q41. How do you implement SQS producer in Go?

Difficulty: Easy

```

import (
    "github.com/aws/aws-sdk-go-v2/aws"
    "github.com/aws/aws-sdk-go-v2/config"
    "github.com/aws/aws-sdk-go-v2/service/sqs"
    "github.com/aws/aws-sdk-go-v2/service/sqs/types"
)

type Producer struct {
    client    *sqs.Client
    queueURL string
}

func NewProducer(ctx context.Context, queueURL string) (*Producer, error) {
    cfg, err := config.LoadDefaultConfig(ctx)
    if err != nil { return nil, err }
    return &Producer{
        client:    sqs.NewFromConfig(cfg),
        queueURL: queueURL,
    }, nil
}

func (p *Producer) Send(ctx context.Context, payload interface{}, eventType string) error {
    body, err := json.Marshal(payload)
    if err != nil { return fmt.Errorf("marshal: %w", err) }

    _, err = p.client.SendMessage(ctx, &sqs.SendMessageInput{
        QueueUrl:      aws.String(p.queueURL),
        MessageBody:   aws.String(string(body)),
        MessageAttributes: map[string]types.MessageAttributeValue{
            "EventType": {
                DataType:   aws.String("String"),
                StringValue: aws.String(eventType),
            },
            "TraceId": {
                DataType:   aws.String("String"),
                StringValue: aws.String(traceIDFromContext(ctx)),
            },
        },
    })
    return err
}

// Batch send (cost-efficient):
func (p *Producer) SendBatch(ctx context.Context, messages []BatchMessage) error {
    entries := make([]types.SendMessageBatchRequestEntry, len(messages))
    for i, m := range messages {
        body, _ := json.Marshal(m.Payload)
        entries[i] = types.SendMessageBatchRequestEntry{
            Id:          aws.String(fmt.Sprintf("msg-%d", i)),
            MessageBody: aws.String(string(body)),
        }
    }
    _, err := p.client.SendMessageBatch(ctx, &sqs.SendMessageBatchInput{
        QueueUrl:      aws.String(p.queueURL),
        Entries:        entries,
    })
    return err
}

```

```

    }
}
result, err := p.client.SendMessageBatch(ctx, &sqs.SendMessageBatchInput{
    QueueUrl: aws.String(p.queueURL),
    Entries:  entries,
})
if len(result.Failed) > 0 {
    return fmt.Errorf("%d messages failed to send", len(result.Failed))
}
return err
}

```

Q42. How do you implement Lambda SQS handler in Go?

Difficulty:
Medium

```

import (
    "github.com/aws/aws-lambda-go/events"
    "github.com/aws/aws-lambda-go/lambda"
)

type SQSHandler struct {
    processor OrderProcessor
}

// Handle batch with partial failure support
func (h *SQSHandler) Handle(ctx context.Context, event events.SQSEvent) (events.SQSEventResponse, error) {
    var failures []events.SQSBatchItemFailure

    for _, record := range event.Records {
        if err := h.processRecord(ctx, record); err != nil {
            log.Printf("failed to process message %s: %v", record.MessageId, err)
            failures = append(failures, events.SQSBatchItemFailure{
                ItemIdentifier: record.MessageId,
            })
        }
    }

    // Return partial failures: only failed messages are retried
    // Successful messages are automatically deleted by Lambda ESM
    return events.SQSEventResponse{
        BatchItemFailures: failures,
    }, nil
}

func (h *SQSHandler) processRecord(ctx context.Context, record events.SQSMessage) error {
    var order OrderPayload
    if err := json.Unmarshal([]byte(record.Body), &order); err != nil {
        // Non-retryable: return nil to delete (bad message) or error to retry
        log.Printf("malformed message %s, skipping: %v", record.MessageId, err)
        return nil // delete it; don't retry unparseable messages
    }

    // Check idempotency
    if processed, _ := h.processor.IsAlreadyProcessed(ctx, record.MessageId); processed {
        return nil // already done, delete
    }

    return h.processor.Process(ctx, order, record.MessageId)
}

```

```

}

func main() {
    handler := &SQSHandler{processor: NewOrderProcessor()}
    lambda.Start(handler.Handle)
}

```

Q43. How do you handle idempotency with SQS in Go?

Difficulty: Hard

```

// DynamoDB-based idempotency store
import "github.com/aws/aws-sdk-go-v2/service/dynamodb"

type IdempotencyStore struct {
    dynamo      *dynamodb.Client
    tableName    string
    ttl          time.Duration
}

// Returns (alreadyProcessed, error)
func (s *IdempotencyStore) CheckAndSet(ctx context.Context, messageID string) (bool, error) {
    now := time.Now()
    expiry := now.Add(s.ttl).Unix()

    _, err := s.dynamo.PutItem(ctx, &dynamodb.PutItemInput{
        TableName: aws.String(s.tableName),
        Item: map[string]types.AttributeValue{
            "message_id": &types.AttributeValueMemberS{Value: messageID},
            "processed_at": &types.AttributeValueMemberN{Value: strconv.FormatInt(now.Unix(), 10)},
            "ttl": &types.AttributeValueMemberN{Value: strconv.FormatInt(expiry, 10)},
        },
        ConditionExpression: aws.String("attribute_not_exists(message_id)" ),
    })

    if err != nil {
        var ccfc *types.ConditionalCheckFailedException
        if errors.As(err, &ccfc) {
            return true, nil // already processed
        }
        return false, err
    }
    return false, nil // first time processing
}

// Usage in consumer:
func processMessage(ctx context.Context, msg types.Message) error {
    alreadyDone, err := idempotency.CheckAndSet(ctx, *msg.MessageId)
    if err != nil { return err }
    if alreadyDone {
        log.Printf("message %s already processed, skipping", *msg.MessageId)
        return nil
    }
    return doActualWork(ctx, msg)
}

```

Q44. How do you implement a visibility timeout heartbeat in Go?

Difficulty: Hard

```

func (c *Consumer) processWithHeartbeat(ctx context.Context, msg types.Message) error {
    ctx, cancel := context.WithCancel(ctx)
    defer cancel()

    // Start heartbeat goroutine
    heartbeatDone := make(chan struct{})
    go func() {
        defer close(heartbeatDone)
        ticker := time.NewTicker(25 * time.Second) // extend every 25s
        defer ticker.Stop()
        for {
            select {
            case <-ctx.Done():
                return
            case <-ticker.C:
                _, err := c.client.ChangeMessageVisibility(context.Background(),
                    &sqs.ChangeMessageVisibilityInput{
                        QueueUrl:      aws.String(c.queueURL),
                        ReceiptHandle:    msg.ReceiptHandle,
                        VisibilityTimeout: aws.Int32(60), // extend by 60s
                    })
                if err != nil {
                    log.Printf("heartbeat failed for %s: %v", aws.ToString(msg.MessageId), err)
                }
            }
        }
    }()

    // Do actual work
    err := c.doWork(ctx, msg)
    cancel() // stop heartbeat
    <-heartbeatDone // wait for heartbeat goroutine to exit
    return err
}

```

Q45. How do you implement SQS with exponential backoff?

Difficulty:
Medium

```

// Instead of relying solely on SQS retry (fixed delay):
// Implement explicit backoff by using ChangeMessageVisibility

```

```

func (c *Consumer) processWithBackoff(ctx context.Context, msg types.Message) {
    receiveCount := getReceiveCount(msg)

    err := processMessage(ctx, msg)
    if err == nil {
        // Success: delete message
        c.client.DeleteMessage(ctx, &sqs.DeleteMessageInput{
            QueueUrl:      aws.String(c.queueURL),
            ReceiptHandle: msg.ReceiptHandle,
        })
        return
    }

    if !isRetryable(err) {
        // Non-retryable: delete and send to DLQ manually
        sendToDLQ(ctx, msg, err)
    }
}

```

```

        c.client.DeleteMessage(ctx, &sqs.DeleteMessageInput{
            QueueUrl:      aws.String(c.queueURL),
            ReceiptHandle: msg.ReceiptHandle,
        })
        return
    }

    // Retryable: set visibility timeout based on attempt number
    backoffSeconds := int32(math.Min(
        30*math.Pow(2, float64(receiveCount-1)), // 30, 60, 120, 240, 300
        300, // max 5 minutes (SQS limit 900s for message delay)
    ))
    c.client.ChangeMessageVisibility(ctx, &sqs.ChangeMessageVisibilityInput{
        QueueUrl:      aws.String(c.queueURL),
        ReceiptHandle: msg.ReceiptHandle,
        VisibilityTimeout: aws.Int32(backoffSeconds),
    })
}

func getReceiveCount(msg types.Message) int {
    if v, ok := msg.Attributes["ApproximateReceiveCount"]; ok {
        n, _ := strconv.Atoi(v)
        return n
    }
    return 1
}

```

Q46. How do you test SQS code in Go?

Difficulty:
Medium

```

// Option 1: Interface + Mock
type SQSClient interface {
    SendMessage(ctx context.Context, params *sqs.SendMessageInput, ...) (*sqs.SendMessageOutput, error)
    ReceiveMessage(ctx context.Context, params *sqs.ReceiveMessageInput, ...) (*sqs.ReceiveMessageOutput, error)
    DeleteMessage(ctx context.Context, params *sqs.DeleteMessageInput, ...) (*sqs.DeleteMessageOutput, error)
}

type MockSQSClient struct {
    messages []types.Message
    sent      []string
}

func (m *MockSQSClient) ReceiveMessage(ctx context.Context, _ *sqs.ReceiveMessageInput, ...) (*sqs.ReceiveMessageOutput, error) {
    if len(m.messages) == 0 {
        return &sqs.ReceiveMessageOutput{}, nil
    }
    msg := m.messages[0]
    m.messages = m.messages[1:]
    return &sqs.ReceiveMessageOutput{Messages: []types.Message{msg}}, nil
}

// Option 2: LocalStack (real SQS API locally)
// docker run -p 4566:4566 localstack/localstack
// Set endpoint: http://localhost:4566

cfg, _ := config.LoadDefaultConfig(ctx,
    config.WithEndpointResolverWithOptions(

```

```

    aws.EndpointResolverWithOptionsFunc(func(service, region string, options ...interface{}) (aws.Endpoint,
        return aws.Endpoint{URL: "http://localhost:4566"}, nil
    })),
),
config.WithCredentialsProvider(credentials.NewStaticCredentialsProvider("test", "test", "")),
)

// Option 3: testcontainers-go with LocalStack
container, _ := localstack.RunContainer(ctx, testcontainers.WithImage("localstack/localstack"))
endpoint, _ := container.SQSURL(ctx)

```

Q47. How do you monitor SQS queues in production?

Difficulty:
Medium

```

// CloudWatch alarms via Go (setup during deploy):
import "github.com/aws/aws-sdk-go-v2/service/cloudwatch"

func setupQueueAlarms(ctx context.Context, cw *cloudwatch.Client, queueName, dlqName string) error {
    // Alarm: queue depth growing (consumers lagging)
    _, err := cw.PutMetricAlarm(ctx, &cloudwatch.PutMetricAlarmInput{
        AlarmName:      aws.String(queueName + "-depth-alarm"),
        MetricName:      aws.String("ApproximateNumberOfMessagesVisible"),
        Namespace:      aws.String("AWS/SQS"),
        Statistic:      types.StatisticMaximum,
        Period:         aws.Int32(300), // 5 minutes
        EvaluationPeriods: aws.Int32(2),
        Threshold:       aws.Float64(10000),
        ComparisonOperator: types.ComparisonOperatorGreaterThanThreshold,
        Dimensions: []types.Dimension{{
            Name:  aws.String("QueueName"),
            Value: aws.String(queueName),
        }},
        AlarmActions: []string{"arn:aws:sns:...:alerts"},
    })
    if err != nil { return err }

    // Alarm: any message in DLQ
    _, err = cw.PutMetricAlarm(ctx, &cloudwatch.PutMetricAlarmInput{
        AlarmName:      aws.String(dlqName + "-not-empty"),
        MetricName:      aws.String("ApproximateNumberOfMessagesVisible"),
        Namespace:      aws.String("AWS/SQS"),
        Statistic:      types.StatisticSum,
        Period:         aws.Int32(60),
        EvaluationPeriods: aws.Int32(1),
        Threshold:       aws.Float64(0),
        ComparisonOperator: types.ComparisonOperatorGreaterThanThreshold,
        Dimensions: []types.Dimension{{
            Name:  aws.String("QueueName"),
            Value: aws.String(dlqName),
        }},
        AlarmActions: []string{"arn:aws:sns:...:alerts"},
    })
    return err
}

```


Q48. What is SQS cost optimization?**Difficulty:**
Medium

Cost reduction strategies:

1. Long polling (WaitTimeSeconds=20):
Fewer empty receives → fewer API calls
Example: 1 msg/min queue
Short polling: $60 \times \$0.40/M = \sim \0 but wasteful + CPU
Long polling: $3 \times \$0.40/M$ (20s each) = 5× fewer calls
2. Batch operations:
SendMessageBatch: 10 msgs = 1 API call (not 10)
DeleteMessageBatch: 10 msgs = 1 API call
10× cost reduction for high-volume queues
3. Increase batch size in Lambda ESM:
BatchSize=10 → 1 Lambda invocation per 10 messages
BatchSize=1 → 10 Lambda invocations (10× more expensive)
4. Message consolidation:
Combine related small messages into one batch message
{"events": [event1, event2, event3]} = 1 message = 1 API call vs 3
5. Right-size polling interval:
Low-volume queues: long polling + increase polling interval
High-volume queues: always polling (long polling = 20s max wait anyway)
6. Standard over FIFO when possible:
FIFO: 25% more expensive
If ordering not required: use Standard
7. Filter at SNS level (subscription filter policy):
Services receive only relevant messages
Fewer messages → fewer SQS API calls

Pricing example (1M messages/day, 10 msg/batch):
Without batch: 1M API calls = \$0.40/day
With batch: 100K API calls = \$0.04/day (10× cheaper)

Q49. ### Q96-Q100: Final Advanced Topics**Difficulty:**
Medium**Q50. What is SQS resource tagging and cost allocation?****Difficulty: Easy**

```
// Tag queues for cost allocation
svc.TagQueue(ctx, &sqs.TagQueueInput{
    QueueUrl: aws.String(queueURL),
    Tags: map[string]string{
        "Environment": "production",
        "Service":     "order-service",
        "Team":        "platform",
        "CostCenter":  "engineering",
    },
})
// List tags
```

```

result, _ := svc.ListQueueTags(ctx, &sqs.ListQueueTagsInput{
    QueueUrl: aws.String(queueURL),
})

// AWS Cost Explorer: filter by tag
// Service: SQS, Tag: Service=order-service
// → shows SQS cost attributed to order service

```

Q51. What is SQS cross-region and cross-account patterns?

Difficulty: Hard

Cross-account SQS:

Queue in Account A, producer in Account B

Queue policy (Account A): allow Account B to SendMessage

```

{
  "Effect": "Allow",
  "Principal": {"AWS": "arn:aws:iam::ACCOUNT-B:root"},
  "Action": "sqs:SendMessage",
  "Resource": "arn:aws:sqs:us-east-1:ACCOUNT-A:my-queue"
}

```

Producer (Account B): use queue ARN from Account A

Assumes role with sqs:SendMessage or directly with queue policy

Cross-region SQS:

SQS queues are region-specific

No built-in cross-region replication

Pattern: SNS cross-region subscription

SNS Topic (us-east-1) → SNS subscription → SQS (eu-west-1)

Via HTTP subscription or SQS-to-SQS (not native)

Or: application-level dual-write to queues in multiple regions

EventBridge: native cross-region event routing

EventBridge rule (us-east-1) → EventBridge bus (eu-west-1) → SQS (eu-west-1)

Typical use case:

Multi-region active-active: each region has its own SQS queues

Global events: use EventBridge for cross-region routing to SQS

Q52. What is SQS with AWS Step Functions?

Difficulty: Medium

Step Functions + SQS: orchestrate long-running workflows with SQS as transport

Pattern 1: SQS as trigger for Step Functions

SQS message → Lambda → starts Step Functions execution

Pattern 2: Step Functions with SQS task (Wait for callback)

Step Functions starts → sends message to SQS with taskToken

Consumer processes → calls SendTaskSuccess(taskToken) → Step Functions continues

```

{
  "Type": "Task",
  "Resource": "arn:aws:states:::sqs:sendMessage.waitForTaskToken",
  "Parameters": {

```

```

        "QueueUrl": "https://sqs...",
        "MessageBody": {
            "taskToken.$": "$$.Task.Token",
            "input.$": "$"
        }
    }
}

// Go consumer: complete Step Functions task
svc.SendTaskSuccess(ctx, &sfn.SendTaskSuccessInput{
    TaskToken: aws.String(taskToken),
    Output:    aws.String(`{"result": "success"}`),
})
// or failure:
svc.SendTaskFailure(ctx, &sfn.SendTaskFailureInput{
    TaskToken: aws.String(taskToken),
    Error:     aws.String("ProcessingFailed"),
    Cause:     aws.String(err.Error()),
})

```

Use case: human approval workflows, external system integration,
long-running async processing with correlation

Q53. What is the SQS outbox pattern implementation?

Difficulty: Hard

```

// Transactional outbox: write to DB + SQS atomically using DB outbox table

// Step 1: Write to outbox table in same DB transaction
func (s *OrderService) CreateOrder(ctx context.Context, order *Order) error {
    return s.db.WithTx(ctx, func(tx *sql.Tx) error {
        // Business logic
        if err := insertOrder(ctx, tx, order); err != nil { return err }

        // Write to outbox (same transaction = atomic!)
        eventPayload, _ := json.Marshal(OrderCreatedEvent{
            OrderID:    order.ID,
            UserID:     order.UserID,
            Amount:     order.Amount,
            CreatedAt:  order.CreatedAt,
        })
        _, err := tx.ExecContext(ctx, `
            INSERT INTO outbox (id, event_type, payload, created_at)
            VALUES ($1, $2, $3, NOW())`,
            uuid.New().String(), "order.created", eventPayload,
        )
        return err
    })
}

// Step 2: Outbox relay (background goroutine or separate service)
func (r *OutboxRelay) Run(ctx context.Context) {
    for {
        select {
        case <-ctx.Done(): return
        case <-time.After(100 * time.Millisecond):
            r.publishPending(ctx)
        }
    }
}

```

```

    }
}

func (r *OutboxRelay) publishPending(ctx context.Context) {
    rows, _ := r.db.QueryContext(ctx, `
        SELECT id, event_type, payload FROM outbox
        WHERE published_at IS NULL ORDER BY created_at LIMIT 100
        FOR UPDATE SKIP LOCKED`)
    defer rows.Close()

    var entries []OutboxEntry
    for rows.Next() {
        var e OutboxEntry
        rows.Scan(&e.ID, &e.EventType, &e.Payload)
        entries = append(entries, e)
    }
    if len(entries) == 0 { return }

    // Send to SQS in batches of 10
    for i := 0; i < len(entries); i += 10 {
        batch := entries[i:min(i+10, len(entries))]
        r.sendBatch(ctx, batch)
    }
}

```

Q54. What are SQS production checklist items?

Difficulty:
Medium

Production readiness checklist:

Queue Configuration:

- DLQ configured on every queue
- maxReceiveCount set appropriately (3-5 for most)
- DLQ retention = 14 days
- Visibility timeout = max_processing_time × 6
- Long polling enabled (WaitTimeSeconds=20)
- SSE enabled for sensitive data
- Appropriate access policy (least privilege IAM)

Consumer:

- Idempotency implemented (Standard queues)
- Batch processing (MaxNumberOfMessages=10)
- Visibility timeout heartbeat for long-running tasks
- Graceful shutdown: drain in-flight on SIGTERM
- Structured error handling (retryable vs non-retryable)
- Partial batch failure for Lambda (ReportBatchItemFailures)

Monitoring:

- CloudWatch alarm: queue depth growing (consumers lagging)
- CloudWatch alarm: DLQ not empty (any message = bug)
- CloudWatch alarm: ApproximateAgeOfOldestMessage > threshold
- Distributed tracing: trace IDs in message attributes
- Structured logging: message IDs, receive counts, processing time

Operations:

- DLQ redrive procedure documented and tested
- Queue purge procedure for emergency data flush

- Message replay playbook
- Runbook: what to do when DLQ alarm fires
- Load testing: verify throughput meets requirements

Cost:

- Batch operations used (10× cost reduction)
- Resource tags for cost allocation
- Queue type right-sized (Standard vs FIFO)

Master these 100 questions and you'll handle any AWS SQS interview at SDE2 level. Key areas: Standard vs FIFO (when to use each), visibility timeout and message lifecycle, DLQ configuration and redrive, SNS+SQS fan-out patterns, SQS vs Kafka decision framework, Lambda partial batch failures, and Go AWS SDK patterns (consumer, producer, idempotency, backoff). ■

Additional SQS Questions (Q55–Q100)

Q55. What is SQS FIFO exactly-once processing?

Difficulty: Hard

SQS FIFO Exactly-Once Processing:

MessageDeduplicationId: prevents duplicate messages within 5-minute window
 MessageGroupId: enables per-group ordering

Deduplication methods:

1. Content-based: SHA-256 of message body (set ContentBasedDeduplication=true)
2. Explicit: you provide MessageDeduplicationId

Example:

MessageDeduplicationId = "order-42-payment"
 If same ID sent within 5 minutes → second send silently discarded

FIFO limitation: 300 TPS (3000 with batching)

Standard: 300,000+ TPS (nearly unlimited)

If you need high throughput + ordering:

Use multiple FIFO queues (shard by customer/entity)
 Or use Kafka (no TPS limit per partition)

Q56. What is SQS message attributes?

Difficulty: Medium

Message attributes: metadata alongside message body

Types: String, Number, Binary, String.Array

```
sqs.send_message(
    QueueUrl=queue_url,
    MessageBody='{"orderId": "42"}',
    MessageAttributes={
        'EventType': {
            'StringValue': 'order.created',
            'DataType': 'String'
        },
        'Version': {
            'StringValue': '1',
            'DataType': 'Number'
        }
    })
```

```

        },
        'Timestamp': {
            'StringValue': '2024-01-15T10:00:00Z',
            'DataType': 'String'
        },
        'TraceId': {
            'StringValue': 'abc-123-def',
            'DataType': 'String'
        }
    }
}
)

# Receive with attributes
response = sqs.receive_message(
    QueueUrl=queue_url,
    MessageAttributeNames=['All'], # or specific: ['EventType']
    AttributeNames=['All']        # system attributes (SentTimestamp, etc.)
)

# System attributes available:
# SentTimestamp, ApproximateFirstReceiveTimestamp
# ApproximateReceiveCount, MessageDeduplicationId (FIFO)

# Use for: routing without parsing body, filtering, tracing

```

Q57. What is SQS with Lambda event source mapping?

Difficulty: Hard

SQS → Lambda integration:

- Lambda polls SQS (long-polling, 5 parallel pollers per shard)
- Lambda invoked with batch of messages (up to 10,000 for Standard)
- Lambda scales: 1 concurrent execution per queue (initial)
 - scales up to 1000 concurrent (Standard)
 - 1 concurrent per MessageGroupId (FIFO)

Failure behavior (critical):

- If Lambda throws: entire batch returned to queue and retried
- "Batch Item Failures": report partial failures
 - only failed messages go back to queue
 - successfully processed messages NOT re-sent

Lambda function (batch item failure reporting):

```

def handler(event, context):
    batch_item_failures = []
    for record in event['Records']:
        try:
            process(record['body'])
        except Exception as e:
            batch_item_failures.append({
                "itemIdentifier": record['messageId']
            })
    return {"batchItemFailures": batch_item_failures}

```

Event source mapping settings:

- BatchSize: 1-10,000 (Standard), 1-10 (FIFO)
- MaximumBatchingWindowInSeconds: wait for more messages (0-300s)
- FunctionResponseType: ["ReportBatchItemFailures"]
- ScalingConfig.MaximumConcurrency: limit concurrent Lambdas

Tip: set Lambda DLQ + SQS DLQ for full failure capture

Q58. What is SQS vs Kinesis?

Difficulty: Hard

SQS:

Ordering: FIFO queue only, no ordering in Standard
 Consumers: competing consumers (each message processed by ONE consumer)
 Retention: up to 14 days
 Replay: NOT possible (consumed message deleted after ACK)
 Throughput: Standard: virtually unlimited; FIFO: 300 TPS
 Message size: max 256KB
 Use: work queue, task queue, microservice decoupling

Kinesis Data Streams:

Ordering: per shard (strict ordering)
 Consumers: multiple consumer groups (same data read by multiple)
 Retention: 24h (default) to 7 days (extended), 365 days (long-term)
 Replay: YES (re-read from any sequence number within retention)
 Throughput: 1MB/s write per shard, 2MB/s read per shard
 Message size: max 1MB per record
 Use: real-time analytics, event sourcing, multiple consumers

Decision:

SQS: "do this task once" (queue semantic)
 Kinesis: "record this event for multiple consumers / replay" (log semantic)

SQS → Kinesis mapping:

SQS consumer group ≈ Kinesis consumer group
 SQS message deleted ≠ Kinesis (pointer moves, data stays)
 SQS DLQ → Kinesis no built-in DLQ (use Lambda DLQ)

Q59. What is SQS Extended Client Library?

Difficulty: Medium

SQS message limit: 256KB

Large message solution: Amazon SQS Extended Client Library
 Stores payload in S3, sends S3 reference in SQS message

Flow:

Producer:

1. Upload payload to S3 (auto-generates key)
2. Send SQS message with S3 pointer: {"s3BucketName":"...", "s3Key":"..."}

Consumer:

1. Receive SQS message
2. Detect S3 pointer (special message attribute)
3. Download payload from S3
4. Process
5. Delete SQS message + S3 object

Java/Python: aws-sqs-java-extended-client-lib

Go: implement manually using same protocol

Manual Go implementation:

```
const maxSQSPayload = 250 * 1024 // 250KB to be safe
```

```
func sendLargeMessage(body []byte) error {
```

```

    if len(body) < maxSQSPayload {
        _, err := sqs.SendMessage(ctx, &sqs.SendMessageInput{
            QueueUrl:    &queueURL,
            MessageBody: aws.String(string(body)),
        })
        return err
    }
    // Upload to S3
    key := uuid.New().String()
    s3.PutObject(ctx, &s3.PutObjectInput{
        Bucket: aws.String(bucket),
        Key:    aws.String(key),
        Body:   bytes.NewReader(body),
    })
    // Send reference
    ref := fmt.Sprintf(`{"bucket":"%s","key":"%s"}`, bucket, key)
    _, err := sqs.SendMessage(ctx, &sqs.SendMessageInput{
        QueueUrl:    &queueURL,
        MessageBody: aws.String(ref),
    })
    return err
}

```

Q60. What is SQS message filtering with SNS?

Difficulty: Hard

SNS → SQS Fan-out with filtering:

```

Topic: OrderEvents
Queue1 (order-service): filter by "eventType": "order.created"
Queue2 (inventory):      filter by "eventType": ["order.created", "order.cancelled"]
Queue3 (analytics):      no filter (receives everything)

```

Subscription filter policy (JSON):

```

{
  "eventType": ["order.created"],
  "region": ["us-east-1", "us-west-2"],
  "amount": [{"numeric": [">=", 100]}]
}

```

Set via CLI:

```

aws sns set-subscription-attributes \
  --subscription-arn arn:aws:sns:... \
  --attribute-name FilterPolicy \
  --attribute-value '{"eventType":["order.created"]}'

```

SNS filter on message attributes:

Publisher sets attributes → SNS evaluates filter before delivery
 Filtered-out messages never reach SQS → cost savings, less noise

Advanced filter operators:

```

Exact match: "status": ["SHIPPED"]
Prefix: "prefix": "order-"
Numeric: {"numeric": [">", 0, "<=", 100]}
Anything-but: {"anything-but": ["ERROR"]}
Exists: {"exists": true}

```


Q61. What is SQS with AWS EventBridge?**Difficulty:**
Medium

EventBridge → SQS:

EventBridge: sophisticated event routing (rules-based)

Target: SQS queue

Use case:

EC2 instance terminates → EventBridge rule → SQS queue → Lambda cleanup

EventBridge advantages over direct SQS publish:

Content-based filtering (complex rules)

Multiple targets (SQS + Lambda + Step Functions simultaneously)

Archive + replay events

Schema registry

Cross-account and cross-region delivery

Rule example:

Event pattern:

```
{
  "source": ["aws.ec2"],
  "detail-type": ["EC2 Instance State-change Notification"],
  "detail": { "state": ["terminated"] }
}
```

Target: SQS queue arn:aws:sqs:us-east-1:...

Comparison:

SNS → SQS: fan-out, simple publish-subscribe, filter on attributes

EventBridge → SQS: complex routing, AWS service events, archive/replay

Both: asynchronous decoupling

EventBridge Pipes:

Source: SQS → enrichment (Lambda) → target (another SQS/API/Step Functions)

Point-to-point pipeline with transformation

Q62. What is SQS cost optimization?**Difficulty:**
Medium

SQS pricing:

Standard: \$0.40 per 1M requests (first 1M free/month)

FIFO: \$0.50 per 1M requests

Data transfer: free within same region

Optimization strategies:

1. Long polling (biggest impact):

Short polling: charges per poll (even empty response)

Long polling (WaitTimeSeconds=20): fewer polls, fewer charges

20s wait = 3 polls/min vs 60 polls/min = 95% cost reduction

2. Batching:

Send: batch up to 10 messages in one API call (charged as 1 request)

Receive: ReceiveMessage returns up to 10 at once

Delete: DeleteMessageBatch (up to 10) = 1 API call

10x cost reduction vs individual operations

3. Message deduplication cost:

FIFO charges for dedup even if message rejected (not sent)
Standard + app-level dedup: cheaper for high-volume

4. Avoid frequent polling of empty queues:
Use SQS → Lambda (Lambda handles polling, you pay per invocation)
5. Extended Client for large payloads:
S3 is cheaper than large SQS messages (over 64KB billing starts)
256KB SQS message = billed as $4 \times 64\text{KB}$ chunks = 4 requests!
6. Monitor:
NumberOfEmptyReceives: high → increase WaitTimeSeconds
ApproximateNumberOfMessages: high → scale consumers

Q63. What are SQS security best practices?

Difficulty:
Medium

1. Encryption:
SSE-SQS: AWS managed keys (default, free)
SSE-KMS: customer managed keys (better audit, cross-account)

Queue policy → Condition: aws:SecureTransport: true
→ Reject unencrypted (non-HTTPS) connections
2. IAM: least privilege
Producer: sqs:SendMessage only
Consumer: sqs:ReceiveMessage, sqs:DeleteMessage, sqs:GetQueueAttributes
Admin: full access (restricted)
3. Queue policy for cross-account:


```
{
  "Effect": "Allow",
  "Principal": {"AWS": "arn:aws:iam::ACCOUNT2:root"},
  "Action": ["sqs:SendMessage"],
  "Resource": "arn:aws:sqs:...:myqueue",
  "Condition": {
    "ArnLike": {"aws:SourceArn": "arn:aws:sns:...:mytopic"}
  }
}
```
4. VPC endpoint: keep traffic within AWS network
Gateway endpoint for S3/DynamoDB (free)
Interface endpoint for SQS (hourly charge)
→ Prevents SQS traffic from going over internet
5. Resource-based policy vs IAM:
IAM: attached to identity (role/user)
Resource policy: attached to queue, controls who can access
Both must allow for cross-account access
6. CloudTrail: log all SQS API calls (SendMessage, DeleteMessage, etc.)
Useful for: compliance, debugging, security audits

Q64. What are SQS consumer scaling strategies?

Difficulty: Hard

Scaling consumers based on queue depth:

1. CloudWatch + Auto Scaling:
 - Metric: ApproximateNumberOfMessages
 - Alarm: > 1000 messages → scale out
 - Scale policy: target tracking (keep queue depth / consumer = X)
2. Step scaling:
 - < 100 messages: 1 consumer
 - 100-1000: 2 consumers
 - > 1000: 4 consumers
3. Kubernetes HPA + KEDA:
 - KEDA: event-driven autoscaler for Kubernetes
 - ScaledObject triggers on SQS depth:

```

apiVersion: keda.sh/v1alpha1
kind: ScaledObject
spec:
  triggers:
  - type: aws-sqs-queue
    metadata:
      queueURL: https://sqs.us-east-1.amazonaws.com/.../queue
      queueLength: "100" # target messages per consumer
      awsRegion: us-east-1
      
```
4. Lambda (no explicit scaling needed):
 - Lambda auto-scales based on queue depth
 - Max concurrency configurable (ReservedConcurrency)
5. Go consumer scaling pattern:
 - Start N goroutines (workers)
 - Each long-polls independently
 - Scale workers based on ApproximateNumberOfMessages metric

Processing rate = workers × (messages per receive / processing time)
 Target: queue depth < MaxReceiveCount × processing time

Q65. What is SQS redrive policy configuration?

Difficulty:
Medium

```

// Redrive policy: configure DLQ for a source queue
// RedrivePolicy on source queue (not DLQ)
{
  "deadLetterTargetArn": "arn:aws:sqs:us-east-1:123:orders-dlq",
  "maxReceiveCount": "5"
}

// maxReceiveCount: after N failed processing attempts, send to DLQ
// A message is "failed" if:
//   - Consumer receives it but doesn't delete within visibility timeout
//   - Consumer explicitly doesn't delete (crash, exception)
// Lambda returns error → SQS makes message visible again (retry)

// SQS DLQ rules:
// - DLQ must be same type (FIFO → FIFO DLQ, Standard → Standard DLQ)
// - DLQ must be in same region
// - DLQ retention can be longer than source (extend it!)
// - Source queue's message retention period affects DLQ messages
  
```

```
// MessageRedriveStatus: check if message is in DLQ or source
// DLQ attributes available: sourceQueueArns

// Testing redrive policy:
// Set maxReceiveCount=1 temporarily in dev
// Observe which messages land in DLQ
// Check message body + attributes for debugging
// Fix bug, then move from DLQ back to source

// Redrive from DLQ back to source (Console or API):
aws sqs start-message-move-task \
  --source-arn arn:aws:sqs:us-east-1:123:orders-dlq \
  --destination-arn arn:aws:sqs:us-east-1:123:orders
```

Q66. What is SQS monitoring and alerting?

Difficulty:
Medium

Key SQS CloudWatch metrics:

ApproximateNumberOfMessages:

Total messages waiting (visible + in-flight + delayed)
Alert: > threshold (backlog building up)

ApproximateAgeOfOldestMessage:

Age in seconds of oldest message in queue
Alert: > SLA (e.g., > 300s = processing falling behind)
Best metric for: consumer lag, SLA violation

NumberOfMessagesSent: rate of incoming messages

NumberOfMessagesReceived: rate consumed

NumberOfMessagesDeleted: rate successfully processed

ApproximateNumberOfMessagesNotVisible: in-flight messages

NumberOfEmptyReceives: idle polls (consumers have nothing to do)

Alarms:

1. ApproximateAgeOfOldestMessage > 300 → PagerDuty alert
2. ApproximateNumberOfMessages > 10000 → scale out consumers
3. DLQ ApproximateNumberOfMessages > 0 → alert (messages failing)
4. NumberOfEmptyReceives > 80% → reduce consumers (over-provisioned)

Dashboard to build:

Queue depth (ApproximateNumberOfMessages)
Message age (ApproximateAgeOfOldestMessage)
Throughput (sent/received/deleted per minute)
DLQ depth (separate alarm, severity=high)
Consumer lag ratio

Tracing:

X-Ray: enable X-Ray tracing on Lambda + SQS
Shows: queue wait time, processing time, downstream calls

Q67. What is SQS with Step Functions?

Difficulty: Hard

SQS → Step Functions integration patterns:

1. SQS as input trigger:

SQS message → Lambda → start Step Functions execution

```

Lambda trigger:
def handler(event, context):
    for record in event['Records']:
        sfn.start_execution(
            stateMachineArn=STATE_MACHINE_ARN,
            input=record['body']
        )

```

2. Step Functions → SQS (send task):

State machine sends message to SQS

Task type: "Resource": "arn:aws:states:::sqs:sendMessage"

```

{
  "Type": "Task",
  "Resource": "arn:aws:states:::sqs:sendMessage.waitForTaskToken",
  "Parameters": {
    "QueueUrl": "...",
    "MessageBody": {
      "taskToken.$": "$$.Task.Token",
      "input.$": "$"
    }
  }
}

```

3. Wait for task token:

Step Functions sends token in SQS message

Worker receives message, processes, calls SendTaskSuccess/Failure

Step Functions resumes execution

Pattern: human approval, long-running external tasks

4. Express vs Standard workflow:

Express: short (< 5 min), high volume (perfect for SQS fan-out)

Standard: long-running (up to 1 year), exactly-once execution

Q68. What is SQS temporary queues?

Difficulty:
Medium

Temporary queues: short-lived queues for request-response over SQS

Pattern: replaces synchronous HTTP with async SQS

Temporary Queue Client (aws-sdk-java-sqs-java-messaging-lib):

Creates virtual queues backed by single physical queue

Each requester gets unique correlationId

Manual Go implementation:

```
const requestQueueURL = "https://sqs.../requests"
```

```

func makeRequest(payload string) (string, error) {
    // Create reply queue
    result, _ := sqs.CreateQueue(ctx, &sqs.CreateQueueInput{
        QueueName: aws.String("reply-" + uuid.New().String()),
        Attributes: map[string]string{
            "MessageRetentionPeriod": "60", // 1 minute
        },
    })
    replyQURL := result.QueueUrl
    defer sqs.DeleteQueue(ctx, &sqs.DeleteQueueInput{QueueUrl: replyQURL})
    // Send request with reply-to

```

```

sqs.SendMessage(ctx, &sqs.SendMessageInput{
    QueueUrl:    aws.String(requestQueueURL),
    MessageBody: aws.String(payload),
    MessageAttributes: map[string]sqsTypes.MessageAttributeValue{
        "replyTo": {StringValue: replyQURL, DataType: aws.String("String")},
    },
})

// Wait for response
for i := 0; i < 10; i++ {
    resp, _ := sqs.ReceiveMessage(ctx, &sqs.ReceiveMessageInput{
        QueueUrl: replyQURL, WaitTimeSeconds: 5,
    })
    if len(resp.Messages) > 0 {
        return *resp.Messages[0].Body, nil
    }
}
return "", errors.New("timeout")
}

```

Q69. What is SQS batching in Go?

Difficulty:
Medium

```

// Efficient batch operations

// Batch send (up to 10 per batch)
type BatchSender struct {
    sqs      *sqs.Client
    queueURL string
    mu       sync.Mutex
    batch    []sqsTypes.SendMessageBatchRequestEntry
}

func (s *BatchSender) Send(ctx context.Context, body string, attrs map[string]sqsTypes.MessageAttributeValue) error {
    entry := sqsTypes.SendMessageBatchRequestEntry{
        Id:            aws.String(uuid.New().String()),
        MessageBody:   aws.String(body),
        MessageAttributes: attrs,
    }

    s.mu.Lock()
    s.batch = append(s.batch, entry)
    if len(s.batch) >= 10 {
        batch := s.batch
        s.batch = nil
        s.mu.Unlock()
        return s.flush(ctx, batch)
    }
    s.mu.Unlock()
    return nil
}

func (s *BatchSender) flush(ctx context.Context, entries []sqsTypes.SendMessageBatchRequestEntry) error {
    result, err := s.sqs.SendMessageBatch(ctx, &sqs.SendMessageBatchInput{
        QueueUrl: &s.queueURL,
        Entries:  entries,
    })
    if err != nil {
        return err
    }
    return result.FailureMessages[0]
}

```

```

    if err != nil { return err }

    // Check failed entries
    for _, failed := range result.Failed {
        log.Printf("message %s failed: %s - %s", *failed.Id, *failed.Code, *failed.Message)
    }
    return nil
}

// Batch delete (after processing)
func deleteMessages(ctx context.Context, sqsClient *sqs.Client, queueURL string, receipts []string) error {
    entries := make([]sqsTypes.DeleteMessageBatchRequestEntry, len(receipts))
    for i, r := range receipts {
        entries[i] = sqsTypes.DeleteMessageBatchRequestEntry{
            Id:          aws.String(strconv.Itoa(i)),
            ReceiptHandle: aws.String(r),
        }
    }
    _, err := sqsClient.DeleteMessageBatch(ctx, &sqs.DeleteMessageBatchInput{
        QueueUrl: aws.String(queueURL),
        Entries:  entries,
    })
    return err
}

```

Q70. What is the SQS visibility timeout best practices?

Difficulty:
Medium

Visibility timeout = time consumer has to process + delete
 If not deleted in time → message reappears in queue (re-processed)

Formula:

$\text{visibility_timeout} > \text{p99_processing_time} + \text{buffer (20\%)}$

Example:

Processing time p99 = 25 seconds
 Set visibility_timeout = 30 seconds minimum

Extending visibility timeout for long tasks:

```

// Heartbeat: extend while processing
func processWithHeartbeat(ctx context.Context, msg Message) error {
    ticker := time.NewTicker(20 * time.Second) // renew every 20s
    defer ticker.Stop()

    done := make(chan error, 1)
    go func() { done <- process(msg) }()

    for {
        select {
        case <- ticker.C:
            sqs.ChangeMessageVisibility(ctx, &sqs.ChangeMessageVisibilityInput{
                QueueUrl:      &queueURL,
                ReceiptHandle: msg.ReceiptHandle,
                VisibilityTimeout: aws.Int32(30), // extend by 30 more seconds
            })
        case err := <- done:
            return err
        }
    }
}

```

```

    }
}
}

```

Queue-level vs per-message:

Queue: default visibility timeout (set during creation)

ChangeMessageVisibility: override for specific messages (max 12 hours)

MaxReceiveCount × visibility_timeout = max time before DLQ

Q71. ### Q71–Q100: Remaining SQS Questions

Difficulty:
Medium

Q72. What is SQS message body size and chunking?

Difficulty:
Medium

Limit: 256KB (262,144 bytes) per message

Includes: message body + message attributes

If payload > 256KB:

1. SQS Extended Client (Java/Python): auto-stores in S3
2. Manual: store payload in S3, put S3 reference in SQS
3. Split message into chunks (if order doesn't matter)
4. Consider Kinesis (1MB per record)

Billing for large messages:

64KB = 1 "request"

128KB = 2 requests

256KB = 4 requests

Combine small messages (batching) to save cost

Q73. What is SNS + SQS vs EventBridge?

Difficulty:
Medium

SNS + SQS fan-out:

Simple publish-subscribe

Message attribute filtering

Push delivery model (SNS pushes to SQS)

Best for: known, fixed subscriber set

EventBridge:

Content-based routing (JSON path filtering)

Native AWS service integration (100+ sources)

Schema registry + discovery

Archive + replay events

Cross-account/region delivery

Higher latency (~0.5s vs SNS <1ms)

Best for: complex routing, AWS service events, schema management

Choose SNS+SQS: simple fan-out, latency sensitive

Choose EventBridge: complex routing, audit/replay, AWS service events

Q74. What is SQS ApproximateNumberOfMessages lag calculation?

Difficulty:
Medium

```
// Calculate processing lag and projected catch-up time
```

```
type QueueMetrics struct {
```



```

    Depth            int64
    AgeOfOldest      int64 // seconds
    MessagesPerMinute float64
}

func checkQueueHealth(ctx context.Context, sqsClient *sqs.Client, queueURL string) (QueueMetrics, error) {
    result, err := sqsClient.GetQueueAttributes(ctx, &sqs.GetQueueAttributesInput{
        QueueUrl: &queueURL,
        AttributeNames: []sqsTypes.QueueAttributeName{
            sqsTypes.QueueAttributeNameApproximateNumberOfMessages,
            sqsTypes.QueueAttributeNameApproximateAgeOfOldestMessage,
        },
    })
    if err != nil { return QueueMetrics{}, err }

    depth, _ := strconv.ParseInt(result.Attributes["ApproximateNumberOfMessages"], 10, 64)
    age, _ := strconv.ParseInt(result.Attributes["ApproximateAgeOfOldestMessage"], 10, 64)

    return QueueMetrics{Depth: depth, AgeOfOldest: age}, nil
}

```

Q75. What is SQS server-side encryption with KMS?

Difficulty:
Medium

SSE with AWS KMS:

Messages encrypted at rest using KMS key

KMS key types:

AWS managed (aws/sqs): automatic rotation, no extra cost

Customer managed: control rotation, cross-account, audit

Cost: \$0.03 per 10,000 KMS API calls

Each send + each receive = 1 KMS call each

```

aws sqs set-queue-attributes \
  --queue-url $URL \
  --attributes KmsMasterKeyId=alias/my-key,\
               KmsDataKeyReusePeriodSeconds=300

```

Data key reuse (300s): reduces KMS calls (cached 5 minutes)

Higher reuse period = fewer KMS calls = lower cost

For Lambda: grant Lambda execution role kms:Decrypt on the key

For cross-account: key policy must allow other account

Q76. What is SQS delay queue per message?

Difficulty:
Medium

```

// Queue-level delay: DelaySeconds on queue (0-900 seconds)
// Message-level delay: override per message

// Use case: retry with delay, scheduled processing

result, err := sqsClient.SendMessage(ctx, &sqs.SendMessageInput{
    QueueUrl:    aws.String(queueURL),
    MessageBody: aws.String(body),
    DelaySeconds: aws.Int32(60), // this message visible after 60 seconds
})

```

```
// FIFO queues: don't support per-message DelaySeconds
// (queue-level delay only for FIFO)

// Pattern: exponential backoff retry with delay
func retryWithDelay(attempt int) error {
    delay := int32(math.Min(float64(1<<attempt), 900)) // 1,2,4,8... max 900s
    _, err := sqsClient.SendMessage(ctx, &sqs.SendMessageInput{
        QueueUrl:      aws.String(retryQueueURL),
        MessageBody:   aws.String(payload),
        DelaySeconds:   &delay,
        MessageAttributes: map[string]sqsTypes.MessageAttributeValue{
            "attempt": {StringValue: aws.String(strconv.Itoa(attempt)), DataType: aws.String("Number")},
        },
    })
    return err
}
```

Q77. ### Q76-Q100: Final SQS Questions Summary

Difficulty:
Medium

Q	Topic
Q76	SQS IAM policy examples with conditions
Q77	SQS with SageMaker for ML pipelines
Q78	SQS SDK retry configuration in Go
Q79	SQS message ordering guarantees
Q80	At-least-once vs at-most-once vs exactly-once
Q81	SQS vs RabbitMQ comparison
Q82	SQS vs Kafka comparison
Q83	SQS local development with LocalStack
Q84	SQS integration testing strategies
Q85	SQS consumer idempotency patterns
Q86	SQS message tracing with X-Ray
Q87	SQS tag-based cost allocation
Q88	SQS queue policies for S3 event notifications
Q89	SQS to trigger ECS tasks
Q90	Multi-region SQS architecture
Q91	SQS for saga pattern implementation
Q92	SQS for outbox pattern
Q93	SQS consumer graceful shutdown in Go
Q94	SQS with API Gateway direct integration
Q95	SQS producer backpressure patterns
Q96	SQS queue depth alerting runbook
Q97	SQS and circuit breaker pattern
Q98	SQS consumer health check implementation

Q99	SQS interview question: design order processing
Q100	SQS production deployment checklist

Extended Questions (Q78–Q100)

Q78. What is SQS dead-letter queue (DLQ) and redrive policy?

Difficulty:
Medium

```
{
  "deadLetterTargetArn": "arn:aws:sqs:us-east-1:123456789:my-queue-dlq",
  "maxReceiveCount": 3
}

// Message goes to DLQ after maxReceiveCount failed attempts
// DLQ must be same type (standard/FIFO) as source queue

// Create DLQ
dlqURL := createQueue("my-queue-dlq")

// Set redrive policy on source queue
svc.SetQueueAttributes(&sqs.SetQueueAttributesInput{
  QueueUrl: &sourceURL,
  Attributes: map[string]*string{
    "RedrivePolicy": aws.String(`{
      "deadLetterTargetArn": "arn:aws:sqs:us-east-1:123:my-queue-dlq",
      "maxReceiveCount": "3"
    }`),
  },
},
})

// Monitor DLQ: set CloudWatch alarm on ApproximateNumberOfMessagesVisible
// Alert when DLQ > 0 (messages failing)

// Redrive from DLQ back to source (after fixing bug)
svc.StartMessageMoveTask(&sqs.StartMessageMoveTaskInput{
  SourceArn:      &dlqARN,
  DestinationArn: &sourceARN,
  MaxNumberOfMessagesPerSecond: aws.Int64(100),
})
```

Q79. What is SQS message filtering with SNS subscription?

Difficulty:
Medium

```
// SNS → SQS with filter policies: only deliver matching messages
// Reduces noise, avoids processing irrelevant messages

// SNS subscription filter policy (JSON attribute matching)
filterPolicy := `{
  "event_type": ["order.placed", "order.cancelled"],
  "priority": [{"numeric": [">=", 5]}],
  "region": ["us-east-1", "eu-west-1"]
}`

// Attach to SNS subscription
```

```

svc.SetSubscriptionAttributes(&sns.SetSubscriptionAttributesInput{
    SubscriptionArn: &subARN,
    AttributeName:   aws.String("FilterPolicy"),
    AttributeValue:  aws.String(filterPolicy),
})

// Message attributes must match for delivery
svc.Publish(&sns.PublishInput{
    TopicArn: &topicARN,
    Message:  aws.String(body),
    MessageAttributes: map[string]*sns.MessageAttributeValue{
        "event_type": {DataType: aws.String("String"), StringValue: aws.String("order.placed")},
        "priority":   {DataType: aws.String("Number"), StringValue: aws.String("7")},
        "region":     {DataType: aws.String("String"), StringValue: aws.String("us-east-1")},
    },
})
// Only queues with matching filter receive this message

```

Q80. What is SQS long polling vs short polling?

Difficulty: Easy

```

// Short polling (default, WaitTimeSeconds=0):
// Returns immediately even if no messages
// May return empty responses → wastes API calls + $$$

// Long polling (WaitTimeSeconds=1-20):
// Waits up to N seconds for messages to arrive
// Fewer empty responses, lower cost, lower latency

// Recommended: always use long polling
result, err := svc.ReceiveMessage(&sqs.ReceiveMessageInput{
    QueueUrl:      &queueURL,
    WaitTimeSeconds: aws.Int64(20), // wait up to 20s
    MaxNumberOfMessages: aws.Int64(10),
})

// Queue-level default (set once, all consumers benefit)
svc.SetQueueAttributes(&sqs.SetQueueAttributesInput{
    QueueUrl: &queueURL,
    Attributes: map[string]*string{
        "ReceiveMessageWaitTimeSeconds": aws.String("20"),
    },
})

// Cost comparison:
// 1M requests short poll: $0.40
// Same throughput long poll: $0.04 (10x cheaper)
// Plus: reduces empty receives from 90% to <1%

```

Q81. What is SQS message visibility and worker crashes?

Difficulty: Hard

```

// If worker crashes after receiving but before deleting:
// Message becomes visible after VisibilityTimeout → reprocessed
// This is SQS's at-least-once delivery guarantee

// Problem: long-running jobs where VisibilityTimeout < processing time
// Message reappears while still being processed → duplicate processing!

```

```
// Solution: extend visibility timeout during processing
func processWithHeartbeat(ctx context.Context, msg *sqs.Message) error {
    done := make(chan struct{})

    // Heartbeat goroutine: extend visibility every 30s
    go func() {
        ticker := time.NewTicker(30 * time.Second)
        defer ticker.Stop()
        for {
            select {
            case <-done:
                return
            case <-ticker.C:
                svc.ChangeMessageVisibility(&sqs.ChangeMessageVisibilityInput{
                    QueueUrl:      &queueURL,
                    ReceiptHandle:    msg.ReceiptHandle,
                    VisibilityTimeout: aws.Int64(60), // extend by 60s
                })
            }
        }
    }()

    err := doLongProcessing(ctx, msg)
    close(done)
    return err
}
```

Q82. What is SQS FIFO queue ordering guarantees?

Difficulty: Hard

```
// FIFO: exactly-once processing within message group
// MessageGroupId: ordering scope (like Kafka partition key)
// MessageDeduplicationId: deduplication within 5-minute window

// Send to FIFO queue
svc.SendMessage(&sqs.SendMessageInput{
    QueueUrl:      aws.String("https://sqs.us-east-1.amazonaws.com/123/orders.fifo"),
    MessageBody:    aws.String(orderJSON),
    MessageGroupId: aws.String(fmt.Sprintf("customer-%d", customerID)),
    MessageDeduplicationId: aws.String(idempotencyKey),
})

// Ordering guarantee:
// SAME MessageGroupId → strict FIFO order
// DIFFERENT MessageGroupId → independent ordering

// Throughput:
// Standard: nearly unlimited TPS
// FIFO: 300 TPS (or 3000 with batching) per queue
// FIFO with high throughput mode: 30,000 TPS

// Deduplication:
// Content-based: SHA-256 of MessageBody (enable on queue)
// Explicit: provide MessageDeduplicationId
// Window: 5 minutes – same ID within 5 min = duplicate (ignored)
```

Q83. What is SQS with Lambda (event source mapping)?**Difficulty:**
Medium

```
// Lambda: triggered by SQS automatically
// Batch size: 1-10000 messages per invocation
// Lambda polls SQS, manages concurrency automatically

// Lambda function handler
func handler(ctx context.Context, event events.SQSEvent) error {
    var failures []events.SQSBatchItemFailure

    for _, record := range event.Records {
        if err := processRecord(ctx, record); err != nil {
            // Report partial failure (only failed messages go to DLQ)
            failures = append(failures, events.SQSBatchItemFailure{
                ItemIdentifier: record.MessageId,
            })
        }
    }

    if len(failures) > 0 {
        return events.SQSBatchResponse{BatchItemFailures: failures}
    }
    return nil
}

// Event source mapping config
{
    "EventSourceArn": "arn:aws:sqs:...:my-queue",
    "FunctionName": "my-processor",
    "BatchSize": 10,
    "MaximumBatchingWindowInSeconds": 5,
    "FunctionResponseType": ["ReportBatchItemFailures"],
    "ScalingConfig": { "MaximumConcurrency": 100 }
}
```

Q84. What is SQS cost optimization?**Difficulty:**
Medium

SQS pricing (us-east-1):
 First 1M requests/month: FREE
 Standard: \$0.40 per million requests
 FIFO: \$0.50 per million requests
 Data transfer in: free
 Data transfer out: first 100GB free

Cost optimizations:

1. Long polling (biggest saving)
 WaitTimeSeconds=20 → 90% fewer empty receives
 100K messages/day short poll: ~\$0.04
 100K messages/day long poll: ~\$0.004
2. Batch operations
 SendMessageBatch: up to 10 messages = 1 request (10x savings)
 ReceiveMessage: up to 10 messages = 1 request
 DeleteMessageBatch: up to 10 = 1 request

3. Message size

Requests billed per 64KB chunk

100KB message = 2 requests

Use SQS Extended Client for messages >256KB (stores in S3)

4. Right-size visibility timeout

Too short → message reprocessed → duplicate requests

Too long → failed messages take long to retry

5. Monitor and alarm on DLQ

Don't let failed messages accumulate (they still cost)

Q85. What is SQS vs SNS vs EventBridge comparison?

Difficulty:
Medium

SQS (Simple Queue Service):

Pattern: point-to-point queue (producer → consumer)

Retention: up to 14 days

Consumers: one consumer group processes each message

Use: task queues, job processing, decoupling services

Example: order processing worker reads from order queue

SNS (Simple Notification Service):

Pattern: pub/sub fan-out (topic → N subscribers)

Retention: no retention (fire-and-forget)

Consumers: ALL subscribers receive EVERY message

Use: notifications, fan-out to multiple queues/lambda

Example: order.placed → email queue + analytics queue + inventory queue

EventBridge (formerly CloudWatch Events):

Pattern: event bus with routing rules

Retention: can archive and replay events

Consumers: rules route events to targets (Lambda, SQS, Step Functions)

Use: complex event routing, cross-account, scheduled events

Example: route EC2 events to different handlers based on event type

Common patterns:

SNS → SQS: fan-out to multiple queues (each processes independently)

EventBridge → SQS: complex routing based on event content

SQS → Lambda: serverless consumer

Q86. What is SQS poison pill message handling?

Difficulty: Hard

```
// Poison pill: message that always fails processing
```

```
// Without DLQ: infinite retry loop, blocks queue
```

```
// Solution 1: DLQ (primary defense)
```

```
// Set maxReceiveCount=3 → after 3 failures → moves to DLQ
```

```
// Solution 2: Track receive count in application
```

```
func processMessage(msg *sqs.Message) error {
```

```
    receiveCount, _ := strconv.Atoi(
```

```
        *msg.Attributes["ApproximateReceiveCount"])
```

```
    if receiveCount > 3 {
```

```
        // Too many retries – log and delete (don't send to DLQ)
```

```
        log.Printf("POISON PILL: msg=%s receive_count=%d body=%s",
```

```

        *msg.MessageId, receiveCount, *msg.Body)
    sendAlert(msg) // alert on-call
    deleteMessage(msg) // remove from queue
    return nil
}

return actualProcessing(msg)
}

// Solution 3: Exponential backoff via ChangeMessageVisibility
func handleWithBackoff(msg *sqs.Message, attempt int) {
    if err := process(msg); err != nil {
        backoff := math.Pow(2, float64(attempt)) * 10 // 10, 20, 40, 80s
        svc.ChangeMessageVisibility(&sqs.ChangeMessageVisibilityInput{
            QueueUrl:      &queueURL,
            ReceiptHandle: msg.ReceiptHandle,
            VisibilityTimeout: aws.Int64(int64(backoff)),
        })
    }
}

```

Q87. What is SQS server-side encryption?

Difficulty: Easy

```

// SSE (Server-Side Encryption): encrypt messages at rest
// Uses AWS KMS (Key Management Service)

// Enable SSE on queue creation
svc.CreateQueue(&sqs.CreateQueueInput{
    QueueName: aws.String("secure-queue"),
    Attributes: map[string]*string{
        "KmsMasterKeyId":      aws.String("alias/aws/sqs"), // AWS managed key
        // OR:
        "KmsMasterKeyId":      aws.String("arn:aws:kms:us-east-1:123:key/abc"), // CMK
        "KmsDataKeyReusePeriodSeconds": aws.String("300"), // reuse data key for 5 min
    },
})

// SQS SSE uses envelope encryption:
// 1. KMS generates data key (DEK)
// 2. DEK encrypts message
// 3. KMS master key encrypts DEK
// 4. Encrypted DEK stored with message

// Cost: KMS API calls ($0.03/10K)
// KmsDataKeyReusePeriodSeconds: reuse DEK to reduce KMS calls

// In-transit encryption: TLS by default (HTTPS endpoint)
// SQS endpoint: https://sqs.us-east-1.amazonaws.com (always HTTPS)

```

Q88. What is SQS batch processing in Go?

Difficulty:
Medium

```

// Process messages efficiently with batching

type SQSWorker struct {
    svc      *sqs.SQS
    queueURL string
}

```



```

    handler func(msg *sqs.Message) error
}

func (w *SQSWorker) Run(ctx context.Context) {
    for {
        select {
        case <-ctx.Done():
            return
        default:
            msgs, err := w.receive()
            if err != nil { time.Sleep(time.Second); continue }

            w.processBatch(msgs)
        }
    }
}

func (w *SQSWorker) processBatch(msgs []*sqs.Message) {
    var (
        wg      sync.WaitGroup
        mu      sync.Mutex
        deletes []*sqs.DeleteMessageBatchRequestEntry
    )

    for _, msg := range msgs {
        wg.Add(1)
        go func(m *sqs.Message) {
            defer wg.Done()
            if err := w.handler(m); err != nil {
                log.Printf("failed msg %s: %v", *m.MessageId, err)
                return
            }
            mu.Lock()
            deletes = append(deletes, &sqs.DeleteMessageBatchRequestEntry{
                Id:          m.MessageId,
                ReceiptHandle: m.ReceiptHandle,
            })
            mu.Unlock()
        }(msg)
    }

    wg.Wait()

    // Batch delete successful messages
    if len(deletes) > 0 {
        w.svc.DeleteMessageBatch(&sqs.DeleteMessageBatchInput{
            QueueUrl: &w.queueURL,
            Entries:  deletes,
        })
    }
}

```

Q89. What is SQS queue depth monitoring?

Difficulty:
Medium

```

// Key metrics to monitor:
// 1. ApproximateNumberOfMessagesVisible: backlog size

```

```
// 2. ApproximateNumberOfMessagesNotVisible: in-flight (being processed)
// 3. ApproximateAgeOfOldestMessage: how stale is the oldest message?
// 4. NumberOfMessagesSent / NumberOfMessagesDeleted: throughput

// Get queue attributes
result, _ := svc.GetQueueAttributes(&sqs.GetQueueAttributesInput{
    QueueUrl: &queueURL,
    AttributeNames: []*string{
        aws.String("ApproximateNumberOfMessages"),
        aws.String("ApproximateNumberOfMessagesNotVisible"),
        aws.String("ApproximateAgeOfOldestMessage"),
    },
})

visible, _ := strconv.Atoi(*result.Attributes["ApproximateNumberOfMessages"])
inFlight, _ := strconv.Atoi(*result.Attributes["ApproximateNumberOfMessagesNotVisible"])
ageSeconds, _ := strconv.Atoi(*result.Attributes["ApproximateAgeOfOldestMessage"])

// CloudWatch alarms:
// Queue depth > 1000 → scale up consumers
// OldestMessage age > 5 min → consumers stuck/crashed
// DLQ depth > 0 → processing failures (page on-call)

// Auto-scaling consumers:
// Target tracking: keep ApproximateNumberOfMessages / running tasks = N
// Scale out when backlog growing, scale in when draining
```

Q90. What is SQS temporary queues (virtual queues)?

Difficulty: Hard

```
// Temporary queues: short-lived queues for request-response pattern
// Without creating thousands of real SQS queues

// SQS Temporary Queue Client (AWS library)
// Creates a "virtual" queue on top of a host queue
// Multiplexes multiple virtual queues over one real queue

// Request-response pattern with SQS
func requestResponse(ctx context.Context, requestQueue string, payload string) (string, error) {
    // Create temporary reply queue
    replyQueueURL := createTempQueue()
    defer deleteTempQueue(replyQueueURL)

    // Send request with reply-to address
    svc.SendMessage(&sqs.SendMessageInput{
        QueueUrl: &requestQueue,
        MessageBody: aws.String(payload),
        MessageAttributes: map[string]*sqs.MessageAttributeValue{
            "ReplyTo": {
                DataType:    aws.String("String"),
                StringValue: &replyQueueURL,
            },
            "CorrelationId": {
                DataType:    aws.String("String"),
                StringValue: aws.String(uuid.New().String()),
            },
        },
    })
}
```

```

    // Wait for response
    ctx, cancel := context.WithTimeout(ctx, 30*time.Second)
    defer cancel()
    return waitForResponse(ctx, replyQueueURL)
}

```

Q91. What is SQS access control with IAM?

Difficulty:
Medium

```

{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {"AWS": "arn:aws:iam::123456789:role/producer-role"},
      "Action": ["sqs:SendMessage"],
      "Resource": "arn:aws:sqs:us-east-1:123:my-queue"
    },
    {
      "Effect": "Allow",
      "Principal": {"AWS": "arn:aws:iam::123456789:role/consumer-role"},
      "Action": ["sqs:ReceiveMessage", "sqs:DeleteMessage", "sqs:ChangeMessageVisibility"],
      "Resource": "arn:aws:sqs:us-east-1:123:my-queue"
    },
    {
      "Effect": "Allow",
      "Principal": {"Service": "sns.amazonaws.com"},
      "Action": "sqs:SendMessage",
      "Resource": "arn:aws:sqs:us-east-1:123:my-queue",
      "Condition": {
        "ArnEquals": {"aws:SourceArn": "arn:aws:sns:us-east-1:123:my-topic"}
      }
    }
  ]
}

```

Q92. What is SQS message attributes?

Difficulty:
Medium

```

// Message attributes: metadata alongside message body
// Up to 10 attributes per message
// Types: String, Number, Binary, String.Array

svc.SendMessage(&sqs.SendMessageInput{
  QueueUrl:    &queueURL,
  MessageBody: aws.String(body),
  MessageAttributes: map[string]*sqs.MessageAttributeValue{
    "EventType": {
      DataType:    aws.String("String"),
      StringValue: aws.String("order.placed"),
    },
    "OrderId": {
      DataType:    aws.String("Number"),
      StringValue: aws.String("42"),
    },
    "TraceId": {
      DataType:    aws.String("String"),

```

```

        StringValue: aws.String(traceID),
    },
},
})

// Receive with attributes
svc.ReceiveMessage(&sqs.ReceiveMessageInput{
    QueueUrl:      &queueURL,
    MessageAttributeNames: []*string{aws.String("All")},
    AttributeNames:  []*string{aws.String("All")},
})
// msg.MessageAttributes["EventType"].StringValue
// msg.Attributes["ApproximateReceiveCount"]

```

Q93. What is SQS delay queues?

Difficulty:
Medium

```

// Delay queue: message invisible for N seconds after send
// Use: schedule future processing, rate limiting, retry with delay

// Queue-level delay (all messages delayed)
svc.CreateQueue(&sqs.CreateQueueInput{
    QueueName: aws.String("delayed-queue"),
    Attributes: map[string]*string{
        "DelaySeconds": aws.String("900"), // 15 min delay
    },
})

// Per-message delay (overrides queue default)
svc.SendMessage(&sqs.SendMessageInput{
    QueueUrl:      &queueURL,
    MessageBody:   aws.String(body),
    DelaySeconds:  aws.Int64(300), // 5 min delay
})

// Max delay: 900 seconds (15 minutes)
// Use case: send "order reminder" email 1 hour after order
// Use case: retry failed webhook after exponential backoff
// Use case: deferred payment processing

// Note: FIFO queues don't support per-message delay
// Only queue-level delay for FIFO

```

Q94. What is SQS vs Kafka: when to use which?

Difficulty:
Medium

SQS:

- Simple setup (fully managed, no ops)
- Auto-scaling consumers
- Pay per use (no idle cost)
- Native AWS integrations (Lambda, SNS, EventBridge)
- At-least-once delivery built-in (DLQ)
- No message replay (once consumed + deleted = gone)
- 256KB message size limit
- No consumer offset tracking
- Max retention: 14 days

Use: task queues, job processing, simple async decoupling

Kafka:

- Message replay (consumers can re-read history)
 - Very high throughput (millions/sec)
 - Long retention (unlimited with tiered storage)
 - Multiple independent consumer groups
 - Stream processing (Kafka Streams, Flink)
 - Operational complexity (or pay for Confluent/MSK)
 - No automatic scaling (manage partitions)
 - Complex consumer offset management
- Use: event sourcing, audit logs, stream processing, analytics

Decision:

< 1M msg/day, AWS-native, simple consumers → SQS
 Event replay needed, analytics, high throughput → Kafka
 Mixed: use SQS for tasks, Kafka for events

Q95. What is SQS extended client for large messages?

Difficulty:
Medium

```
// SQS max message size: 256KB
// Extended client: store large payloads in S3, pointer in SQS

import "github.com/awsdocs/amazon-sqs-developer-guide"

// All messages → S3
extClient := sqsextendedclient.New(svc, s3Svc, "my-bucket",
    sqsextendedclient.AlwaysUseS3(true))

// Only messages > threshold → S3
extClient := sqsextendedclient.New(svc, s3Svc, "my-bucket",
    sqsextendedclient.PayloadSizeThreshold(100*1024)) // 100KB threshold

// Send (auto-uploads to S3 if needed)
extClient.SendMessage(&sqs.SendMessageInput{
    QueueUrl:    &queueURL,
    MessageBody: aws.String(largePayload),
})

// Receive (auto-downloads from S3)
result, _ := extClient.ReceiveMessage(&sqs.ReceiveMessageInput{
    QueueUrl: &queueURL,
})
// msg.Body contains original large payload (downloaded from S3)

// Cleanup: extended client deletes S3 object when SQS message deleted
extClient.DeleteMessage(&sqs.DeleteMessageInput{
    QueueUrl:    &queueURL,
    ReceiptHandle: msg.ReceiptHandle,
})
```

Q96. What is SQS cross-account access?

Difficulty:
Medium

```
// Account A: producer
// Account B: queue owner

// Step 1: Account B adds resource policy to allow Account A
```

```

policy := `{
    "Statement": [{
        "Effect": "Allow",
        "Principal": {"AWS": "arn:aws:iam::ACCOUNT-A:root"},
        "Action": "sqs:SendMessage",
        "Resource": "arn:aws:sqs:us-east-1:ACCOUNT-B:shared-queue"
    }]
}`

// Step 2: Account A assumes role or uses direct access
// Producer in Account A:
sess, _ := session.NewSession(&aws.Config{
    Region: aws.String("us-east-1"),
})
svc := sqs.New(sess)
svc.SendMessage(&sqs.SendMessageInput{
    QueueUrl:    aws.String("https://sqs.us-east-1.amazonaws.com/ACCOUNT-B/shared-queue"),
    MessageBody: aws.String(body),
})

// Cross-account patterns:
// Org-wide event bus: central SQS queue receives from all accounts
// SNS cross-account: topic in one account, SQS subscription in another

```

Q97. What is SQS approximate vs exact message counts?

Difficulty:
Medium

SQS counts are APPROXIMATE:

- ApproximateNumberOfMessages: may lag by seconds
- ApproximateNumberOfMessagesNotVisible: may be inaccurate
- ApproximateAgeOfOldestMessage: approximate

Why approximate?

- SQS is distributed across multiple servers
- Counting is eventually consistent
- Trade-off: high availability over exact consistency

Implications:

- Don't use for billing/exact counts
- Use for: autoscaling decisions (approximate is fine)
- Don't rely on count=0 for "queue is empty" (use polling)

Getting better accuracy:

- Purge + wait 60s → count more accurate
- Use CloudWatch metrics (15-min resolution)
- Count in your application code (track sent/processed)

FIFO queue counts:

- More accurate because FIFO queues don't use same distributed model
- But still marked as "approximate" in docs

Q98. What is SQS with AWS CDK/Terraform?

Difficulty:
Medium

```

// AWS CDK (TypeScript)
import * as sqs from 'aws-cdk-lib/aws-sqs';

const dlq = new sqs.Queue(this, 'OrdersDLQ', {

```

```

    retentionPeriod: Duration.days(14),
});

const queue = new sqs.Queue(this, 'OrdersQueue', {
    visibilityTimeout: Duration.seconds(300),
    retentionMessagePeriod: Duration.days(4),
    deadLetterQueue: {
        queue: dlq,
        maxReceiveCount: 3,
    },
    encryption: sqs.QueueEncryption.KMS_MANAGED,
});

# Terraform
resource "aws_sqs_queue" "orders_dlq" {
    name                = "orders-dlq"
    message_retention_seconds = 1209600 # 14 days
}

resource "aws_sqs_queue" "orders" {
    name                = "orders"
    visibility_timeout_seconds = 300
    message_retention_seconds = 345600 # 4 days
    redrive_policy = jsonencode({
        deadLetterTargetArn = aws_sqs_queue.orders_dlq.arn
        maxReceiveCount     = 3
    })
    kms_master_key_id = "alias/aws/sqs"
}

```

Q99. What is SQS message deduplication strategies?

Difficulty:
Medium

```

// Problem: SQS is at-least-once → consumer may receive duplicates
// Solution: idempotent consumers

// Strategy 1: Database unique constraint
func processOrder(ctx context.Context, msg *sqs.Message) error {
    var order Order
    json.Unmarshal([]byte(*msg.Body), &order)

    _, err := db.ExecContext(ctx,
        "INSERT INTO processed_messages (message_id) VALUES ($1) ON CONFLICT DO NOTHING",
        *msg.MessageId,
    )
    if err != nil { return err }

    // Check if already processed (0 rows inserted = duplicate)
    // ... or use INSERT ... ON CONFLICT DO NOTHING RETURNING id
    return processOrderBusiness(ctx, order)
}

// Strategy 2: Redis SET NX
func isProcessed(ctx context.Context, messageID string) bool {
    result, _ := rdb.SetNX(ctx, "processed:"+messageID, "1", 24*time.Hour).Result()
    return !result // false = not processed (we set it), true = already exists
}

```

```
// Strategy 3: FIFO queue with MessageDeduplicationId
// SQS itself deduplicates (5-minute window)
// Best for: preventing duplicates at source

// Strategy 4: Outbox + message ID
// Producer writes message_id to outbox → consumer checks processed table
```

Q100. What is SQS production deployment checklist?

Difficulty:
Medium

Before go-live with SQS in production:

Queue configuration:

- Appropriate VisibilityTimeout (> max processing time)
- DLQ configured with maxReceiveCount=3-5
- Message retention period set (default 4 days)
- Long polling enabled (WaitTimeSeconds=20)
- SSE enabled (KmsMasterKeyId set)
- Resource policy: least privilege

Consumer:

- Idempotent message processing
- Batch processing (up to 10 messages)
- Heartbeat for long-running jobs (ChangeMessageVisibility)
- Graceful shutdown (finish in-flight before exit)
- Error handling (don't delete on failure)
- Structured logging with message IDs

Monitoring:

- CloudWatch alarm: ApproximateAgeOfOldestMessage > threshold
- CloudWatch alarm: DLQ depth > 0
- CloudWatch alarm: consumer error rate
- Queue depth dashboard for capacity planning

Operations:

- DLQ redrive procedure documented
- Runbook for consumer outage
- Autoscaling policy based on queue depth
- IAM roles (not access keys) for authentication